

## Practice Test 3

Completed on 14-June-2020



### Attempt

01



### Marks Obtained

0 / 55



### Your score

0.0%



### Time Taken

00 H 00 M 03 S



### Result

Failed

## Domains wise Quiz Performance Report

1

Determine workload requirements

13

0

0

13

0

2

Design a business continuity strategy

3

0

0

3

0

3

Design an infrastructure strategy

21

0

0

21

0

4

Design for deployment, migration, and integration

7

0

0

7

0

5

Design for identity and security

9

0

0

9

0

6

Design a data platform solution

2

0

0

2

0

Total

All Domain

55

0

0

55

0

## Review the Answers

Sorting by

All



Question 1

Unattempted

A company has setup a set of Virtual Machines in a network in Azure. They have connected the Virtual Network to their on-premise network using ExpressRoute. There is an issue with an application hosted on the Virtual Machines in the network. A team needs to inspect the packets flowing into the Virtual Machines. The company decides to use Azure Advisor for the packet analysis. Would this fulfil the requirement?

- A. Yes
- B. No

#### Explanation:

Answer – B

Azure Advisor is a recommendations-based tool and can't be used to perform the packet analysis.

For more information on Azure Advisor, please visit the below URL

<https://docs.microsoft.com/en-us/azure/advisor/advisor-overview>

[Ask our Experts](#)

Rate this Question?  

[View Queries](#)

open ▾

Question 2

Unattempted

Domain :Determine workload requirements

A company has setup a set of Virtual Machines in a network in Azure. They have connected the Virtual Network to their on-premise network using ExpressRoute. There is an issue with an application hosted on the Virtual Machines in the network. A team needs to inspect the packets flowing into the Virtual Machines. The company decides to use Azure Traffic Analytics for the packet analysis. Would this fulfil the requirement?

- A. Yes
- B. No

#### Explanation:

Answer – A

Azure Traffic Analytics is used for traffic monitoring and is the correct solution to this business scenario.

The Microsoft documentation mentions the following

# Traffic Analytics

06/15/2018 • 14 minutes to read • Contributors  all

Traffic Analytics is a cloud-based solution that provides visibility into user and application activity in cloud networks. Traffic analytics analyzes Network Watcher network security group (NSG) flow logs to provide insights into traffic flow in your Azure cloud. With traffic analytics, you can:

- Visualize network activity across your Azure subscriptions and identify hot spots.
- Identify security threats to, and secure your network, with information such as open-ports, applications attempting internet access, and virtual machines (VM) connecting to rogue networks.
- Understand traffic flow patterns across Azure regions and the internet to optimize your network deployment for performance and capacity.
- Pinpoint network misconfigurations leading to failed connections in your network.

For more information on Traffic Analysis, please visit the below URL

<https://docs.microsoft.com/en-us/azure/network-watcher/traffic-analytics>

[Ask our Experts](#)

Rate this Question?  

[View Queries](#)

open ▾

Question 3

Unattempted

Domain :Determine workload requirements

A company has setup a set of Virtual Machines in a network in Azure. They have connected the Virtual Network to their on-premise network using ExpressRoute. There is an issue with an application hosted on the Virtual Machines in the network. A team needs to inspect the packets flowing into the Virtual Machines. The company decides to use Azure Traffic Manager for the packet analysis. Would this fulfil the requirement?

A. Yes

B. No

## Explanation:

Answer – B

Azure Traffic Manager is a DNS based load balancing tool and can't be used for the purpose of network capture.

# What is Traffic Manager?

03/05/2019 • 2 minutes to read • Contributors 

Azure Traffic Manager is a DNS-based traffic load balancer that enables you to distribute traffic optimally to services across global Azure regions, while providing high availability and responsiveness.

Traffic Manager uses DNS to direct client requests to the most appropriate service endpoint based on a traffic-routing method and the health of the endpoints. An endpoint is any Internet-facing service hosted inside or outside of Azure. Traffic Manager provides a range of [traffic-routing methods](#) and [endpoint monitoring options](#) to suit different application needs and automatic failover models. Traffic Manager is resilient to failure, including the failure of an entire Azure region.

For more information on Azure Traffic Manager, please visit the below URL

<https://docs.microsoft.com/en-us/azure/traffic-manager/traffic-manager-overview>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 4

Unattempted

Domain :Design a business continuity strategy

A company is planning on deploying an Azure Web App to 2 regions.

One of the key requirements is to ensure that the web app is always running if an Azure region fails. You need to ensure deployment costs are minimized.

Which of the following service would you include in the deployment of the solution?

- A. Azure Functions
- B. Azure Traffic Manager 
- C. Azure Application Gateway
- D. Azure Load Balancer

## Explanation:

Answer – B

You can use the Azure Traffic Manager to switch traffic over to a failover region. The Microsoft documentation mentions the following on the Azure Traffic Manager

# What is Traffic Manager?

03/05/2019 • 2 minutes to read • Contributors 

Azure Traffic Manager is a DNS-based traffic load balancer that enables you to distribute traffic optimally to services across global Azure regions, while providing high availability and responsiveness.

Traffic Manager uses DNS to direct client requests to the most appropriate service endpoint based on a traffic-routing method and the health of the endpoints. An endpoint is any Internet-facing service hosted inside or outside of Azure. Traffic Manager provides a range of [traffic-routing methods](#) and [endpoint monitoring options](#) to suit different application needs and automatic failover models. Traffic Manager is resilient to failure, including the failure of an entire Azure region.

Option A is incorrect since this is a compute service in Azure

Options C and D are incorrect since these services can't be used in region failover scenarios

For more information on Azure Traffic Manager, please visit the below URL

<https://docs.microsoft.com/en-us/azure/traffic-manager/traffic-manager-overview>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 5

Unattempted

Domain :Design a business continuity strategy

A company is planning on deploying an Azure Web App to 2 regions.

One of the key requirements is to ensure that the web app is always running if the primary region fails. You need to ensure deployment costs are minimized.

Which of the following feature would be used to ensure failover in the service?

- A. Session Affinity
- B. Performance based routing
- C. Priority Routing 
- D. URL routing

**Explanation:**

Answer – C

You can change the routing method for the Traffic Manager to the Priority routing method for implementing failover. The Microsoft documentation mentions the following on the routing method.

# Traffic Manager routing methods

09/17/2018 • 12 minutes to read • Contributors  all

Azure Traffic Manager supports six traffic-routing methods to determine how to route network traffic to the various service endpoints. For any profile, Traffic Manager applies the traffic-routing method associated to it to each DNS query it receives. The traffic-routing method determines which endpoint is returned in the DNS response.

The following traffic routing methods are available in Traffic Manager:

- **Priority:** Select **Priority** when you want to use a primary service endpoint for all traffic, and provide backups in case the primary or the backup endpoints are unavailable.

Options A and D are incorrect since these features are more pertinent to the Application gateway.

Option B is incorrect since the priority routing method needs to be used for the Azure Traffic Manager.

For more information on Azure Traffic Manager routing methods, please visit the below URL

<https://docs.microsoft.com/en-us/azure/traffic-manager/traffic-manager-routing-methods>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 6

Unattempted

Domain :Design an infrastructure strategy

A company currently has an on-premise network with an IP address space of 186.16.0.0/16. The company is going to deploy 20 Virtual machines to Azure. The Virtual machines will be placed in a subnet in an Azure virtual network. The requirement is to ensure the on-premise servers can communicate with the virtual machines hosted in Azure via a site-to-site VPN connection. You have to design the subnet for the virtual network in Azure which will be used to host the virtual machines.

Which of the following address space would you assign for the subnet in the Virtual Network?

- A. **186.16.0.0/16**
- B. **186.16.1.0/28**
- C. **192.168.0.0/24**
- D. **192.168.1.0/28**

## Explanation:

Answer – C

The address space for the Virtual Network should not conflict with the address space for the on-premise network. So, in this case the ideal option to choose as the address space is 192.168.0.0/24. Also, if we look at the question it clearly mentions about having "**20 VMs will be deployed in Azure**". So taking this into consideration Option C is the correct answer, reason we get **256 IP addresses** to work with.

A note on this is also given in the Microsoft documentation

## 1. Create a virtual network

To create a VNet in the Resource Manager deployment model by using the Azure portal, follow the steps below. Use the [example values](#) if you are using these steps as a [tutorial](#). If you are not doing these steps as a tutorial, be sure to replace the values with your own. For more information about working with virtual networks, see the [Virtual Network Overview](#).

### ⓘ Note

In order for this VNet to connect to an on-premises location you need to coordinate with your on-premises network administrator to carve out an IP address range that you can use specifically for this virtual network. If a duplicate address range exists on both sides of the VPN connection, traffic does not route the way you may expect it to. Additionally, if you want to connect this VNet to another VNet, the address space cannot overlap with other VNet. Take care to plan your network configuration accordingly.

Options A and B are incorrect since these address spaces would conflict with the on-premise address space.

Option D is incorrect since this address space should ideally be used for the gateway subnet

For more information on site-to-site VPN, please visit the below URL

<https://docs.microsoft.com/en-us/azure/vpn-gateway/vpn-gateway-howto-site-to-site-resource-manager-portal>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 7

Unattempted

A company currently has an on-premise network with an IP address space of 186.16.0.0/16. The company is going to deploy 20 Virtual machines to Azure. The Virtual machines will be placed in a subnet in an Azure virtual network. The requirement is to ensure the on-premise servers can communicate with the virtual machines hosted in Azure via a site-to-site VPN connection. You have to design the subnet for the virtual network in Azure which will be used to host the virtual machines.

Which of the following address space would you assign for the gateway subnet in the Virtual Network?

- A. **186.16.0.0/16**
- B. **186.16.1.0/28**
- C. **192.168.0.0/24**
- D. **192.168.1.0/28**

#### Explanation:

Answer – D

The address space for the Virtual Network should not conflict with the address space for the on-premise network. So, in this case the ideal option to choose as the address space is 192.168.0.0/24 for the subnet in the virtual network. And then use 192.168.1.0/28 as the address space for the gateway subnet.

The Microsoft documentation mentions the following on the gateway subnet

### 3. Create the gateway subnet

The virtual network gateway uses specific subnet called the gateway subnet. The gateway subnet is part of the virtual network IP address range that you specify when configuring your virtual network. It contains the IP addresses that the virtual network gateway resources and services use. The subnet must be named 'GatewaySubnet' in order for Azure to deploy the gateway resources. You can't specify a different subnet to deploy the gateway resources to. If you don't have a subnet named 'GatewaySubnet', when you create your VPN gateway, it will fail.

When you create the gateway subnet, you specify the number of IP addresses that the subnet contains. The number of IP addresses needed depends on the VPN gateway configuration that you want to create. Some configurations require more IP addresses than others. We recommend that you create a gateway subnet that uses a /27 or /28.

Options A and B are incorrect since these address spaces would conflict with the on-premise address space.

Option C is incorrect since this address space should ideally be used for the subnet to host the virtual machines.

For more information on site-to-site VPN, please visit the below URL

<https://docs.microsoft.com/en-us/azure/vpn-gateway/vpn-gateway-howto-site-to-site-resource-manager-portal>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 8

Unattempted

Domain :Design for deployment, migration, and integration

A company has an API service that currently returns XML data to its internal users. The API is going to be migrated onto Azure. It will sit behind an API Management instance. Below are the requirements for the API when it is moved to Azure

- The API must send data in JSON format to its internal users
- When external consultants access the API, the header information must be stripped before the data is received

What is the minimum number of API's that need to be added to Azure API management?

- A. 1 
- B. 2
- C. 3
- D. 4

#### Explanation:

Answer – A

Since you just have one API, you can place that behind the API management instance

Because of this reasoning all other options are incorrect

For more information on creating an API management instance, please visit the below URL

<https://docs.microsoft.com/en-us/azure/api-management/get-started-create-service-instance>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 9

Unattempted

Domain :Design for deployment, migration, and integration

A company has an API service that currently returns XML data to its internal users. The API is going to be migrated onto Azure. It will sit behind an API Management instance. Below are the requirements for the API when it is moved to Azure

- The API must send data in JSON format to its internal users
- When external consultants access the API, the header information must be stripped before the data is received

What is the minimum number of products to publish in Azure API management?

- A. 1
- B. 2
- C. 3
- D. 4

#### Explanation:

Answer – A

You can have one product that is published for the Internal development team. The Microsoft documentation mentions the following on products for API management

## Products

Products are how APIs are surfaced to developers. Products in API Management have one or more APIs, and are configured with a title, description, and terms of use. Products can be **Open** or **Protected**. Protected products must be subscribed to before they can be used, while open products can be used without a subscription. When a product is ready for use by developers, it can be published. Once it is published, it can be viewed (and in the case of protected products subscribed to) by developers. Subscription approval is configured at the product level and can either require administrator approval, or be auto-approved.

Groups are used to manage the visibility of products to developers. Products grant visibility to groups, and developers can view and subscribe to the products that are visible to the groups in which they belong.

Because of this reasoning all other options are incorrect

For more information on the key concepts for API Management, please visit the below URL

<https://docs.microsoft.com/en-us/azure/api-management/api-management-key-concepts>

Ask our Experts

Rate this Question?  

Domain :Design for deployment, migration, and integration

A company has an API service that currently returns XML data to its internal users. The API is going to be migrated onto Azure. It will sit behind an API Management instance. Below are the requirements for the API when it is moved to Azure

- The API must send data in JSON format to its internal users
- When external consultants access the API, the header information must be stripped before the data is received

What is the minimum number of policy elements that need to be added to the API?

- A. 1
- B. 2
- C. 3
- D. 4

#### Explanation:

Answer – B

You can have one policy element to ensure that XML data is transformed to JSON for the Internal users when it is published to Azure

Below is the policy statement given in the Microsoft documentation

## Convert XML to JSON

The `xml-to-json` policy converts a request or response body from XML to JSON. This policy can be used to modernize APIs based on XML-only backend web services.

### Policy statement

XML	
<pre>&lt;xml-to-json kind="javascript-friendly   direct" apply="always   content-type-xml" consider-accept-header="true  </pre>	

And then you can have one policy element to set the header of the response, so that it is sent as per the requirement to external consultants

Below is the policy statement given in the Microsoft documentation

# Set HTTP header

The `set-header` policy assigns a value to an existing response and/or request header or adds a new response and/or request header.

Inserts a list of HTTP headers into an HTTP message. When placed in an inbound pipeline, this policy sets the HTTP headers for the request being passed to the target service. When placed in an outbound pipeline, this policy sets the HTTP headers for the response being sent to the gateway's client.

## Policy statement

XML

 Copy

```
<set-header name="header name" exists-action="override | skip | append | delete">
    <value>value</value> <!--for multiple headers with the same name add additional value elements-->
</set-header>
```

Because of this reasoning all other options are incorrect

For more information on API Management policies, please visit the below URL

<https://docs.microsoft.com/en-us/azure/api-management/api-management-policies>

[Ask our Experts](#)

Rate this Question?  

[View Queries](#)

open 

Question 11

Unattempted

Domain :Design for identity and security

A company currently has resources deployed to their on-premise network and to Azure AD. There is a requirement to ensure that the Azure AD tenant can only be managed from workstations on the on-premise network. Which of the following needs to be part of the implementation of this requirement?

- A. Azure AD roles and administrators
- B. Conditional access policy 
- C. Role Based Access Control
- D. Azure AD Privileged Identity Management

## Explanation:

Answer - B

This can be managed by conditional access policies ensuring that the Locations is set in the policy.

The screenshot shows the 'whizlabpolicy' configuration page in the Azure portal. On the left, under 'Assignments', there are sections for 'Users and groups' and 'Cloud apps or actions', both showing 1 app included. Under 'Access controls', there are sections for 'Grant' and 'Session', both showing 0 controls selected. At the bottom, the 'Enable policy' switch is set to 'Off'. On the right, the 'Conditions' tab is selected, showing a list of conditions: 'Sign-in risk' (Not configured), 'Device platforms' (Not configured), 'Locations' (Not configured), 'Client apps (preview)', and 'Device state (preview)'. A red arrow points to the 'Locations' condition entry. At the bottom right of the conditions panel is a blue 'Done' button.

Since this is clear from the implementation, all other options are incorrect

For more information on conditional access, please visit the below URL

<https://docs.microsoft.com/en-us/azure/active-directory/conditional-access/overview>

## View Queries

open ▾

## Question 12

Unattempted

Domain :Design an infrastructure strategy

A team has an application that receives data from IoT based devices. The data is sent to CosmosDB which uses the SQL API. A notification needs to be sent when data is received from the IoT devices. Which of the following can be part of the implementation? Choose 2 answers from the options give below

- A. Deploy an Azure Logic App that has an Azure CosmosDB connector 
- B. Deploy a function App that has an Azure CosmosDB connector
- C. Ensure the Azure Logic App uses a SendGrid action 
- D. Ensure the Azure Function App uses a SendGrid action

**Explanation:**

Answer - A and C

You can use the CosmosDB connector for Azure Logic App to trigger a workflow when data is sent to CosmosDB. The Microsoft documentation mentions the following

# Azure Cosmos DB



Azure Cosmos DB is a globally distributed multi-model database service designed to enable you to elastically and independently scale throughput and storage across any number of geographical regions with a comprehensive SLA.

This connector is available in the following products and regions:

Service	Class	Regions
Logic Apps	Standard	All Logic Apps regions
Flow	Standard	All Flow regions
PowerApps	Standard	All PowerApps regions

To use this integration, you will need a Cosmos DB SQL API account configured in the Azure Portal. Note that MongoDB API accounts are not currently supported.

You can use SendGrid to send emails from the Log App. The Microsoft documentation mentions the following

## Send emails and manage mailing lists in SendGrid by using Azure Logic Apps

08/24/2018 • 2 minutes to read • Contributors A row of small, semi-transparent circular icons representing different users.

With Azure Logic Apps and the SendGrid connector, you can create automated tasks and workflows that send emails and manage your recipient lists, for example:

- Send email.
- Add recipients to lists.
- Get, add, and manage global suppression.

You can use SendGrid actions in your logic apps to perform these tasks. You can also have other actions use the output from SendGrid actions.

Options B and D are incorrect since you would need to use Azure Logic Apps to build a workflow

For more information on the connectors, please visit the below URL

<https://docs.microsoft.com/en-us/connectors/documentdb/>

Ask our Experts

Rate this Question?  

View Queries

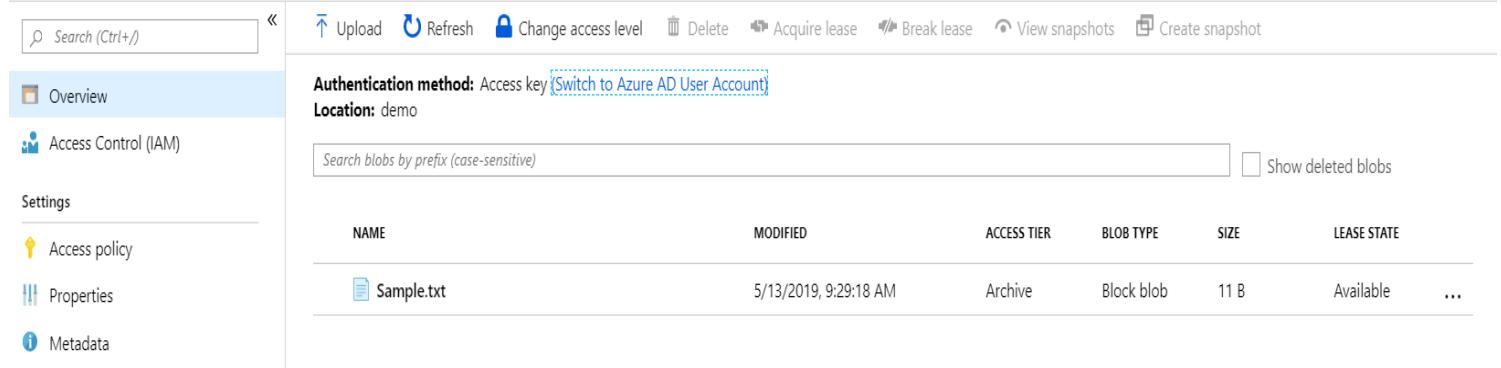
open ▾

Question 13

Unattempted

Domain :Design an infrastructure strategy

A team has created a storage account in Azure. They also have the following object available in the storage account



The screenshot shows the Azure Storage Explorer interface. On the left, there's a sidebar with options: Search (Ctrl+), Overview, Access Control (IAM), Settings, Access policy, Properties, and Metadata. The main area shows a blob named 'Sample.txt' in a container named 'demo'. The blob details are: NAME: Sample.txt, MODIFIED: 5/13/2019, 9:29:18 AM, ACCESS TIER: Archive, BLOB TYPE: Block blob, SIZE: 11 B, LEASE STATE: Available. A red box highlights the 'Archive' access tier.

NAME	MODIFIED	ACCESS TIER	BLOB TYPE	SIZE	LEASE STATE
Sample.txt	5/13/2019, 9:29:18 AM	Archive	Block blob	11 B	Available

In order to access the Sample.txt file, which of the following must be done first?

- A. Generate a snapshot
- B. Modify the access tier 
- C. Generate a shared access signature
- D. Modify the type of blob

#### Explanation:

Answer – B

In order to access the BLOB, since it is in the archive access tier, you need to first change the access tier for the blob object. The Microsoft documentation mentions the following

## Blob rehydration

To read data in archive storage, you must first change the tier of the blob to hot or cool. This process is known as rehydration and can take up to 15 hours to complete. Large blob sizes are recommended for optimal performance. Rehydrating several small blobs concurrently may add additional time.

During rehydration, you may check the **Archive Status** blob property to confirm if the tier has changed. The status reads "rehydrate-pending-to-hot" or "rehydrate-pending-to-cool" depending on the destination tier. Upon completion, the archive status property is removed, and the **Access Tier** blob property reflects the new hot or cool tier.

Since this is clearly mentioned in the documentation, all other options are incorrect

For more information on the storage tiers, please visit the below URL

<https://docs.microsoft.com/en-us/azure/storage/blobs/storage-blob-storage-tiers>

Ask our Experts

Rate this Question?

View Queries

open ▾

Question 14

Unattempted

Domain :Design an infrastructure strategy

A team has created a storage account in Azure. They also have the following object available in the storage account

The screenshot shows the Azure Storage Blob service interface. On the left, there's a sidebar with navigation links: Search (Ctrl+P), Overview, Access Control (IAM), Settings, Access policy, Properties, and Metadata. The 'Overview' link is currently selected. At the top, there are several action buttons: Upload, Refresh, Change access level, Delete, Acquire lease, Break lease, View snapshots, and Create snapshot. Below these buttons, there's a search bar labeled 'Search blobs by prefix (case-sensitive)' and a checkbox for 'Show deleted blobs'. The main content area displays a table of blobs. The table has columns: NAME, MODIFIED, ACCESS TIER, BLOB TYPE, SIZE, and LEASE STATE. There is one entry: 'Sample.txt' with a modified date of '5/13/2019, 9:29:18 AM', an access tier of 'Archive', a blob type of 'Block blob', a size of '11 B', and a lease state of 'Available'. A 'More' button is at the bottom right of the table.

NAME	MODIFIED	ACCESS TIER	BLOB TYPE	SIZE	LEASE STATE
Sample.txt	5/13/2019, 9:29:18 AM	Archive	Block blob	11 B	Available

Currently the Sample.txt file will be stored

- A. At the highest storage cost
- B. At the lowest storage cost
- C. At the lowest data retrieval cost
- D. At the premium storage cost

## Explanation:

Answer – B

The Microsoft documentation mentions that the object will be set at a lower storage cost.

The following considerations apply to the different access tiers:

- The archive access tier is available only at the blob level and not at the storage account level.
- Data in the cool access tier can tolerate slightly lower availability, but still requires high durability and similar time-to-access and throughput characteristics as hot data. For cool data, a slightly lower availability service-level agreement (SLA) and higher access costs compared to hot data are acceptable trade-offs for lower storage costs.
- Archive storage is offline and offers the lowest storage costs but also the highest access costs.
- Only the hot and cool access tiers can be set at the account level.
- Hot, cool, and archive tiers can be set at the object level.

Since this is clearly mentioned in the documentation, all other options are incorrect

For more information on the storage tiers, please visit the below URL

<https://docs.microsoft.com/en-us/azure/storage/blobs/storage-blob-storage-tiers>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 15

Unattempted

Domain :Design a data platform solution

A company has the following on-premise data stores

- A Microsoft SQL Server 2012 database
- A Microsoft SQL Server 2008 database

The data needs to be migrated to Azure.

- **Requirement 1** - The data in the Microsoft SQL Server 2012 database needs to be migrated to an Azure SQL database
- **Requirement 2** - The data in a table in the Microsoft SQL Server 2008 database needs to be migrated to an Azure CosmosDB account that uses the SQL API

Which of the following should be used to accomplish Requirement1?

- A. AzCopy

B. Azure CosmosDB Data Migration tool

C. Data Management Gateway

D. Data Migration Assistant 

#### Explanation:

Answer – D

The Data Migration assistant can be used to migrate the data. It has support for various versions of Microsoft SQL Server as shown below

## Supported source and target versions

DMA replaces all previous versions of SQL Server Upgrade Advisor and should be used for upgrades for most SQL Server versions. Supported source and target versions are:

### Sources

- SQL Server 2005
- SQL Server 2008
- SQL Server 2008 R2
- SQL Server 2012
- SQL Server 2014
- SQL Server 2016
- SQL Server 2017 on Windows

### Targets

- SQL Server 2012
- SQL Server 2014
- SQL Server 2016
- SQL Server 2017 on Windows and Linux
- Azure SQL Database
- Azure SQL Database Managed Instance

Option A is incorrect since this works with data in Azure storage accounts

Option B is incorrect since this is used for migration of data to CosmosDB

Option C is incorrect since this is used for building a gateway with the on-premise infrastructure

For more information on the data migration assistant, please visit the below URL

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 16

Unattempted

Domain :Design a data platform solution

A company has the following on-premise data stores

- A Microsoft SQL Server 2012 database
- A Microsoft SQL Server 2008 database

The data needs to be migrated to Azure.

- **Requirement 1** - The data in the Microsoft SQL Server 2012 database needs to be migrated to an Azure SQL database
- **Requirement 2** - The data in a table in the Microsoft SQL Server 2008 database needs to be migrated to an Azure CosmosDB account that uses the SQL API

Which of the following should be used to accomplish Requirement2?

- A. AzCopy
- B. Azure CosmosDB Data Migration tool 
- C. Data Management Gateway
- D. Data Migration Assistant

#### Explanation:

Answer – B

This tool can be used for migrating data onto CosmosDB. The Microsoft documentation mentions the following

# Use Data migration tool to migrate your data to Azure Cosmos DB

02/22/2019 • 24 minutes to read • Contributors  all

This tutorial provides instructions on using the Azure Cosmos DB Data Migration tool, which can import data from various sources into Azure Cosmos DB collections and tables. You can import from JSON files, CSV files, SQL, MongoDB, Azure Table storage, Amazon DynamoDB, and even Azure Cosmos DB SQL API collections. You migrate that data to collections and tables for use with Azure Cosmos DB. The Data Migration tool can also be used when migrating from a single partition collection to a multi-partition collection for the SQL API.

Option A is incorrect since this works with data in Azure storage accounts

Option C is incorrect since this is used for building a gateway with the on-premise infrastructure

Option D is incorrect since this tool does not have support for CosmosDB

For more information on importing data into CosmosDB, please visit the below URL

<https://docs.microsoft.com/en-us/azure/cosmos-db/import-data>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 17

Unattempted

Domain :Design an infrastructure strategy

Your team is planning on hosting an application that will be hosted on 2 virtual machines in Azure named demovm1 and demovm2. You have to load balance the traffic from the Internet to the Virtual Machines using one Azure Load balancer.

You need to recommend the minimum number of Public IP addresses that would be required

What is the minimum number of Public IP addresses that would be required for the Load Balancer?

- A. 0
- B. 1 
- C. 2
- D. 3

## Explanation:

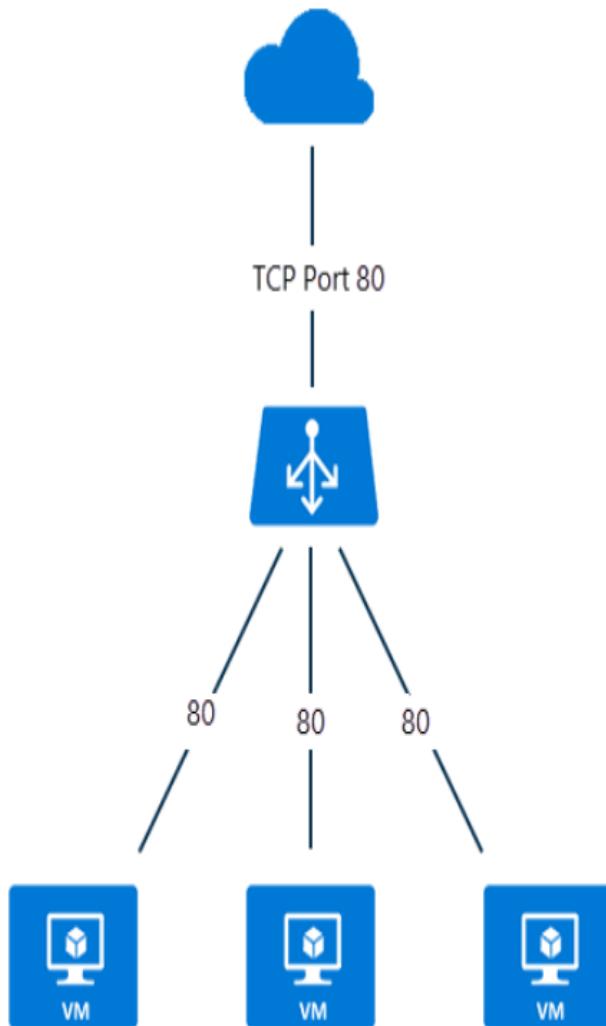
Answer – B

You would just need one Public IP address assigned to the Load Balancer. If you look at the example give in the Microsoft documentation on the public load balancer, you can see that the request comes from the clients to the public IP and port of the load balancer

## Public Load Balancer

A public Load Balancer maps the public IP address and port number of incoming traffic to the private IP address and port number of the VM, and vice versa for the response traffic from the VM. By applying load-balancing rules, you can distribute specific types of traffic across multiple VMs or services. For example, you can spread the load of web request traffic across multiple web servers.

The following figure shows a load-balanced endpoint for web traffic that is shared among three VMs for the public and TCP port 80. These three VMs are in a load-balanced set.



Since this is a clear concept of the Load balancer, all other options are incorrect

For more information on the concepts of the load balancer, please visit the below URL

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 18

Unattempted

Domain :Design an infrastructure strategy

Your team is planning on hosting an application that will be hosted on 2 virtual machines in Azure named demovm1 and demovm2. You have to load balance the traffic from the Internet to the Virtual Machines using one Azure Load balancer.

You need to recommend the minimum number of Public IP addresses that would be required  
What is the minimum number of Public IP addresses that would be required for demovm1?

- A. 0
- B. 1
- C. 2
- D. 3

#### Explanation:

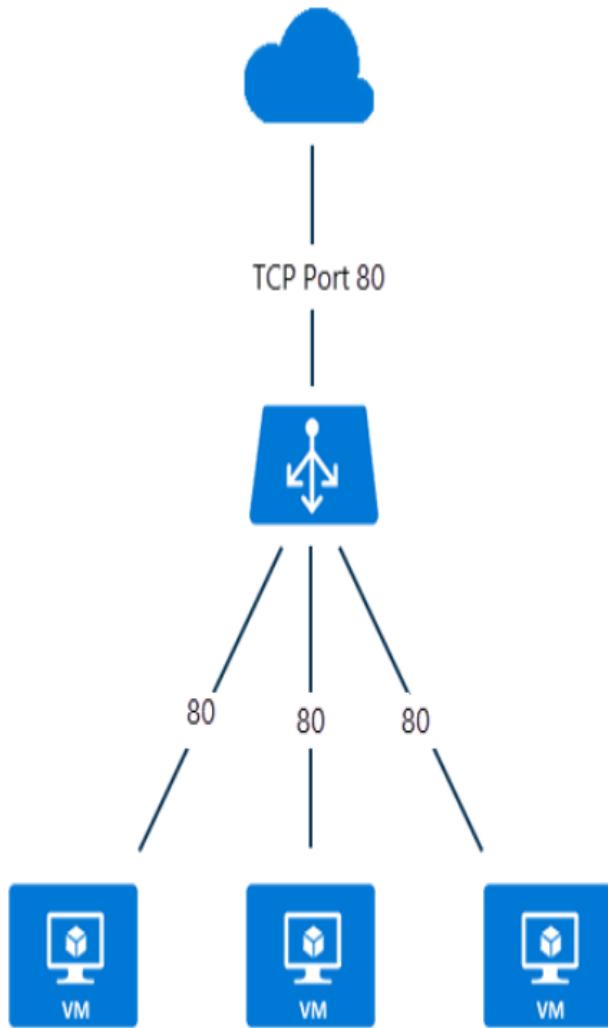
Answer – A

When the request is directed from the load balancer to the back end virtual machine, it is made to the Private IP address of the backend virtual machine. This is why you don't need to have a public IP address assigned to the Virtual Machine. This is also given in the Microsoft documentation

## Public Load Balancer

A public Load Balancer maps the public IP address and port number of incoming traffic to the private IP address and port number of the VM, and vice versa for the response traffic from the VM. By applying load-balancing rules, you can distribute specific types of traffic across multiple VMs or services. For example, you can spread the load of web request traffic across multiple web servers.

The following figure shows a load-balanced endpoint for web traffic that is shared among three VMs for the public and TCP port 80. These three VMs are in a load-balanced set.



Since this is a clear concept of the Load balancer, all other options are incorrect

For more information on the concepts of the load balancer, please visit the below URL

<https://docs.microsoft.com/en-us/azure/load-balancer/load-balancer-overview>

Ask our Experts

Rate this Question?

**Domain :Design an infrastructure strategy**

Your team is planning on hosting an application that will be hosted on 2 virtual machines in Azure named demovm1 and demovm2. You have to load balance the traffic from the Internet to the Virtual Machines using one Azure Load balancer.

You need to recommend the minimum number of Public IP addresses that would be required  
What is the minimum number of Public IP addresses that would be required for demovm2?

- A. 0
- B. 1
- C. 2
- D. 3

**Explanation:**

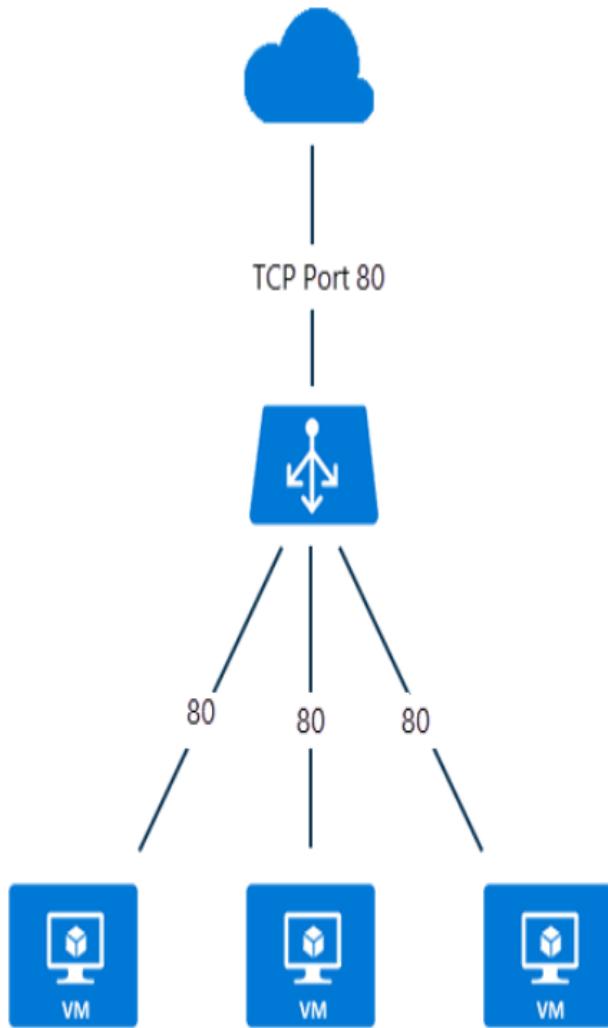
Answer – A

When the request is directed from the load balancer to the back end virtual machine, it is made to the Private IP address of the backend virtual machine. This is why you don't need to have a public IP address assigned to the Virtual Machine. This is also given in the Microsoft documentation

## Public Load Balancer

A public Load Balancer maps the public IP address and port number of incoming traffic to the private IP address and port number of the VM, and vice versa for the response traffic from the VM. By applying load-balancing rules, you can distribute specific types of traffic across multiple VMs or services. For example, you can spread the load of web request traffic across multiple web servers.

The following figure shows a load-balanced endpoint for web traffic that is shared among three VMs for the public and TCP port 80. These three VMs are in a load-balanced set.



Since this is a clear concept of the Load balancer, all other options are incorrect

For more information on the concepts of the load balancer, please visit the below URL

<https://docs.microsoft.com/en-us/azure/load-balancer/load-balancer-overview>

Ask our Experts

Rate this Question?

Domain :Design an infrastructure strategy

A company wants to run a series of tasks using the Azure Batch service. They have the following key requirements

- Large-scale parallel execution of Azure Batch jobs.
- Minimize costs to run the solution

Which of the following would you consider implementing for the Virtual Machines for running the Batch jobs?

- A. Basic A series Virtual Machines
- B. Use Virtual Machines that support Message Passing Interface (MPI)
- C. Low priority virtual machines
- D. Burstable virtual machines

#### Explanation:

Answer – C

The Azure Batch service already gives you the capability of running processes in parallel. If you want to ensure minimum costs for your solution, then you can consider using low priority nodes.

The number of nodes and VM size are set using defined constants. Batch supports dedicated nodes and [low-priority nodes](#), and you can use either or both in your pools. Dedicated nodes are reserved for your pool. Low-priority nodes are offered at a reduced price from surplus VM capacity in Azure. Low-priority nodes become unavailable if Azure does not have enough capacity. The sample by default creates a pool containing only 5 low-priority nodes in size *Standard\_A1\_v2*.

Option A is incorrect since this is not supported as per the Microsoft documentation as given below

## Supported VM families and sizes

### Pools in Virtual Machine configuration

Batch pools in the Virtual Machine configuration support all VM sizes ([Linux](#), [Windows](#)) except for the following:

Family	Unsupported sizes
Basic A-series	Basic_A0 (A0)

Option B is incorrect since this is used if you have tightly coupled workloads that need to run on the Azure Batch service.

Option D is incorrect since this is used if you have unpredictable workloads to run in Azure..

For more information on an example of using the Azure Batch service, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/batch/tutorial-parallel-dotnet>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 21

Unattempted

Domain :Design for identity and security

A company has deployed several applications to Azure based virtual machines. These virtual machines need to be managed by a set of administrators by using point-to-site VPN connections. The certificates for the connections will be generated by an on-premise certification authority.

You need to provide recommendations on what certificates are required for the deployment  
Which of the following needs to be setup on the Personal certificate store on each laptop?

- A. The root CA certificate that has the private key
- B. The root CA certificate that has the public key
- C. The user certificate that has the private key 
- D. The user certificate that has the public key

#### Explanation:

Correct Answer: C

#### Explanation:

At Server: need to install "Root Certificate having public Key" normally \*.crt file

At Client Computer: need to install "Client Certificate having Private key"

Point-to-Site certificate authentication connections require the following prerequisites:

- A Dynamic VPN gateway.
- The public key (.cer file) for a root certificate, which is uploaded to Azure. This key is considered a trusted certificate and is used for authentication
- A client certificate generated from the root certificate, and installed on each client computer that will connect. This certificate is used for client authentication
- A VPN client configuration package must be generated and installed on every client computer that connects. The client configuration package configures the native VPN client that's already on the operating

system with the necessary information to connect to the VNet

Point-to-Site connections don't require a VPN device or an on-premises public-facing IP address. The VPN connection is created over SSTP (Secure Socket Tunneling Protocol). On the server side, we support SSTP versions 1.0, 1.1, and 1.2. The client decides which version to use. For Windows 8.1 and above, SSTP uses 1.2 by default

Detail explanation at <https://docs.microsoft.com/en-us/azure/vpn-gateway/vpn-gateway-certificates-point-to-site>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 22

Unattempted

Domain :Design for identity and security

A company has deployed several applications to Azure based virtual machines. These virtual machines need to be managed by a set of administrators by using point-to-site VPN connections. The certificates for the connections will be generated by an on-premise certification authority.

You need to provide recommendations on what certificates are required for the deployment  
Which of the following needs to be setup on the Computer Personal store on each laptop?

- A. The root CA certificate that has the private key
- B. The root CA certificate that has the public key
- C. The user certificate that has the private key 
- D. The user certificate that has the public key

#### Explanation:

Answer – C

The user's personal store will have the user certificate that has the private key.

For more information on point to site connections, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/vpn-gateway/vpn-gateway-howto-point-to-site-resource-manager-portal>

## View Queries

open ▾

## Question 23

Unattempted

Domain :Design for identity and security

A company has deployed several applications to Azure based virtual machines. These virtual machines need to be managed by a set of administrators by using point-to-site VPN connections. The certificates for the connections will be generated by an on-premise certification authority.

You need to provide recommendations on what certificates are required for the deployment

Which of the following needs to be setup in the Azure VPN gateway?

- A. The root CA certificate that has the private key
- B. The root CA certificate that has the public key
- C. The user certificate that has the private key
- D. The user certificate that has the public key

**Explanation:**

Answer – B

This is mentioned in the Microsoft documentation as one of the steps in setting up a point to site connection.

## 9. Upload the root certificate public certificate data

You can upload additional trusted root certificates up to a total of 20. Once the public certificate data is uploaded, Azure can use it to authenticate clients that have installed a client certificate generated from the trusted root certificate. Upload the public key information for the root certificate to Azure.

1. Certificates are added on the **Point-to-site configuration** page in the **Root certificate** section.
2. Make sure that you exported the root certificate as a Base-64 encoded X.509 (.cer) file. You need to export the certificate in this format so you can open the certificate with text editor.
3. Open the certificate with a text editor, such as Notepad. When copying the certificate data, make sure that you copy the text as one continuous line without carriage returns or line feeds. You may need to modify your view in the text editor to 'Show Symbol/Show all characters' to see the carriage returns and line feeds. Copy only the following section as one continuous line:



```
P2SRootCert.cer - Notepad
File Edit Format View Help
-----BEGIN CERTIFICATE-----
MIIC6zCAQD0AwIBAgIQUVU0/09T3qJGMbd6rc9zCTANBgkqhkiG9w0BAQsFADAY
MRYwFAYDVQQDADQ1Qm1NSb290Q2VydDEwIB4XDTE3MDgwNzIxNTgBN1oXDTE4MDgw
NzIyMTg0N1owGOEWIBQGA1UEAwwNU0JTUm9vdEN1cnQxDCCASIwOQYJKoZIhvcN
AQEBBQAQDggEPADCCAQoCggEBANW4PjxpJkPnYHbToxn4+YE178cP8HzIsZqvzqv
vJvgovweHQ2mQnxweU127arHaZ9fjaJ9ACOUgT/XKC2gnq3Dej4ZCdPzG7Hgpf
eVZZuAUDeeUh1D9nqnxpsVCrRiuHyT9Yh9zwRYDHQa12/ttTjh3fp7cXPj1
KSpvdms5esZpwvPphVN83KAHuWGMk4eVCX2kS9fRGte3iR9RjGo/Ueqj/I/pVmUN
bIEte4AJEKmmjd8Lg6rdqdd+h1eky9u3fxZTPCwoqTE4TzL6972mzU1p1lyV8q5L
XbmLQPUXaMKGjIvZ67k14xqc5+0z8pRq0jIwmZk03N10ECAwEAAaMx/MC8wOgYD
VR0PAQH/BAAQAgEMB0GA1udggQwBRxEryqxzhULzGcfna3QbPoKSSTANBpkq
hkig9w0BAQsFAACQAfqlxeuzss+EU24pBrPVqB99QyFyfJHAZ3n3kawIxhiHTQ
+hu6tDoemScv9u+aTRRj8j2CRhDec6SeuD3aptw-PVtUNew7MOpivHpyXiiwpHL
FpyoUCqhK7X31zwaZIAFp90/+Cns0Wz18b1RgagY7x4pYighwhCvJHttB0fczz
CX2jppjeHbecJ8Kfhd1NWxByJEFXkf/vA1hu1qOKgPGV03l2icNVGLywG7xb6b
lkQoKTCRTvHYA9wd9vCERS5mhHBc5jboQJ0T1m7jgSeciLC11kyMC7LR2Qkc0NyB
+SPkthQa3ky0Ebd3DG7Rdzdr3Ic0Zuj6E1D1EJhpg==
-----END CERTIFICATE-----
```

Since this is clearly mentioned, all other options are incorrect

For more information on point to site connections, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/vpn-gateway/vpn-gateway-howto-point-to-site-resource-manager-portal>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 24

Unattempted

Domain :Design an infrastructure strategy

A company has deployed several applications across Windows and Linux Virtual machines in Azure. Log Analytics are being used to send the required data for alerting purposes for the Virtual Machines.

You need to recommend which tables need to be queried for security related queries.

Which of the following would you query for events from Windows Event Logs?

- A. Azure Activity

B. Azure Diagnostics

C. Event 

D. Syslog

#### Explanation:

Answer – C

This is also given in the Microsoft documentation, wherein you would use the Event Table for the queries on events from Windows Virtual machines

## Log queries with Windows Events

The following table provides different examples of log queries that retrieve Windows Event records.

Query	Description
Event	All Windows events.
Event   where EventLevelName == "error"	All Windows events with severity of error.
Event   summarize count() by Source	Count of Windows events by source.
Event   where EventLevelName == "error"   summarize count() by Source	Count of Windows error events by source.

Since this is clearly mentioned, all other options are incorrect

For more information on collecting event data from windows virtual machines, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/azure-monitor/platform/data-sources-windows-events>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 25

Unattempted

Domain :Design an infrastructure strategy

A company has deployed several applications across Windows and Linux Virtual machines in Azure. Log Analytics are being used to send the required data for alerting purposes for the Virtual Machines.

You need to recommend which tables need to be queried for security related queries.

Which of the following would you query for events from Linux system logging?

- A. Azure Activity
- B. Azure Diagnostics
- C. Event
- D. Syslog 

#### Explanation:

Answer – D

This is also given in the Microsoft documentation, wherein you would use the Syslog Table for the queries on events from Linux Virtual machines

## Log queries with Syslog records

The following table provides different examples of log queries that retrieve Syslog records.

Query	Description
Syslog	All Syslogs.
Syslog   where SeverityLevel == "error"	All Syslog records with severity of error.
Syslog   summarize AggregatedValue = count() by Computer	Count of Syslog records by computer.
Syslog   summarize AggregatedValue = count() by Facility	Count of Syslog records by facility.

Since this is clearly mentioned, all other options are incorrect

For more information on collecting event data from linux virtual machines, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/azure-monitor/platform/data-sources-syslog>

Ask our Experts

Rate this Question?  

[View Queries](#)

open 

Your company needs to generate a monthly report to determine the what, who, and when for any write operations on all resources that were deployed to the Azure subscription. Which of the following would help achieve this requirement?

- A. Azure Log Analytics
- B. Azure Activity Log
- C. Azure Monitor action groups
- D. Azure Advisor

### Explanation:

Answer – B

You can actually go to the

- 1) Activity Logs section in Azure Monitor
- 2) Choose the timespan for the time period required
- 3) Download the report as CSV

OPERATION NAME	STATUS	TIME	TIME STAMP	SUBSCRIPTION	EVENT INITIATED BY
Update Storage Account Create	Succeeded	55 min ago	Mon May 1...	Pay-As-You-Go	demousr2016@gmail.com
List Storage Account Keys	Succeeded	1 h ago	Mon May 1...	Pay-As-You-Go	demousr2016@gmail.com
Put blob container	Succeeded	1 h ago	Mon May 1...	Pay-As-You-Go	demousr2016@gmail.com
Validate Deployment	Succeeded	1 h ago	Mon May 1...	Pay-As-You-Go	demousr2016@gmail.com
Validate Deployment	Succeeded	1 h ago	Mon May 1...	Pay-As-You-Go	demousr2016@gmail.com

Since this is clear from the implementation, all other options are incorrect

For more information on Activity logs, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/azure-monitor/platform/activity-logs-overview>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 27

Unattempted

Domain :Determine workload requirements

A company is planning on deploying an application onto Azure. The application will be based on the .Net core programming language. The application would be hosted using Azure Web apps. Below is part of the various requirements for the application

- Give the ability to correlate Azure resource usage and the performance data with the actual application configuration and performance data
- Give the ability to visualize the relationships between application components
- Give the ability to track requests and exceptions to specific lines of code from within the application
- Give the ability to actually analyse how users return to an application and see how often they only select a particular drop-down value

Which of the following service would be best suited for fulfilling the requirement of

**"Give the ability to correlate Azure resource usage and the performance data with the actual application configuration and performance data"**

- A. Azure Application Insights
- B. Azure Service Map
- C. Azure Log Analytics
- D. Azure Activity Log

#### Explanation:

Answer – C

You can send data about the application and resource usage to Azure Log Analytics. You can then build queries on the stored data.

For more information on Azure Log Analytics, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/azure-monitor/learn/tutorial-viewdata>

Ask our Experts

Rate this Question?  

## Question 28

Unattempted

Domain :Determine workload requirements

A company is planning on deploying an application onto Azure. The application will be based on the .Net core programming language. The application would be hosted using Azure Web apps. Below is part of the various requirements for the application

- Give the ability to correlate Azure resource usage and the performance data with the actual application configuration and performance data
- Give the ability to visualize the relationships between application components
- Give the ability to track requests and exceptions to specific lines of code from within the application
- Give the ability to actually analyse how users return to an application and see how often they only select a particular drop-down value

Which of the following service would be best suited for fulfilling the requirement of  
"Give the ability to visualize the relationships between application components"

- A. Azure Application Insights 
- B. Azure Service Map
- C. Azure Log Analytics
- D. Azure Activity Log

**Explanation:**

Answer – A

question is talking about "application would be hosted using Azure Web apps"

Means PaaS service.

A. Azure Application Insights

This can be applied to PaaS and hence correct answer

<https://docs.microsoft.com/en-us/azure/azure-monitor/app/app-insights-overview>

B. Azure Service Map

This can be applied to IaaS. Our requirement is for Web App which is not IaaS

Microsoft Azure introduced a new service called Service Map, which is a tool that enables us to map all communication flow to and from a monitored server.

<https://docs.microsoft.com/en-us/azure/azure-monitor/insights/service-map>

C. Azure Log Analytics

This is good for Analysis of Logs of the component bit not for **ability to visualize**.

<https://docs.microsoft.com/en-us/azure/azure-monitor/log-query/get-started-portal>

#### D. Azure Activity Log

Activity Log is storing the data what activity done in Azure Portal or by CLI or PowerShell on a particular Service. Hence this is wrong

Ask our Experts

Rate this Question?  

#### View Queries

open ▾

#### Question 29

Unattempted

Domain :Determine workload requirements

A company is planning on deploying an application onto Azure. The application will be based on the .Net core programming language. The application would be hosted using Azure Web apps. Below is part of the various requirements for the application

- Give the ability to correlate Azure resource usage and the performance data with the actual application configuration and performance data
- Give the ability to visualize the relationships between application components
- Give the ability to track requests and exceptions to specific lines of code from within the application
- Give the ability to actually analyse how users return to an application and see how often they only select a particular drop-down value

Which of the following service would be best suited for fulfilling the requirement of

**"Give the ability to track requests and exceptions to specific lines of code from within the application"**

- A. Azure Application Insights 
- B. Azure Service Map
- C. Azure Log Analytics
- D. Azure Activity Log

#### Explanation:

Answer – A

This feature is part of the Application Insights tool. An example of this is given in the Microsoft documentation

You can use the CodeLens feature in Application Insights to get a deep dive look into exceptions at the code level. An example from the Microsoft documentation is given below

# Exceptions in CodeLens

The screenshot shows a tooltip for an exception named 'ConnectionTimeoutException' at the 'ReportsController.AllForReport' method. The tooltip includes the following information:

- Count: 47
- Type: ConnectionTimeoutException at ReportsController.AllForReport
- Description: Timeout expired. The timeout period elapsed prior to obtaining a connection from the pool. This may have occ...
- Change: (-83%)
- Actions: Search | Trend | Go to code
- Link: View all exceptions in this app | Explore exception trends

The tooltip also displays a snippet of the source code:

```
47
48     public ActionResult Tickets()
49     {
50         var report = this.AllForReport(
51             serviceticket => serviceticket.Customer,
```

The exception CodeLens indicator shows the number of exceptions that have occurred in the past 24 hours from the 15 most frequently occurring exceptions in your application during that period, while processing the request served by the method.

To see more details, click the exceptions CodeLens indicator:

- The percentage change in number of exceptions from the most recent 24 hours relative to the prior 24 hours
- Choose **Go to code** to navigate to the source code for the function throwing the exception
- Choose **Search** to query all instances of this exception that have occurred in the past 24 hours
- Choose **Trend** to view a trend visualization for occurrences of this exception in the past 24 hours
- Choose **View all exceptions in this app** to query all exceptions that have occurred in the past 24 hours
- Choose **Explore exception trends** to view a trend visualization for all exceptions that have occurred in the past 24 hours.

Since this is clearly mentioned in the documentation, all other options are incorrect

For more information on the CodeLens feature, please visit the below URL

<https://docs.microsoft.com/en-us/azure/azure-monitor/app/visual-studio-codelens>

Ask our Experts

Rate this Question?

View Queries

open ▾

Question 30

Unattempted

Domain :Determine workload requirements

A company is planning on deploying an application onto Azure. The application will be based on the .Net core programming language. The application would be hosted using Azure Web apps. Below is part of the various requirements for the application

- Give the ability to correlate Azure resource usage and the performance data with the actual application configuration and performance data
- Give the ability to visualize the relationships between application components
- Give the ability to track requests and exceptions to specific lines of code from within the application
- Give the ability to actually analyse how users return to an application and see how often they only select a particular drop-down value

Which of the following service would be best suited for fulfilling the requirement of

**"Give the ability to actually analyse how users return to an application and see how often they only select a particular drop-down value"**

- A. Azure Application Insights 
- B. Azure Service Map
- C. Azure Log Analytics
- D. Azure Activity Log

#### Explanation:

Answer – A

This feature is part of the Application Insights tool. An example of this is given in the Microsoft documentation

## User retention analysis for web applications with Application Insights

05/03/2017 • 2 minutes to read • Contributors 

The retention feature in [Azure Application Insights](#) helps you analyze how many users return to your app, and how often they perform particular tasks or achieve goals. For example, if you run a game site, you could compare the numbers of users who return to the site after losing a game with the number who return after winning. This knowledge can help you improve both your user experience and your business strategy.

Since this is clearly mentioned in the documentation, all other options are incorrect

For more information on the retention feature of Application Insights, please visit the below URL

<https://docs.microsoft.com/en-us/azure/azure-monitor/app/usage-retention>

Ask our Experts

Rate this Question?  

## Question 31

Unattempted

Domain :Design for identity and security

A company is setting up the required privileges for users in their Azure AD tenant. They need to assign a group of users with the below mentioned privileges

- The users should be able to manage virtual networks
- They should not be allowed to manage role assignments

You need to ensure the right role is assigned based on the least privilege access.

You decide to provide the Owner role to the group

Would this fulfil the requirement?

A. Yes

B. No

**Explanation:**

Answer – B

This would also allow the users to have the ability to manage all resources and this would provide too many privileges. Below is a snippet from the Microsoft documentation on the permissions for the role

## Owner

---

Description

Lets you manage everything, including access to resources.

---

Id

8e3af657-a8ff-443c-a75c-2fe8c4bcb635

---

Actions



Create and manage resources of all types

For more information on in-built roles, please visit the below URL

<https://docs.microsoft.com/en-us/azure/role-based-access-control/built-in-roles>

Ask our Experts

Rate this Question?  

## Question 32

Unattempted

Domain :Design for identity and security

A company is setting up the required privileges for users in their Azure AD tenant. They need to assign a group of users with the below mentioned privileges

- The users should only be able to manage virtual networks
- They should not be allowed to manage role assignments

You need to ensure the right role is assigned based on the least privilege access.

You decide to provide the Contributor role to the group

Would this fulfil the requirement?

A. Yes

B. No

**Explanation:**

Answer - B

This would also allow the users to have the ability to manage all resources and this would provide too many privileges. Below is a snippet from the Microsoft documentation on the permissions for the role

## Contributor

Description

Lets you manage everything except access to resources.

Id

b24988ac-6180-42a0-ab88-20f7382dd24c

Actions

Create and manage resources of all types



For more information on in-built roles, please visit the below URL

<https://docs.microsoft.com/en-us/azure/role-based-access-control/built-in-roles>

Ask our Experts

Rate this Question?  

Domain :Design for identity and security

A company is setting up the required privileges for users in their Azure AD tenant. They need to assign a group of users with the below mentioned privileges

- The users should be able to manage virtual networks
- They should not be allowed to manage role assignments

You need to ensure the right role is assigned based on the least privilege access.

You decide to provide the Network Contributor role to the group

Would this fulfil the requirement?

- A. Yes
- B. No

#### Explanation:

Answer – A

Yes, this is the ideal role which allows users to manage virtual networks but does not give them the access to manage role assignments. Below is a snippet from the Microsoft documentation on the permissions for the role

# Network Contributor

Description	Lets you manage networks, but not access to them.
Id	4d97b98b-1d4f-4787-a291-c67834d212e7
Actions	
Microsoft.Authorization/*/read	Read roles and role Assignments
Microsoft.Insights/alertRules/*	Create and manage alert rules
Microsoft.Network/*	Create and manage networks
Microsoft.ResourceHealth/availabilityStatuses/read	Gets the availability statuses for all resources in the specified scope
Microsoft.Resources/deployments/*	Create and manage resource group deployments
Microsoft.Resources/subscriptions/resourceGroups/read	Gets or lists resource groups.
Microsoft.Support/*	Create and manage support tickets

For more information on in-built roles, please visit the below URL

<https://docs.microsoft.com/en-us/azure/role-based-access-control/built-in-roles>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 34

Unattempted

Domain :Determine workload requirements

A company is planning to migrate an on-premise application to Azure. One component of the application is a windows native application that would be used to perform image processing. It needs to be ensured that when the component is not running, no Azure compute resources are consumed. You need to also ensure that the image processing application runs every hour.

You decide to use Azure Web Jobs for the implementation

Would this fulfil the requirement?

A. Yes

B. No 

### Explanation:

Answer - B

Azure Web jobs are used along with Azure Web Apps for performing background tasks. Here there is no mention of the application being a web application, so using Azure Web job would not be practical.

For more information on Azure Web Jobs, please visit the below URL

<https://docs.microsoft.com/en-us/azure/app-service/webjobs-create>

[Ask our Experts](#)

Rate this Question?  

[View Queries](#)

open ▾

Question 35

Unattempted

Domain :Determine workload requirements

A company is planning to migrate an on-premise application to Azure. One component of the application is a windows native application that would be used to perform image processing. It needs to be ensured that when the component is not running, no Azure compute resources are consumed. You need to also ensure that the image processing application runs every hour.

You decide to use Azure Functions for the implementation

Would this fulfil the requirement?

A. Yes 

B. No

### Explanation:

Answer – A

This would fulfil the requirement. So Azure Functions can run as part of the Consumption plan wherein they only consume resources when they are run.

For more information on Azure Functions, please visit the below URL

<https://docs.microsoft.com/en-us/azure/azure-functions/functions-overview>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 36

Unattempted

Domain :Determine workload requirements

A company is planning to migrate an on-premise application to Azure. One component of the application is a windows native application that would be used to perform image processing. It needs to be ensured that when the component is not running, no Azure compute resources are consumed. You need to also ensure that the image processing application runs every hour.

You decide to use Azure Logic App for the implementation

Would this fulfil the requirement?

A. Yes

B. No 

#### Explanation:

Answer – B

Azure Logic Apps does perform the task, provided if recurrence triggers is used. Since, it's not mentioned in the question the answer is No..

For more information on Azure Logic Apps, please visit the below URL

<https://docs.microsoft.com/en-us/azure/logic-apps/logic-apps-overview>

**Note** - We can run recurring automated tasks, processes and workflows with Azure Logic Apps on a schedule by creating a logic app workflow that starts with a built-in Recurrence trigger or Sliding Window trigger. Refer the link to know further.

<https://docs.microsoft.com/en-us/azure/logic-apps/concepts-schedule-automated-recurring-tasks-workflows>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 37

Unattempted

A company currently has a Web application deployed in classic ASP. The application makes use of third-party DLL's. The deployment process for the application is too prone to errors and scaling and high availability are always issues.

You need to devise an approach to modernize the application that could be deployed to Azure and also ensure deployment and administrative overheads are reduced.

Which of the following would you use as the Modernization approach?

- A. Use Azure Logic Apps
- B. Use Azure Functions
- C. Use Azure Virtual Machines
- D. Use Azure Container Services

#### Explanation:

Answer – D

The ideal approach for modernization of applications is to use Containers. And then to run in the cloud, you can use Azure Container Services. The Microsoft documentation mentions the following on Azure Container Instances

## What is Azure Container Instances?

04/25/2019 • 2 minutes to read • Contributors  all

Containers are becoming the preferred way to package, deploy, and manage cloud applications. Azure Container Instances offers the fastest and simplest way to run a container in Azure, without having to manage any virtual machines and without having to adopt a higher-level service.

Azure Container Instances is a great solution for any scenario that can operate in isolated containers, including simple applications, task automation, and build jobs. For scenarios where you need full container orchestration, including service discovery across multiple containers, automatic scaling, and coordinated application upgrades, we recommend [Azure Kubernetes Service \(AKS\)](#).

Option A is incorrect since this is more of a workflow-based service

Option B is incorrect since is a compute service that ideally should not run web-based applications

Option C is incorrect since this just ports the application but does not present a modern approach to application design

For more information on Azure Container instances, please visit the below URL

<https://docs.microsoft.com/en-us/azure/container-instances/container-instances-overview>

[Ask our Experts](#)Rate this Question?  [View Queries](#)

open ▾

**Question 38****Unattempted****Domain :Design for deployment, migration, and integration**

A company currently has a Web application deployed in classic ASP. The application makes use of third-party DLL's. The deployment process for the application is prone to errors and scaling and high availability are always issues.

You need to devise an approach to modernize the application that could be deployed to Azure and also ensure deployment and administrative overheads are reduced.

Which of the following would you recommend as the deployment approach?

- A. Configure the application as a service
- B. Configure the application into a web role
- C. Package the application into a container
- D. Package the application into one DLL

**Explanation:**

Answer – C

You should package your application into containers which can then be deployed onto Azure

For more information on an example of building a container image, please visit the below URL

<https://docs.microsoft.com/en-us/azure/container-instances/container-instances-tutorial-prepare-app>

[Ask our Experts](#)Rate this Question?  [View Queries](#)

open ▾

**Question 39****Unattempted****Domain :Design an infrastructure strategy**

A team needs to deploy an application onto Azure. The application will be used to perform intensive and long compute calculations. You need to deploy the solution to Azure. Below are the key requirements

- Ensure every compute node can communicate with each other
- Maximize the number of nodes to perform the computations faster
- Least amount of effort to implement the entire solution

Which of the following would you use as the service for hosting the solution?

- A. Azure Load Balancer
- B. Azure Application Gateway
- C. Azure Virtual Machine scale sets
- D. Azure Batch

#### Explanation:

Answer – D

The ideal solution to use for this is the Azure Batch Service. The Microsoft documentation mentions the following

## What is Azure Batch?

01/19/2018 • 5 minutes to read • Contributors  all

Use Azure Batch to run large-scale parallel and high-performance computing (HPC) batch jobs efficiently in Azure. Azure Batch creates and manages a pool of compute nodes (virtual machines), installs the applications you want to run, and schedules jobs to run on the nodes. There is no cluster or job scheduler software to install, manage, or scale. Instead, you use [Batch APIs and tools](#), command-line scripts, or the Azure portal to configure, manage, and monitor your jobs.

Developers can use Batch as a platform service to build SaaS applications or client apps where large-scale execution is required. For example, build a service with Batch to run a Monte Carlo risk simulation for a financial services company, or a service to process many images.

There is no additional charge for using Batch. You only pay for the underlying resources consumed, such as the virtual machines, storage, and networking.

Options A and B are incorrect since these are load balancing solutions

Option C is incorrect since this should be used for applications that need to scale on demand

For more information on the Azure Batch service, please visit the below URL

<https://docs.microsoft.com/en-us/azure/batch/batch-technical-overview>

[Ask our Experts](#)Rate this Question?  [View Queries](#)

open ▾

**Question 40****Unattempted****Domain :Design an infrastructure strategy**

A team needs to deploy an application onto Azure. The application will be used to perform intensive and long compute calculations. You need to deploy the solution to Azure. Below are the key requirements

- Ensure every compute node can communicate with each other
- Maximize the number of nodes to perform the computations faster
- Least amount of effort to implement the entire solution

Which of the following would you implement within the service?

- A. Ensure Session Affinity
- B. Enable parallel task execution 
- C. Enable containers
- D. Enable URL routing

**Explanation:**

Answer – B

Since we are going to be using the Azure Batch Service to run the workloads, we need to ensure the tasks run in parallel. The Microsoft documentation mentions the following

# Run parallel workloads

Batch works well with intrinsically parallel (also known as "embarrassingly parallel") workloads. Intrinsically parallel workloads are those where the applications can run independently, and each instance completes part of the work. When the applications are executing, they might access some common data, but they do not communicate with other instances of the application. Intrinsically parallel workloads can therefore run at a large scale, determined by the amount of compute resources available to run applications simultaneously.

Some examples of intrinsically parallel workloads you can bring to Batch:

- Financial risk modeling using Monte Carlo simulations
- VFX and 3D image rendering
- Image analysis and processing
- Media transcoding
- Genetic sequence analysis
- Optical character recognition (OCR)
- Data ingestion, processing, and ETL operations
- Software test execution

Since this is the only ideal approach, all other options are incorrect

For more information on the Azure Batch service, please visit the below URL

<https://docs.microsoft.com/en-us/azure/batch/batch-technical-overview>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 41

Unattempted

Domain :Design an infrastructure strategy

**View Case Study**

The IT Security department has the requirement to ensure they get notified if there are any changes made to the configuration of the underlying Virtual Machines. A workflow needs to be designed for this requirement. Which of the following would you use for this purpose?

- A. Azure Notification Hub
- B. Azure Event Hub

C. Azure Logic App

D. Azure services Bus

#### Explanation:

Answer – C

If you are looking at creating workflows, then you have to use Azure Logic Apps

The Microsoft documentation additionally mentions the following

Azure Logic Apps is a cloud service that helps you automate and orchestrate tasks, business processes, and workflows when you need to integrate apps, data, systems, and services across enterprises or organizations. Logic Apps simplifies how you design and build scalable solutions for app integration, data integration, system integration, enterprise application integration (EAI), and business-to-business (B2B) communication, whether in the cloud, on premises, or both.

Option A is incorrect since this is used for notifications

Option B is incorrect since this is used for listening to events emitted by Azure resources

Option D is incorrect since this is used to work with Azure queues and topics

For more information on Azure Logic Apps, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/logic-apps/logic-apps-overview>

Ask our Experts

Rate this Question?

View Queries

open ▾

Question 42

Unattempted

Domain :Design an infrastructure strategy

**View Case Study**

Which of the following service would you use to fulfil the below case study requirement

**"The Virtual Machines need to have a central location for storage of files. They would connect to these file shares using SMB"**

A. Azure BLOB storage

B. Azure Site Recovery

C. Azure File Service

D. Azure Table Service

**Explanation:**

Answer – C

You would use the Azure File service which can be used to work with shares via the SMB protocol

The screenshot shows the Azure Files landing page. It features a header with the Azure logo and the word "Azure Files". Below the header, there's a section titled "Simple, secure, and fully managed cloud file shares". A bulleted list follows, with the fourth item highlighted by a large red arrow pointing to it. The list items are:

- ✓ Extend your servers to Azure with Sync for on-premises performance and capability
- ✓ Secure data at rest and in-transit using **SMB 3.0 and HTTPS**
- ✓ Simplify cloud file share management using familiar tools
- ✓ Create high-performance file shares using the Premium Files storage tier

Option A is incorrect since this is used for Object level storage

Option B is incorrect since this is used for backup and recovery purposes

Option D is incorrect since this is used for hosting tables in Azure

For more information on Azure Files, please go ahead and visit the below URL

<https://azure.microsoft.com/en-us/services/storage/files/>

Ask our Experts

Rate this Question?

View Queries

open ▾

Question 43

Unattempted

Domain :Design an infrastructure strategy

**View Case Study**

You need to comply with the below case study requirement

**"There should be an encrypted connection between the On-premise data centre and the Virtual Network whizlabs-net2"**

You decide to provision a point to site VPN connection

Would this fulfil the requirement?

A. Yes

B. No 

### Explanation:

Answer – B

Point to site connections are used when you want to connect to an Azure Virtual Network from workstation PC's

For more information on point to site VPN connections, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/vpn-gateway/vpn-gateway-howto-point-to-site-resource-manager-portal>

[Ask our Experts](#)

Rate this Question?  

[View Queries](#)

open ▾

Question 44

Unattempted

Domain :Design an infrastructure strategy

[View Case Study](#)

You need to comply with the below case study requirement

**"There should be an encrypted connection between the On-premise data centre and the Virtual Network whizlabs-net2"**

You decide to implement Virtual Network Peering

Would this fulfil the requirement?

A. Yes

B. No 

### Explanation:

Answer – B

Virtual Network peering is used to connect multiple virtual networks together

For more information on Virtual Network Peering, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/virtual-network/virtual-network-peering-overview>

## View Queries

open ▾

## Question 45

Unattempted

Domain :Design an infrastructure strategy

## View Case Study

You need to comply with the below case study requirement

**"There should be an encrypted connection between the On-premise data centre and the Virtual Network whizlabs-net2"**

You decide to implement a site-to-site VPN Connection.

Would this fulfil the requirement?

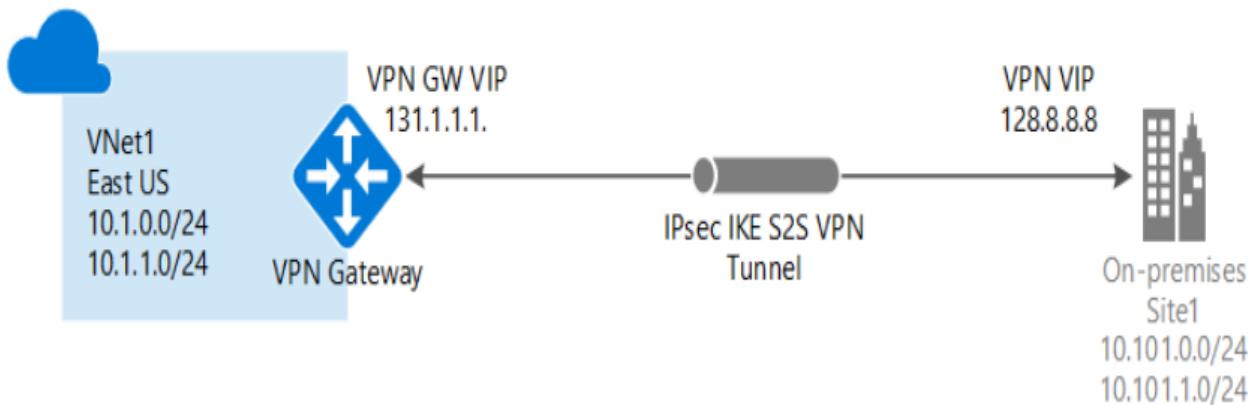
- A. Yes 
- B. No

**Explanation:**

Answer – A

This is the ideal and correct way to connect an Azure Virtual Network with an on-premise data center. The Microsoft documentation mentions the following

A Site-to-Site VPN gateway connection is used to connect your on-premises network to an Azure virtual network over an IPsec/IKE (IKEv1 or IKEv2) VPN tunnel. This type of connection requires a VPN device located on-premises that has an externally facing public IP address assigned to it. For more information about VPN gateways, see [About VPN gateway](#).



For more information on site to site VPN connections, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/vpn-gateway/vpn-gateway-howto-site-to-site-resource-manager-portal>

[Ask our Experts](#)Rate this Question?  [View Queries](#)

open ▾

**Question 46****Unattempted**

Domain :Design for identity and security

[View Case Study](#)

You need to decide on the implementation strategy for the synchronization of user passwords between the on-premise AD and the Azure AD tenant. You need to ensure that the sign-in is completely managed in the cloud. The company currently does not have the need to apply any policies from their on-premise AD for the authentication process. Which of the following would you consider for the configuration of the synchronization?

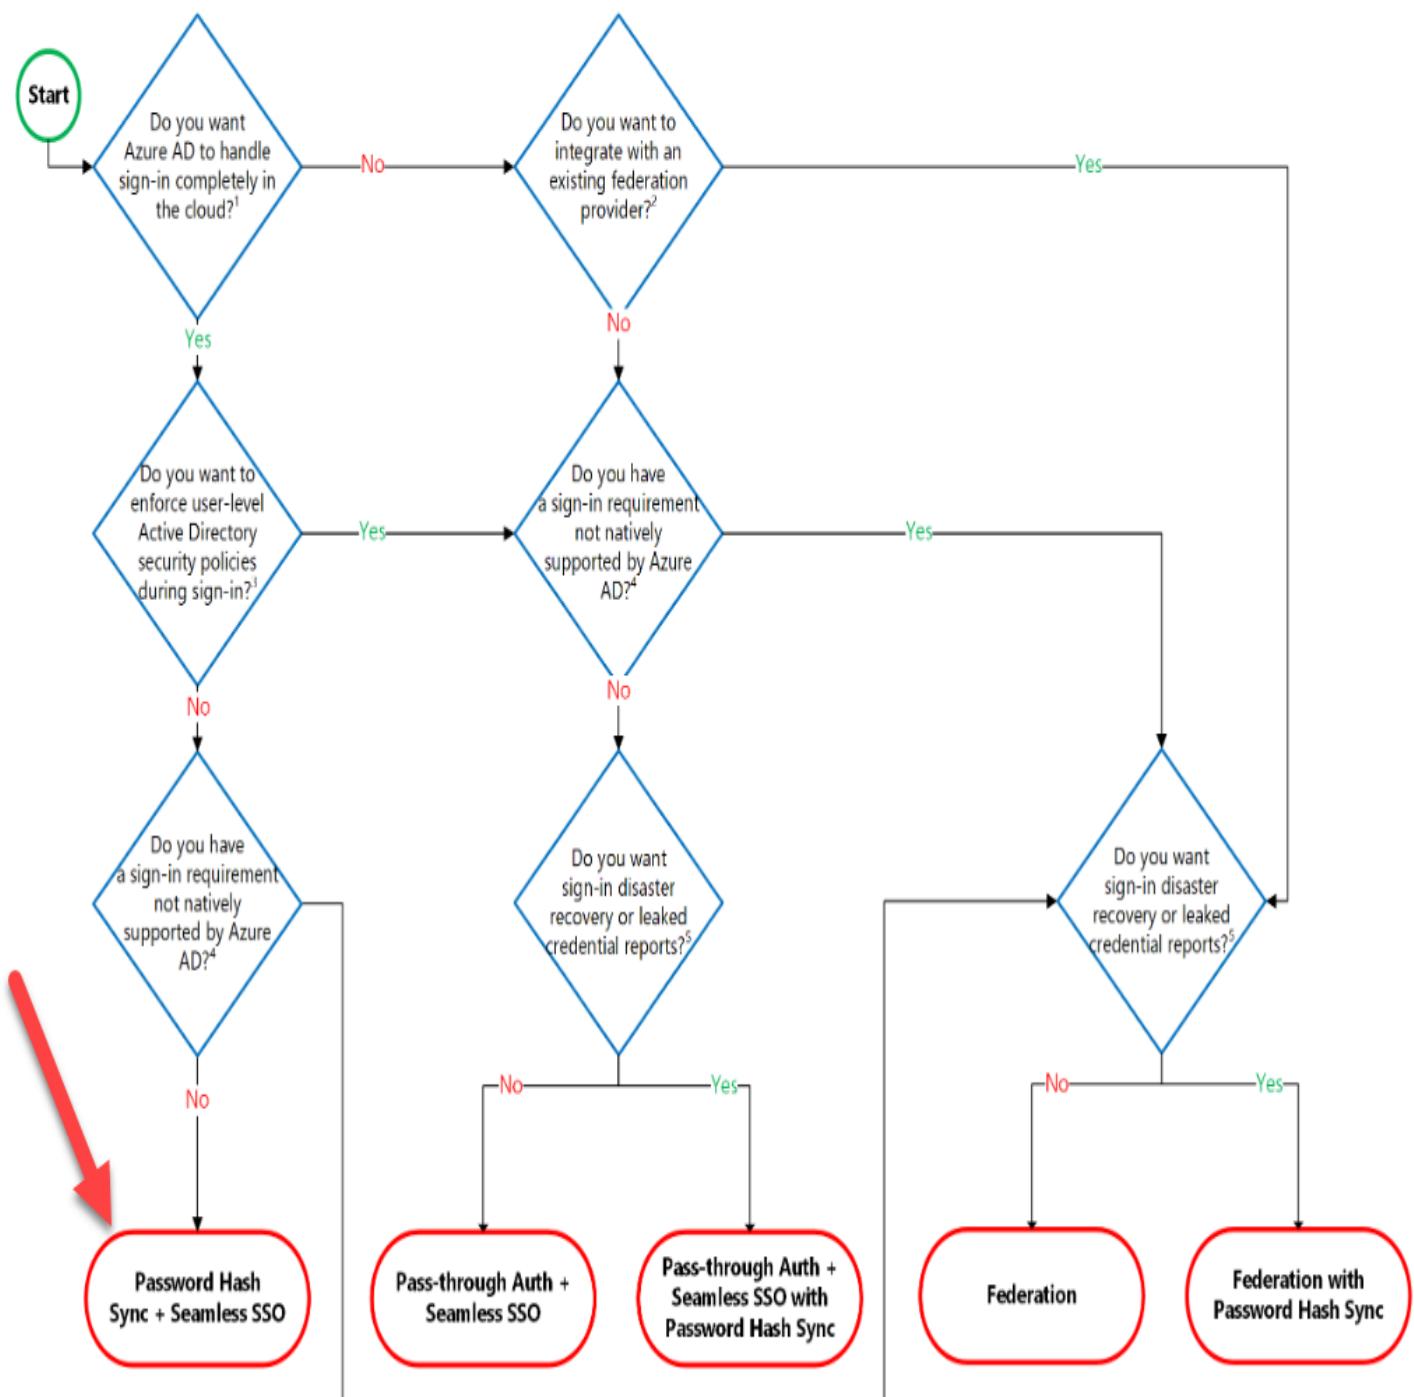
- A. Password Hash Sync 
- B. Pass-through authentication
- C. Federation
- D. Federation with Password Hash Sync

**Explanation:**

Answer – A

A clear diagram on the decision for the type of synchronization to choose is given in the Microsoft documentation as shown below

# Decision tree



Based on the decision tree we have to choose "Password Hash Sync" as the option and hence all other options are incorrect

For more information on choosing the right authentication mechanism, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/active-directory/hybrid/choose-ad-authn>

Ask our Experts

Rate this Question?

View Queries

open ▾

Domain :Design for identity and security

A company is developing an application that will be hosted using the Azure Kubernetes Service. The application modules need to transmit files to a destination location and must make use of an encryption key for the encryption of the files before transmission. Which of the following will be used for managing the storage of the encryption keys?

- A. Azure Key Vault 
- B. Azure Storage Service Encryption
- C. A Kubernetes deployment YAML file
- D. A ConfigureMap service

**Explanation:**

Answer – A

In order to store encryption keys in this scenario, we work with Azure Key Vault.

- Keys are stored in a vault and invoked by URI when needed.
- Keys are safeguarded by Azure, using industry standard algorithms, key lengths, and hardware security modules.
- Keys are processed in HSMs that reside in the same Azure datacenters as the applications.

To know how, do refer the following link.

<https://github.com/Azure/acs-engine/blob/master/docs/kubernetes/features.md#azure-key-vault-data-encryption>

[Ask our Experts](#)

Rate this Question?  

[View Queries](#)[open ▾](#)

Domain :Determine workload requirements

A company is planning on deploying a stateless based application based on microservices using the Azure Service Fabric service. You need to design the infrastructure that would be required in the Azure Service Fabric service. Which of the following should you consider? Choose 2 answers from the applications given below

- A. The number of node types in the cluster 
- B. The properties for each node type 

C. The network connectivity

D. The service tier

### Explanation:

Answer – A and B

The Microsoft documentation specifies what are the Service Fabric cluster capacity planning considerations

# Service Fabric cluster capacity planning considerations

06/27/2018 • 16 minutes to read • Contributors  all

For any production deployment, capacity planning is an important step. Here are some of the items that you have to consider as a part of that process.

- The number of node types your cluster needs to start out with
- The properties of each of node type (size, primary, internet facing, number of VMs, etc.)
- The reliability and durability characteristics of the cluster

Since this is clearly mentioned in the Microsoft documentation, all other options are incorrect

For more information on the cluster capacity for Azure Service Fabric, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/service-fabric/service-fabric-cluster-capacity>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 49

Unattempted

Domain :Design for deployment, migration, and integration

A company currently has the following systems running on their on-premise environment

- An ASP.NET application running on Internet Information Services
- A MongoDB database

The company wants to migrate the systems onto Azure. They want to ensure to use managed services to reduce the administrative overhead. They want to minimize the time for migration as well and also reduce

costs wherever possible.

Which of the following Azure service would you use for the ASP.Net application?

- A. Azure Virtual Machines
- B. Azure Functions
- C. Azure App Service
- D. Azure Virtual Machine Scale sets

#### Explanation:

Answer – C

The ideal service to use is the Azure Web App in the Azure App Service. The Microsoft documentation mentions the following

## App Service overview

01/04/2017 • 2 minutes to read • Contributors 

Azure App Service is a service for hosting web applications, REST APIs, and mobile back ends. You can develop in your favorite language, be it .NET, .NET Core, Java, Ruby, Node.js, PHP, or Python. Applications run and scale with ease on both Windows and Linux-based environments. For Linux-based environments, see [App Service on Linux](#).

App Service not only adds the power of Microsoft Azure to your application, such as security, load balancing, autoscaling, and automated management. You can also take advantage of its DevOps capabilities, such as continuous deployment from Azure DevOps, GitHub, Docker Hub, and other sources, package management, staging environments, custom domain, and SSL certificates.

With App Service, you pay for the Azure compute resources you use. The compute resources you use is determined by the *App Service plan* that you run your apps on. For more information, see [Azure App Service plans overview](#).

Options A and D are incorrect since these would less cost effective and not a managed service

Option B is incorrect since this is used for running small pieces of code or functions on Azure cloud

For more information on the Azure App service, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/app-service/overview>

Ask our Experts

Rate this Question?  

**Question 50****Unattempted****Domain :Design for deployment, migration, and integration**

A company currently has the following systems running on their on-premise environment

- An ASP.Net application running on Internet Information Services
- A MongoDB database

The company wants to migrate the systems onto Azure. They want to ensure to use managed services to reduce the administrative overhead. They want to minimize the time for migration as well and also reduce costs wherever possible.

Which of the following Azure service would you use for the MongoDB database?

- A. **CosmosDB** 
- B. **Azure SQL Database**
- C. **Virtual Machines**
- D. **Azure SQL Data warehouse**

**Explanation:**

Answer – A

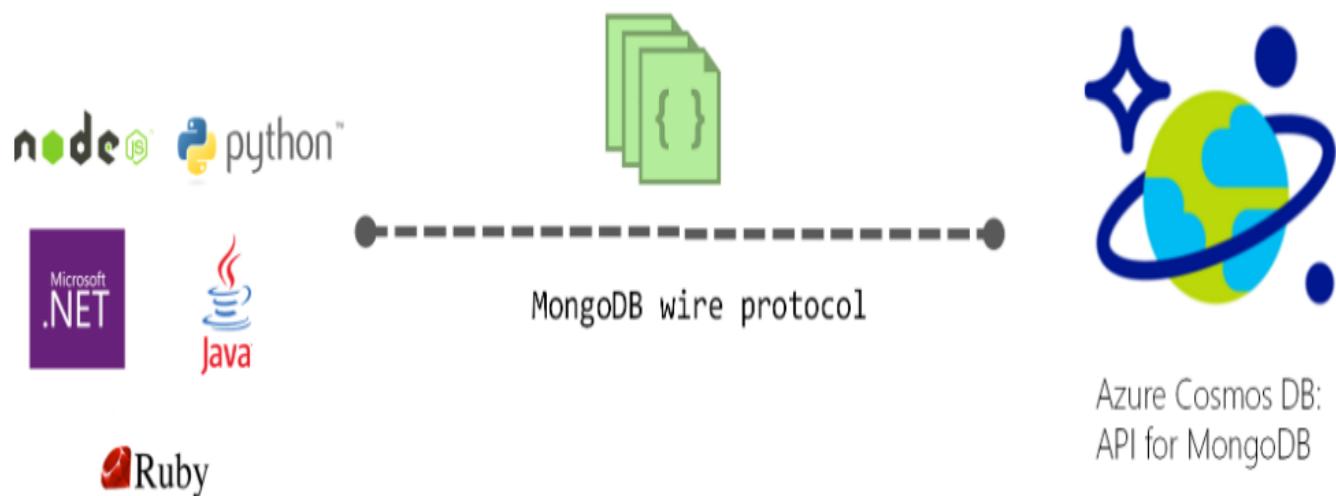
You can use the MongoDB API which is available as part of CosmosDB

The Microsoft documentation mentions the following

# Wire protocol compatibility

Azure Cosmos DB implements wire protocols of common NoSQL databases including Cassandra, MongoDB, Gremlin, and Azure Tables Storage. By providing a native implementation of the wire protocols directly and efficiently inside Cosmos DB, it allows existing client SDKs, drivers, and tools of the NoSQL databases to interact with Cosmos DB transparently. Cosmos DB does not use any source code of the databases for providing wire-compatible APIs for any of the NoSQL databases.

By default, Azure Cosmos DB's API for MongoDB is compatible with version 3.2 of the MongoDB's wire protocol. Features or query operators added in version 3.4 of the wire protocol are currently available as a preview feature. Any MongoDB client driver that understands these protocol versions should be able to natively connect to Cosmos DB.



Options B and D are incorrect since these are SQL based data stores

Option C is incorrect since this would less cost effective and not a managed service

For more information on CosmosDB and the MongoDB API, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/cosmos-db/mongodb-introduction>

Ask our Experts

Rate this Question?

View Queries

open ▾

Question 51

Unattempted

Domain :Design an infrastructure strategy

A company needs to create a storage account as shown below

## PROJECT DETAILS

Select the subscription to manage deployed resources and costs. Use resource groups like folders to organize and manage all your resources.

* Subscription	Pay-As-You-Go
* Resource group	whizlabs-rg
	<a href="#">Create new</a>

## INSTANCE DETAILS

The default deployment model is Resource Manager, which supports the latest Azure features. You may choose to deploy using the classic deployment model instead. [Choose classic deployment model](#)

* Storage account name <a href="#">?</a>	whizlabsstore
* Location	UK South

The storage account must meet the following requirements

- Ensure that the documents can be stored in the storage account
- The documents must be made accessible via drive mappings from Azure virtual machines that run Windows Server 2016.
- Provide the highest possible redundancy for the documents.
- Minimize storage access costs.

### Objective - Design an infrastructure strategy

Which of the following would you choose for the performance of the Storage account?

- A. Standard
- B. Premium
- C. General Purpose
- D. BLOB

### Explanation:

Answer – A

You have to use the Standard performance tier. The premium tier only allows storage of Page BLOB's.

In the below diagram you can see that if we choose Premium as the Performance option, we will have no option for files shares. And that is the main requirement. File shares allow to have drive mappings from Azure Virtual Machines.

## INSTANCE DETAILS

The default deployment model is Resource Manager, which supports the latest Azure features. You may choose to deploy using the classic deployment model instead. [Choose classic deployment model](#)

\* Storage account name  ✓

\* Location  ↗

Performance  Standard  Premium 1

Account kind  ↗

Replication  ↗

Accounts with the selected kind, replication and performance type only support page blobs. Block blobs, append blobs, file shares, tables, and queues will not be available. 2

Since this is clear from the implementation, all other options are incorrect

For more information on Azure storage accounts, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/storage/common/storage-account-overview>

Ask our Experts

Rate this Question? 😊 😢

View Queries

open ↗

Question 52

Unattempted

Domain :Design an infrastructure strategy

A company needs to create a storage account as shown below

### PROJECT DETAILS

Select the subscription to manage deployed resources and costs. Use resource groups like folders to organize and manage all your resources.

\* Subscription  ↗

\* Resource group  ↗

[Create new](#)

### INSTANCE DETAILS

The default deployment model is Resource Manager, which supports the latest Azure features. You may choose to deploy using the classic deployment model instead. [Choose classic deployment model](#)

\* Storage account name  ✓

\* Location  ↗

The storage account must meet the following requirements

- Ensure that the documents can be stored in the storage account
- The documents must be made accessible via drive mappings from Azure virtual machines that run Windows Server 2016.
- Provide the highest possible redundancy for the documents.
- Minimize storage access costs.

Which of the following would you choose for the Account kind?

- A. General Purpose v1
- B. General Purpose v2
- C. BLOBStorage
- D. ObjectStorage

#### Explanation:

Answer – B

If you need to use a storage account to host file shares, the consider General Purpose v2 which has support for Page BLOB's. Microsoft recommends this storage account kind which has more features than General Purpose v2.

## Types of storage accounts

Azure Storage offers three types of storage accounts. Each type supports different features and has its own pricing model. Consider these differences before you create a storage account to determine the type of account that is best for your applications. The three types of storage accounts are:

- **General-purpose v2 accounts:** Basic storage account type for blobs, files, queues, and tables. Recommended for most scenarios using Azure Storage.
- **General-purpose v1 accounts:** Legacy account type for blobs, files, queues, and tables. Use general-purpose v2 accounts instead when possible.
- **Blob storage accounts:** Blob-only storage accounts. Use general-purpose v2 accounts instead when possible.

In the Azure portal, you would configure this as shown below

## INSTANCE DETAILS

The default deployment model is Resource Manager, which supports the latest Azure features. You may choose to deploy using the classic deployment model instead. [Choose classic deployment model](#)

\* Storage account name  ✓

\* Location

Performance  Standard  Premium

Account kind  ▾

Replication  ▾

Access tier (default)  Cool  Hot

Since this is clear from the implementation, all other options are incorrect

For more information on Azure storage accounts, please go ahead and visit the below URL

- <https://docs.microsoft.com/en-us/azure/storage/common/storage-account-overview>

Ask our Experts

Rate this Question?  

**View Queries**

open ▾

**Question 53**

Unattempted

Domain :Design a business continuity strategy

A company needs to create a storage account as shown below

## PROJECT DETAILS

Select the subscription to manage deployed resources and costs. Use resource groups like folders to organize and manage all your resources.

* Subscription	Pay-As-You-Go
* Resource group	whizlabs-rg
	<a href="#">Create new</a>

## INSTANCE DETAILS

The default deployment model is Resource Manager, which supports the latest Azure features. You may choose to deploy using the classic deployment model instead. [Choose classic deployment model](#)

* Storage account name <a href="#">?</a>	whizlabsstore
* Location	UK South

The storage account must meet the following requirements

- Ensure that the documents can be stored in the storage account
- The documents must be made accessible via drive mappings from Azure virtual machines that run Windows Server 2016.
- Provide the highest possible redundancy for the documents.
- Minimize storage access costs.

Which of the following would you choose for the replication?

- A. Locally redundant storage (LRS)
- B. Zone-redundant storage (ZRS)
- C. Geo-redundant storage (GRS) ✓
- D. Read-access geo-redundant storage (RA-GRS)

## Explanation:

Answer – C

Please refer to

<https://docs.microsoft.com/en-us/azure/storage/common/storage-introduction#redundancy>

## Redundancy

In order to ensure that your data is durable, Azure Storage replicates multiple copies of your data. When you set up your storage account, you select a redundancy option.

Replication options for a storage account include:

- **Locally-redundant storage (LRS):** A simple, low-cost replication strategy. Data is replicated within a single storage scale unit.
- **Zone-redundant storage (ZRS):** Replication for high availability and durability. Data is replicated synchronously across three availability zones.
- **Geo-redundant storage (GRS):** Cross-regional replication to protect against region-wide unavailability.

- **Read-access geo-redundant storage (RA-GRS):** Cross-regional replication with read access to the replica.

Means redundancy of RA-GRS and GRS: both has same but availability of data is more in RA-GRS than GRS. In GRS, in case of region failure, data is unavailable till Microsoft trigger failover but in case of RA-GRS: Data for read-only, is available from secondary region all the time

Since question is also talking about "Minimize cost and as we know cost of GRS is comparatively lower than RA-GRS."

Now please refer to

- <https://docs.microsoft.com/en-us/azure/storage/common/storage-redundancy-grs#read-access-geo-redundant-storage>

?

## Geo-redundant storage (GRS): Cross-regional replication for Azure Storage

10/19/2018 • 5 minutes to read • 

Geo-redundant storage (GRS) is designed to provide at least 99.99999999999999% (16 9's) durability of objects over a given year by replicating your data to a secondary region that is hundreds of miles away from the primary region. If your storage account has GRS enabled, then your data is durable even in the case of a complete regional outage or a disaster in which the primary region isn't recoverable.

If you opt for GRS, you have two related options to choose from:

- GRS replicates your data to another data center in a secondary region, but that data is available to be read only if Microsoft initiates a failover from the primary to secondary region.
- Read-access geo-redundant storage (RA-GRS) is based on GRS. RA-GRS replicates your data to another data center in a secondary region, and also provides you with the option to read from the secondary region. With RA-GRS, you can read from the secondary region regardless of whether Microsoft initiates a failover from the primary to secondary region.

For a storage account with GRS or RA-GRS enabled, all data is first replicated with locally redundant storage (LRS). An update is first committed to the primary location and replicated using LRS. The update is then replicated asynchronously to the secondary region using GRS. When data is written to the secondary location, it's also replicated within that location using LRS.

Both the primary and secondary regions manage replicas across separate fault domains and upgrade domains within a storage scale unit. The storage scale unit is the basic replication unit within the datacenter. Replication at this level is provided by LRS; for more information, see [Locally redundant storage \(LRS\): Low-cost data redundancy for Azure Storage](#).

[Ask our Experts](#)

Rate this Question?  

[View Queries](#)

open ▾

Question 54

Unattempted

A testing team needs to create resources in an Azure subscription.

These resources will be pertinent to applications that need to be tested in Azure. These applications will make use of shared resources and application specific resources.

You have to create a deployment solution that meets the following requirements

- Minimize the administrative effort for the testing team to create the application environment
- Ensure that the application environment can be created consistently multiple times
- When the testing of the application is complete, all resources except for the shared resources should be deleted

Which of the following should be used to implement this requirement?

- A. Use JSON templates to create the resources. Place the shared resources in one resource group and the application specific resources in a separate resource group
- B. Use JSON templates to create the resources. Place the shared resources and application specific resources in one resource group.
- C. Use Azure Powershell scripts to create the resources. Place the shared resources in one resource group and the application specific resources in a separate resource group
- D. Use Azure Powershell scripts to create the resources. Place the shared resources and application specific resources in one resource group.

#### **Explanation:**

Answer – A

The ideal solution is to use Template deployment. Below is why Microsoft recommends a template deployment.

## **Template deployment**

With Resource Manager, you can create a template (in JSON format) that defines the infrastructure and configuration of your Azure solution. By using a template, you can repeatedly deploy your solution throughout its lifecycle and have confidence your resources are deployed in a consistent state.

Also ensure separate resource group deployments. So, when the application needs to be deleted, you can just delete the resource group, so that all resources in the resource group get deleted.

Option B is incorrect since you should deploy shared resources and application specific resources to separate resources groups

Options C and D are incorrect since you should use JSON based templates for deployment

For more information on Resources Groups, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/azure-resource-manager/resource-group-overview>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 55

Unattempted

Domain :Determine workload requirements

A company currently has an Azure account and subscription. They want to host an application using Virtual Machines and a load balancer. There is a requirement to ensure that the application is made available 99.99% of the time. Which of the following would need to be in place? You also have to minimize costs associated with the solution. Choose 2 answers from the options given below

- A. Create a Basic Load balancer
- B. Create a Standard Load balancer 
- C. Add 2 Virtual Machines to the backend pool 
- D. Add a Virtual Machine to the backend pool

#### Explanation:

Answer – B and C

This is clearly mentioned in the Microsoft documentation

## SLA for Load Balancer

Last updated: March 2018

We guarantee that a Load Balanced Endpoint using Azure Standard Load Balancer, serving two or more Healthy Virtual Machine Instances, will be available 99.99% of the time.

Basic Load Balancer is excluded from this SLA.

Since this is clearly given in the Microsoft documentation, all other options are incorrect

For more information on the SLA for the Load balancer, please go to the below URL

[https://azure.microsoft.com/en-us/support/legal/sla/load-balancer/v1\\_0/](https://azure.microsoft.com/en-us/support/legal/sla/load-balancer/v1_0/)

Ask our Experts

Rate this Question?  

View Queries

open ▾

## Certification

Cloud Certification

Java Certification

PM Certification

Big Data Certification

## Company

Support

Discussions

Blog

Business

## Follow us



---

© Copyright 2020. Whizlabs Software Pvt. Ltd. All Right Reserved.