

Practice Test 5

Completed on 14-June-2020



Attempt

01



Marks Obtained

0 / 55



Your score

0.0%



Time Taken

00 H 00 M 03 S



Result

Failed

Domains wise Quiz Performance Report

1

Design a data platform solution

11

0

0

11

0

2

Design a business continuity strategy

14

0

0

14

0

3

Design for deployment, migration, and integration

5

0

0

5

0

4

Design for identity and security

8

0

0

8

0

5

Determine workload requirements

10

0

0

10

0

6

Design an infrastructure strategy

7

0

0

7

0

Total

All Domain

55

0

0

55

0

Review the Answers

Sorting by

All



Question 1

Unattempted

A company currently has an on-premise setup in place. They want to migrate a set of their on-premise SQL Servers to Azure. They want to make use of their existing SQL Server license that they have as part of their Software Assurance contract with Microsoft. This would help them make use of their existing investment. They decide to set up Azure Virtual Machines and install the SQL Server engine on those machines. Would this fulfill the requirement?

- A. Yes ✓
- B. No

Explanation:

Answer – A

Since here you are fully in control of the underlying virtual machine, you can install the SQL Server engine and make use of your existing licences.

For more information on Azure virtual machines, please visit the below URL

<https://azure.microsoft.com/en-us/services/virtual-machines/>

[Ask our Experts](#)

Rate this Question?  

View Queries

open ▾

Question 2

Unattempted

Domain :Design a data platform solution

A company currently an on-premise setup in place. They want to migrate a set of their on-premise SQL Servers to Azure. They want to make use of their existing SQL server licences that they have as part of their Software Assurance contract with Microsoft. This would help them make use of their existing investment. They decide to use the Azure SQL database service along with the DTU licensing model. Would this fulfil the requirement?

- A. Yes
- B. No ✓

Explanation:

Answer – B

The DTU licensing model would be less effective as you won't be able to make use of your existing licences.

For more information on the DTU licensing model, please visit the below URL

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 3

Unattempted

Domain :Design a data platform solution

A company currently has an on-premise setup in place. They want to migrate a set of their on-premise SQL Servers to Azure. They want to make use of their existing SQL server licences that they have as part of their Software Assurance contract with Microsoft. This would help them make use of their existing investment. They decide to use the Azure SQL database service along with the v-core licensing model. Would this fulfil the requirement?

- A. Yes 
- B. No

Explanation:

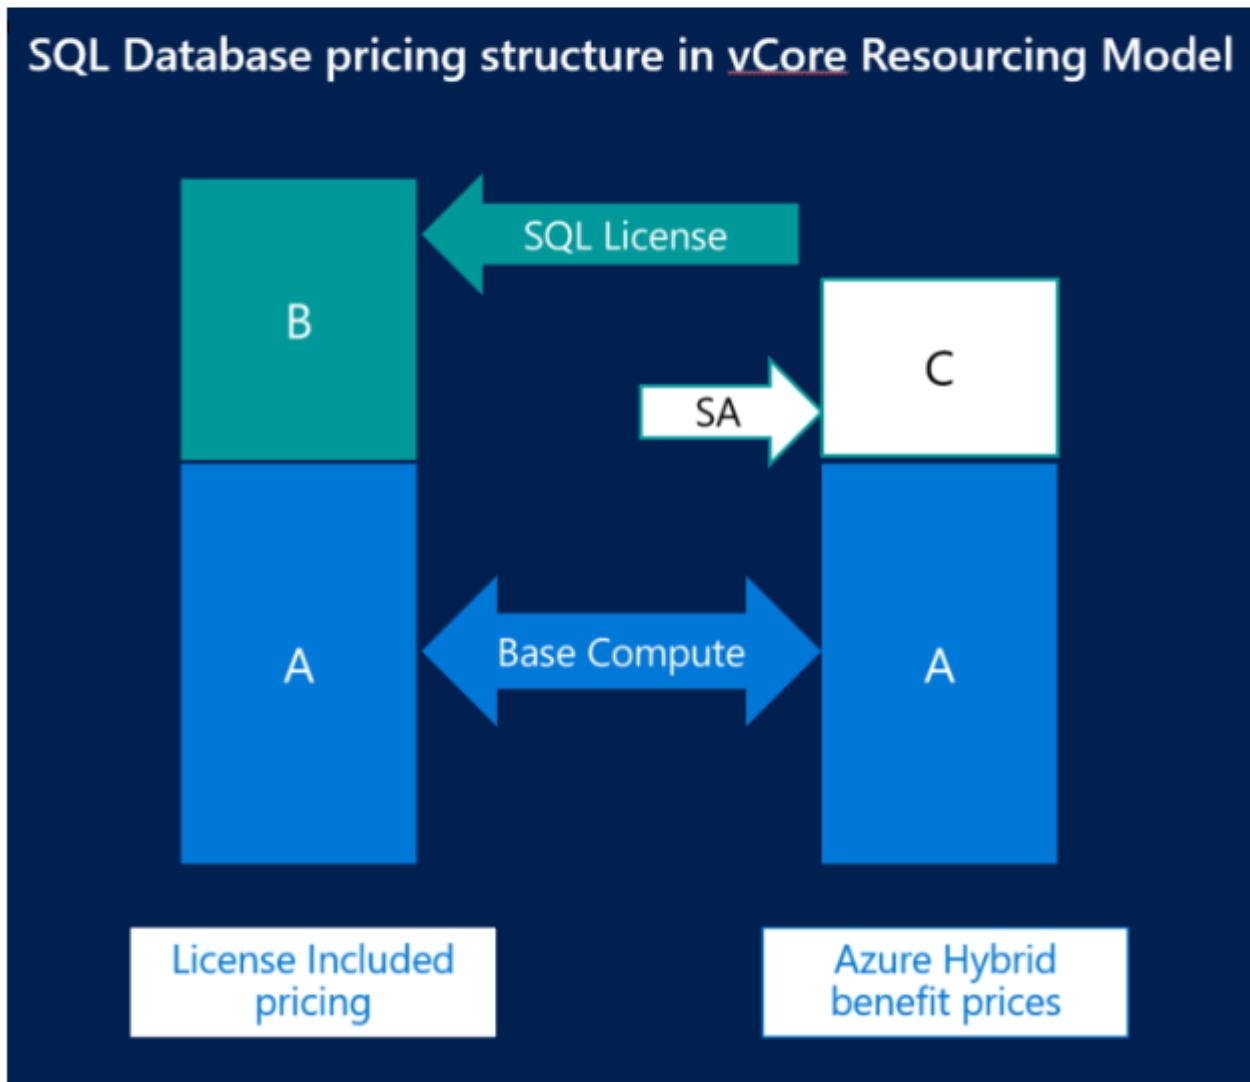
Answer – A

Yes, this would work. In this model, you can make use of the Azure Hybrid Benefit. Here you can make use of your existing licences.

The Microsoft documentation mentions the following

Azure Hybrid Benefit

In the provisioned compute tier of the vCore-based purchasing model, you can exchange your existing licenses for discounted rates on SQL Database by using [Azure Hybrid Benefit for SQL Server](#). This Azure benefit allows you to save up to 30 percent on Azure SQL Database by using your on-premises SQL Server licenses with Software Assurance.



For more information on the v-core licensing model, please visit the below URL

<https://docs.microsoft.com/en-us/azure/sql-database/sql-database-service-tiers-vcore>

Ask our Experts

Rate this Question?

View Queries

open ▾

Question 4

Unattempted

Domain :Design a business continuity strategy

A company is planning on deploying an application onto a set of virtual machines. They are going to spin it up as part of a virtual machine scale set



- Operating System disk image – Windows Server 2016
- Availability Zone – None
- Instance count – 2
- Instance Size – Standard DS1 v2
- Load Balancing Options – None
- Autoscale - Enabled

Below are the key requirements of the application

- Scale Instances based on CPU threshold going beyond 60%
- Provide redundancy if the primary Azure region goes down
- Distribute incoming user traffic

Does this architecture support the scaling requirement?

A. Yes

B. No

Explanation:

Answer – A

Azure virtual machine scale sets let you create and manage a group of load balanced VMs. The number of VM instances can automatically increase or decrease in response to demand or a defined schedule. Scale sets provide high availability to your applications, and allow you to centrally manage, configure, and update many VMs.

The Microsoft documentation mentions the following

What are virtual machine scale sets?

05/21/2018 • 3 minutes to read •

Azure virtual machine scale sets let you create and manage a group of identical, load balanced VMs. The number of VM instances can automatically increase or decrease in response to demand or a defined schedule. Scale sets provide high availability to your applications, and allow you to centrally manage, configure, and update a large number of VMs. With virtual machine scale sets, you can build large-scale services for areas such as compute, big data, and container workloads.

For more information on virtual machine scale sets, please visit the below URL

<https://docs.microsoft.com/en-us/azure/virtual-machine-scale-sets/overview>

[Ask our Experts](#)Rate this Question?  [View Queries](#)

open ▾

Question 5**Unattempted****Domain :Design a business continuity strategy**

A company is planning on deploying an application onto a set of virtual machines. They are going to spin it up as part of a virtual machine scale set

- 
- Operating System disk image – Windows Server 2016
 - Availability Zone – None
 - Instance count – 2
 - Instance Size – Standard DS1 v2
 - Load Balancing Options – None
 - Autoscale - Enabled

Below are the key requirements of the application

- Scale Instances based on CPU threshold going beyond 60%
- Provide redundancy if the primary Azure region goes down
- Distribute incoming user traffic

Does this architecture support the redundancy requirement?

A. Yes

B. No 

Explanation:

Answer – B

The virtual machine scale sets can only provide the scaling feature and not redundancy. You can use the Traffic Manager service to provide redundancy by routing traffic to a secondary location.

For more information on the Azure Traffic Manager service, please visit the below URL

<https://docs.microsoft.com/en-us/azure/traffic-manager/traffic-manager-overview>

[Ask our Experts](#)Rate this Question?  

Question 6

Unattempted

Domain :Design a business continuity strategy

A company is planning on deploying an application onto a set of virtual machines. They are going to spin it up as part of a virtual machine scale set



- Operating System disk image – Windows Server 2016
- Availability Zone – None
- Instance count – 2
- Instance Size – Standard DS1 v2
- Load Balancing Options – None
- Autoscale - Enabled

Below are the key requirements of the application

- Scale Instances based on CPU threshold going beyond 60%
- Provide redundancy if the primary Azure region goes down
- Distribute incoming user traffic

Does this architecture support the traffic distribution requirement?

A. Yes

B. No

Explanation:

Answer – B

Since there is no Load balancer or Application Gateway placed in front of the scale set, traffic distribution is not possible with this setup.

For more information on virtual machine scale sets, please visit the below URL

<https://docs.microsoft.com/en-us/azure/virtual-machine-scale-sets/overview>

Ask our Experts

Rate this Question?

Domain :Design a business continuity strategy

A team is setting the below policy for the backup of Azure virtual machines

DefaultPolicy

Associated items Delete Save Discard

Backup schedule

* Frequency * Time * Timezone
Daily 17:00 (UTC) Coordinated Universal Time

Instant Restore

Retain instant recovery snapshot(s) for

2 Day(s)

Retention range

Retention of daily backup point.

* At For
17:00 30 Day(s)

Retention of weekly backup point.

* On * At For
Sunday 17:00 12 Week(s)

Retention of monthly backup point.

Week Based Day Based

* On * Day * At For
First Sunday 17:00 12 Month(s)

Retention of yearly backup point.

Week Based Day Based

* In * On * Day * At For
January First Sunday 17:00 2 Year(s)

What is the minimum recovery point objective for virtual machines backed up by this policy?

A. 1 hour

B. 1 day

C. 12 weeks

D. 12 months

Explanation:

Answer – B

Since there is a backup being taken every day, the minimum recovery point objective is one day.

Recovery point objective is the point to which you can recover back to based on the backup's being taken. Since the backups are being taken every day, there would be a recovery point available on a daily basis.

Since this is evident from the settings of the Azure policy, all other options are incorrect

For more information on the Azure backup architecture, please visit the below URL

<https://docs.microsoft.com/en-us/azure/backup/backup-architecture>

[Ask our Experts](#)

Rate this Question?  

[View Queries](#)

open ▾

Question 8

Unattempted

Domain :Design a business continuity strategy

A team is setting the below policy for the backup of Azure virtual machines

DefaultPolicy

[Associated items](#) [Delete](#) [Save](#) [Discard](#)

Backup schedule

* Frequency * Time * Timezone

Daily 17:00 (UTC) Coordinated Universal Time

Instant Restore ⓘ

Retain instant recovery snapshot(s) for

2 Day(s)

Retention range

Retention of daily backup point.

* At For

17:00 30 Day(s)

Retention of weekly backup point.

* On * At For

Sunday 17:00 12 Week(s)

Retention of monthly backup point.

[Week Based](#) [Day Based](#)

* On * Day * At For

First Sunday 17:00 12 Month(s)

Retention of yearly backup point.

[Week Based](#) [Day Based](#)

* In * On * Day * At For

January First Sunday 17:00 2 Year(s)

Which of the following is the maximum time that a virtual machine would be available for recovery based on the policy?

- A. 30 days
- B. 12 weeks
- C. 12 months
- D. 2 years

Explanation:

Answer – D

Since there is a yearly recovery point in place, that would be retained for 2 years, this would be the maximum time over which the backup of the virtual machine would be available for.

DefaultPolicy

Associated items Delete Save Discard

Backup schedule

* Frequency * Time * Timezone
Daily 17:00 (UTC) Coordinated Universal Time

Instant Restore

Retain instant recovery snapshot(s) for

2 Day(s)

Retention range

Retention of daily backup point.

* At For
17:00 30 Day(s)

Retention of weekly backup point.

* On * At For
Sunday 17:00 12 Week(s)

Retention of monthly backup point.

Week Based Day Based

* On * Day * At For
First Sunday 17:00 12 Month(s)

Retention of yearly backup point.

Week Based Day Based

* In * On * Day * At For
January First Sunday 17:00 2 Year(s)



Since this is evident from the settings of the Azure policy, all other options are incorrect

For more information on the Azure backup architecture, please visit the below URL

<https://docs.microsoft.com/en-us/azure/backup/backup-architecture>

View Queries

open ▾

Question 9

Unattempted

Domain :Design for deployment, migration, and integration

A company wants to migrate a set of resources from their on-premise data center onto Azure. Below are the requirements for each application

Name
Requirement

whizlabappA

This application is based entirely on docker container. The underlying service must provide the ability to manage the container and also provide automatically scaling.

whizlabappB

This is a web-based application. The application needs to be ported to a service which can scale automatically. Minimum administrative overhead is expected.

whizlabappC

This application needs to run on demand. The solution must focus on minimizing costs.

Which of the following would you consider for hosting the application "whizlabappA"?

- A. Azure Functions
- B. Azure Logic Apps
- C. Azure Kubernetes
- D. Azure Virtual Machines
- E. Azure Web Apps

Explanation:

Answer – C

If you are looking for a fully managed service to host docker based applications, then look towards using the Azure Kubernetes service

The Microsoft documentation mentions the following

Home / Kubernetes

What is Kubernetes?

Kubernetes is open-source orchestration software for deploying, managing and scaling containers

Option A is incorrect because this is a serverless service that is used to run code on Azure

Option B is incorrect because this is a workflow-based service

Option D is incorrect because even though you can host Kubernetes on a virtual machine, it wouldn't be a fully managed service

Option E is incorrect because this is used for hosting web applications. It can host docker based web applications, but if you want a service to manage the contains, using Kubernetes is more ideal

For more information on Kubernetes, please visit the below URL

<https://azure.microsoft.com/en-in/topic/what-is-kubernetes/>

Ask our Experts

Rate this Question?

View Queries

open ▾

Question 10

Unattempted

Domain :Design for deployment, migration, and integration

A company wants to migrate a set of resources from their on-premise data center onto Azure. Below are the requirements for each application

Name

Requirement

whizlabappA

This application is based entirely on docker containers. The underlying service must provide the ability to manage the containers and also provide automatic scaling.

whizlabappB

This is a web-based application. The application needs to be ported to a service which can scale automatically. Minimum administrative overhead is expected.

whizlabappC

This application needs to run on demand. The solution must focus on minimizing costs.

Which of the following would you consider for hosting the application "whizlabappB"?

- A. Azure Functions
- B. Azure Logic Apps
- C. Azure Kubernetes
- D. Azure Virtual Machines
- E. Azure Web Apps

Explanation:

Answer – E

You can use the managed service of Azure App Service to deploy your web-based application

The Microsoft documentation mentions the following

App Service overview

01/04/2017 • 2 minutes to read • 

Azure App Service is an HTTP-based service for hosting web applications, REST APIs, and mobile backends. You can develop in your favorite language, be it .NET, .NET Core, Java, Ruby, Node.js, PHP, or Python. Applications run and scale with ease on both Windows and Linux-based environments. For Linux-based environments, see [App Service on Linux](#).

App Service not only adds the power of Microsoft Azure to your application, such as security, load balancing, autoscaling, and automated management. You can also take advantage of its DevOps capabilities, such as continuous deployment from Azure DevOps, GitHub, Docker Hub, and other sources, package management, staging environments, custom domain, and SSL certificates.

Option A is incorrect because this is a serverless service that is used to run code on Azure

Option B is incorrect because this is a workflow-based service

Option C is incorrect because this is an orchestration-based service for container-based applications

Option D is incorrect because even though you can host Web applications on a virtual machine, it wouldn't be a fully managed service

For more information on Azure App service, please visit the below URL

<https://docs.microsoft.com/en-us/azure/app-service/overview>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 11

Unattempted

Domain :Design for deployment, migration, and integration

A company wants to migrate a set of resources from their on-premise data center onto Azure. Below are the requirements for each application

Name
Requirement

whizlabappA

This application is based entirely on docker containers. The underlying service must provide the ability to manage the containers and also provide automatic scaling.

whizlabappB

This is a web-based application. The application needs to be ported to a service which can scale automatically. Minimum administrative overhead is expected.

whizlabappC

This application needs to run on demand. The solution must focus on minimizing costs.

Which of the following would you consider for hosting the application "whizlabappC"?

- A. Azure Functions 
- B. Azure Logic Apps
- C. Azure Kubernetes
- D. Azure Virtual Machines
- E. Azure Web Apps

Explanation:

Answer – A

This is the ideal cost-effective solution to run code on demand on Azure.

The Microsoft documentation mentions the following

An introduction to Azure Functions

10/03/2017 • 4 minutes to read •  +13

Azure Functions is a solution for easily running small pieces of code, or "functions," in the cloud. You can write just the code you need for the problem at hand, without worrying about a whole application or the infrastructure to run it. Functions can make development even more productive, and you can use your development language of choice, such as C#, F#, Node.js, Java, or PHP. Pay only for the time your code runs and trust Azure to scale as needed. Azure Functions lets you develop serverless applications on Microsoft Azure.

Option B is incorrect because this is a workflow-based service

Option C is incorrect because this is an orchestration-based service for container-based applications

Option D is incorrect because even though you can host applications on a virtual machine, it would be less cost-effective if you just want to run code on demand.

Option E is incorrect because this is used as a platform for hosting web-based applications. We could use the web jobs feature, but this would not be effective from a costing perspective.

For more information on Azure Functions, please visit the below URL

<https://docs.microsoft.com/en-us/azure/azure-functions/functions-overview>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 12

Unattempted

Domain :Design for identity and security

A company currently has an on-premise setup that contains an Active directory instance. They have recently purchased an Azure subscription and an Azure tenant. They are planning on setting up Azure AD Connect to synchronize the users from their on-premise Active Directory onto Azure AD. They have the following requirements for this setup

- Enforcement of on-premise Active Directory security policies
- Provide the ability for Single-Sign On

Which of the following would you setup as the authentication mechanism?

- A. Active Directory authentication
- B. Pass-through authentication 
- C. Password hash synchronization
- D. Seamless authentication

Explanation:

Answer – B

If you need to enforce your on-premise Active Directory security policies, then you should choose "Pass-through authentication" as the authentication technique.

The Microsoft documentation mentions the following

Azure Active Directory (Azure AD) Pass-through Authentication allows your users to sign in to both on-premises and cloud-based applications using the same passwords. This feature provides your users a better experience - one less password to remember, and reduces IT helpdesk costs because your users are less likely to forget how to sign in. When users sign in using Azure AD, this feature validates users' passwords directly against your on-premises Active Directory.



Azure AD Pass-through Authentication and Seamless Single Sign-on

Pass-through Authentication

Azure AD Pass-through Authentication

and Seamless Single Sign-on. Full demos and how it all works

This feature is an alternative to [Azure AD Password Hash Synchronization](#), which provides the same benefit of cloud authentication to organizations. However, certain organizations wanting to enforce their on-premises Active Directory security and password policies, can choose to use Pass-through Authentication instead. Review [this guide](#) for a comparison of the various Azure AD sign-in methods and how to choose the right sign-in method for your organization.

Since this is clearly mentioned in the Microsoft documentation, all other options are incorrect

For more information on Pass-through authentication, please visit the below URL

<https://docs.microsoft.com/en-us/azure/active-directory/hybrid/how-to-connect-pta>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 13

Unattempted

Domain :Design for identity and security

A company currently has an on-premise setup that contains an Active directory instance. They have recently purchased an Azure subscription and an Azure tenant. They are planning on setting up Azure AD Connect to synchronize the users from their on-premise Active Directory onto Azure AD. They have the following requirements for this setup

- Enforcement of on-premise Active Directory security policies
- Provide the ability for Single-Sign On

Where would you configure the "Seamless Sign-on" setting?

- A. Active Directory
- B. Azure AD
- C. Azure AD Connect
- D. Azure AD Connect Health

Explanation:

Answer – C

You would configure this setting in Azure AD Connect. The Microsoft documentation mentions the steps for enabling this.

Step 2: Enable the feature

Enable Seamless SSO through [Azure AD Connect](#).

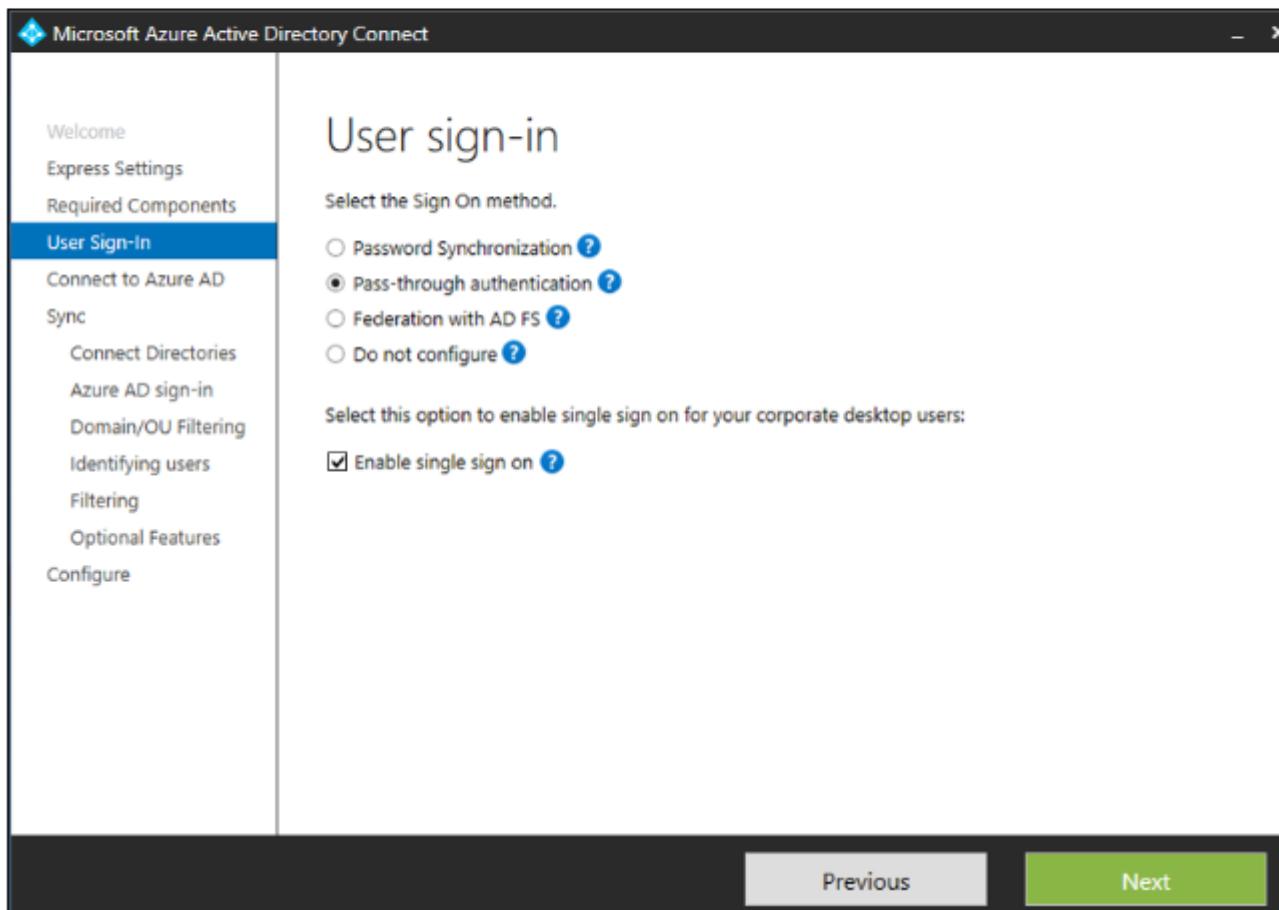
① Note

You can also [enable Seamless SSO using PowerShell](#) if Azure AD Connect doesn't meet your requirements. Use this option if you have more than one domain per Active Directory forest, and you want to be more targeted about the domain you want to enable Seamless SSO for.

If you're doing a fresh installation of Azure AD Connect, choose the [custom installation path](#). At the **User sign-in** page, select the **Enable single sign on** option.

① Note

The option will be available for selection only if the Sign On method is **Password Hash Synchronization** or **Pass-through Authentication**.



Since this is clearly mentioned in the Microsoft documentation, all other options are incorrect

For more information on configuring SSO, please visit the below URL

<https://docs.microsoft.com/en-us/azure/active-directory/hybrid/how-to-connect-sso-quick-start>

Ask our Experts

Rate this Question?

Question 14

Unattempted

Domain :Design a data platform solution

A company currently has a set of Azure SQL databases defined. They want to ensure the following when it comes to the Azure SQL Database instance

- Gain an insight onto all the database activity
- Help detect any suspected security violations

Which of the following feature within Azure SQL database could help fulfil these requirements?

- A. Locks
- B. Auditing
- C. Dynamic Data Masking
- D. Transparent Data Encryption

Explanation:

Answer – B

This could be accomplished with the Auditing feature. The Microsoft documentation mentions the following

Get started with SQL database auditing

04/16/2019 • 12 minutes to read •  +12

Auditing for Azure [SQL Database](#) and [SQL Data Warehouse](#) tracks database events and writes them to an audit log in your Azure storage account, OMS workspace or Event Hubs. Auditing also:

- Helps you maintain regulatory compliance, understand database activity, and gain insight into discrepancies and anomalies that could indicate business concerns or suspected security violations.
- Enables and facilitates adherence to compliance standards, although it doesn't guarantee compliance. For more information about Azure programs that support standards compliance, see the [Azure Trust Center](#) where you can find the most current list of SQL Database compliance certifications.

Since this is clearly mentioned in the Microsoft documentation, all other options are incorrect

For more information on database auditing, please visit the below URL

<https://docs.microsoft.com/en-us/azure/sql-database/sql-database-auditing>

[Ask our Experts](#)Rate this Question?  [View Queries](#)

open ▾

Question 15**Unattempted****Domain :Design a data platform solution**

A company currently has a set of Azure SQL databases defined. They want to ensure the following when it comes to the Azure SQL Database instance

- Gain an insight onto all the database activity
- Help detect any suspected security violations

Which of the following could be used to log all the database activity based on the solution chosen? Choose 3 answers from the options given below

- A. Azure Storage 
- B. Azure SQL Datawarehouse
- C. Azure Log Analytics 
- D. Azure Event Hub 

Explanation:

Answer – A, C and D

When you look at the options available when enabling auditing, you can see the options available.

- [Overview](#)
- [Activity log](#)
- [Tags](#)
- [Diagnose and solve problems](#)
- [Quick start](#)
- [Query editor \(preview\)](#)

- Settings**
- [Configure](#)
- [Geo-Replication](#)
- [Connection strings](#)
- [Sync to other databases](#)
- [Add Azure Search](#)
- [Properties](#)
- [Locks](#)
- [Export template](#)

- Security**
- [Advanced Data Security](#)
- [Auditing](#)

If Blob Auditing is enabled on the server, it will always apply to the database, regardless of the database settings.

[View server settings](#)

Server-level Auditing: **Disabled**

Auditing

ON **OFF**

Audit log destination (choose at least one):

Storage

Log Analytics (Preview)

Event Hub (Preview)

Turn on Advanced Data Security to receive security alerts upon suspicious events.

For more information on database auditing, please visit the below URL

<https://docs.microsoft.com/en-us/azure/sql-database/sql-database-auditing>

[Ask our Experts](#)

Rate this Question?

[View Queries](#)

open ▾

Question 16

Unattempted

Domain :Determine workload requirements

A company is planning on hosting a system in Azure. It would consist of a web application module and a processing module as shown below



The Web Application would be hosted using the Azure Web App service. The Processing module would be implemented using Azure Functions.

The company wants a system to be in place which could be used for exchange of messages between the Web Application and the Processing module. Which of the following could be used for this purpose?

- A. Azure Blob Storage
- B. Azure Table Storage
- C. Azure Queue Storage
- D. Azure File Storage

Explanation:

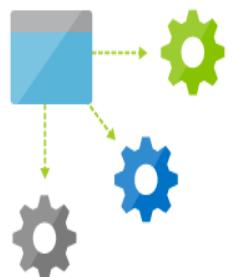
Answer – C

You can use Azure Queue storage for managing the messaging between application components.

The Microsoft documentation mentions the following

Decouple components

Use Azure Queue Storage to build flexible applications and separate functions for better durability across large workloads. When you design applications for scale, application components can be decoupled, so that they can scale independently. Queue storage gives you asynchronous message queuing for communication between application components, whether they are running in the cloud, on the desktop, on-premises or on mobile devices.



Since this is clear from the Microsoft documentation, all other options are incorrect

For more information on storage queues, please visit the below URL

<https://azure.microsoft.com/en-in/services/storage/queues/>

Ask our Experts

Rate this Question?

View Queries

open ▾

A company is planning on hosting a system in Azure. It would consist of a web application module and a processing module as shown below



The Web Application would be hosted using the Azure Web App service. The Processing module would be implemented using Azure Functions.

The company wants to have a cost-effective storage platform for storing user data. Which of the following could be used for this purpose?

- A. Azure Blob Storage
- B. Azure Table Storage
- C. Azure Queue Storage
- D. Azure File Storage

Explanation:

Answer – B

You can use Azure table storage as a cost-effective data store for the user data.

The Microsoft documentation mentions the following

Azure Table storage is a service that stores structured NoSQL data in the cloud, providing a key/attribute store with a schemaless design. Because Table storage is schemaless, it's easy to adapt your data as the needs of your application evolve. Access to Table storage data is fast and cost-effective for many types of applications, and is typically lower in cost than traditional SQL for similar volumes of data.

You can use Table storage to store flexible datasets like user data for web applications, address books, device information, or other types of metadata your service requires. You can store any number of entities in a table, and a storage account may contain any number of tables, up to the capacity limit of the storage account.

Since this is clear from the Microsoft documentation, all other options are incorrect

For more information on Table storage, please visit the below URL

<https://docs.microsoft.com/en-us/azure/cosmos-db/table-storage-overview>

View Queries

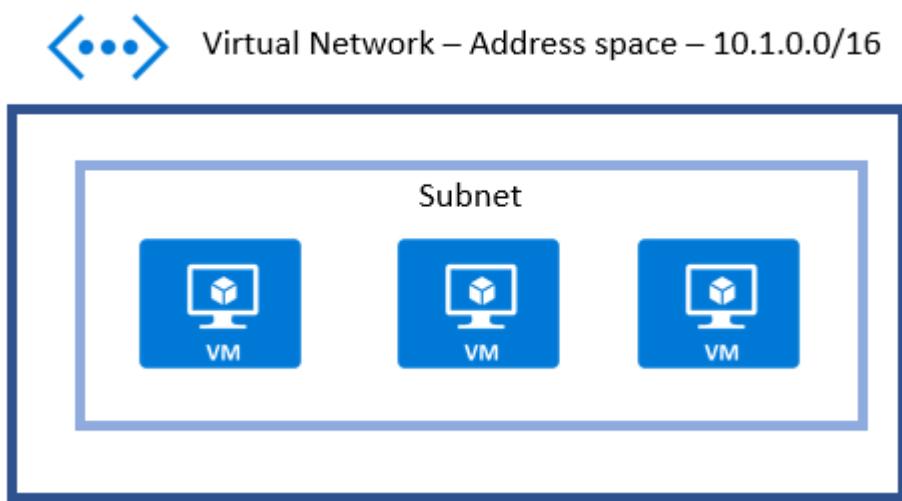
open ▾

Question 18

Unattempted

Domain :Design an infrastructure strategy

A company currently has the following deployment for an application in Azure



Network Security Groups are applied at both a subnet level and at the network interface level. The Networking team wants to use the Network Watcher tool to diagnose various issues when it comes to the networking aspect.

Which of the following would they use for the below requirement?

"Find out if a network security rule is preventing a network packet from reaching a virtual machine hosted in an Azure virtual network"

- A. IP Flow Verify 
- B. Next Hop
- C. Packet Capture
- D. Traffic Analysis

Explanation:

Answer – A

This can be done with the IP Flow Verify feature. The Microsoft documentation mentions the following

Introduction to IP flow verify in Azure Network Watcher

11/30/2017 • 2 minutes to read • Contributors 

IP flow verify checks if a packet is allowed or denied to or from a virtual machine. The information consists of direction, protocol, local IP, remote IP, local port, and remote port. If the packet is denied by a security group, the name of the rule that denied the packet is returned. While any source or destination IP can be chosen, IP flow verify helps administrators quickly diagnose connectivity issues from or to the internet and from or to the on-premises environment.

Option B is incorrect since this feature is used to get the next hop type and IP address of a packet from a specific VM

Option C is incorrect since this feature is used for deep dive network packet capture

Option D is incorrect since this feature is a cloud-based solution that provides visibility into user and application activity in cloud networks

For more information on the IP Flow Verify feature, please visit the below URL

<https://docs.microsoft.com/en-us/azure/network-watcher/network-watcher-ip-flow-verify-overview>

Ask our Experts

Rate this Question?  

View Queries

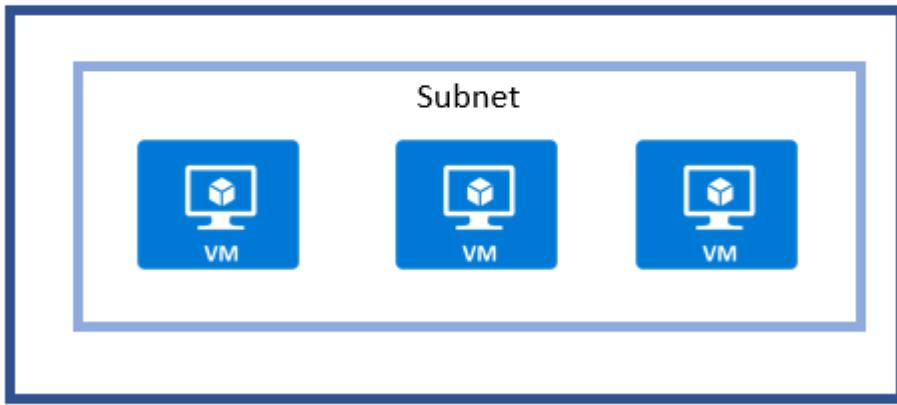
open ▾

Question 19

Unattempted

Domain :Design an infrastructure strategy

A company currently has the following deployment for an application in Azure



Network Security Groups are applied at both a subnet level and at the network interface level. The Networking team wants to use the Network Watcher tool to diagnose various issues when it comes to the networking aspect.

Which of the following would they use for the below requirement?

"Find out if there is outbound connectivity between an Azure virtual machine and an external host"

- A. IP Flow Verify
- B. Next Hop
- C. Connection Monitor
- D. Traffic Analytics

Explanation:

Answer – C

This can be done with the Connection Monitor feature. The Microsoft documentation mentions the following

Monitor communication between a virtual machine and an endpoint [🔗](#)

Endpoints can be another virtual machine (VM), a fully qualified domain name (FQDN), a uniform resource identifier (URI), or IPv4 address. The *connection monitor* capability monitors communication at a regular interval and informs you of reachability, latency, and network topology changes between the VM and the endpoint. For example, you might have a web server VM that communicates with a database server VM. Someone in your organization may, unknown to you, apply a custom route or network security rule to the web server or database server VM or subnet.

If an endpoint becomes unreachable, connection troubleshoot informs you of the reason. Potential reasons are a DNS name resolution problem, the CPU, memory, or firewall within the operating system of a VM, or the hop type of a custom route, or security rule for the VM or subnet of the outbound connection. Learn more about [security rules](#) and [route hop types](#) in Azure.

Option A is incorrect since this feature is used to verify the flow of traffic based on security group rules

Option B is incorrect since this feature is used to get the next hop type and IP address of a packet from a specific VM

Option D is incorrect since this feature is a cloud-based solution that provides visibility into user and application activity in cloud networks

For more information on the network watcher tool, please visit the below URL

<https://docs.microsoft.com/en-us/azure/network-watcher/network-watcher-ip-flow-verify-overview>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 20

Unattempted

Domain :Design an infrastructure strategy

A team wants to deploy a set of Linux based virtual machines onto an Azure subscription. They want to ensure that a set of packages get automatically installed when the machines are provisioned in Azure. Which of the following could they use for this purpose?

- A. Azure Policies
- B. Cloud-init 
- C. Scale sets
- D. Function Apps

Explanation:

Answer - B

The Cloud-init feature can be used to install packages on Linux based virtual machines. The Microsoft documentation mentions the following

Cloud-init

[Cloud-init](#) is a widely used approach to customize a Linux VM as it boots for the first time. You can use cloud-init to install packages and write files, or to configure users and security. Because cloud-init is called during the initial boot process, there are no additional steps or required agents to apply your configuration. For more information on how to properly format your `#cloud-config` files, see the [cloud-init documentation site](#). `#cloud-config` files are text files encoded in base64.

Cloud-init also works across distributions. For example, you don't use **apt-get install** or **yum install** to install a package. Instead you can define a list of packages to install. Cloud-init automatically uses the native package management tool for the distro you select.

Since this is clear from the Microsoft documentation, all other options are incorrect

For more information on Infrastructure automation, please visit the below URL

<https://docs.microsoft.com/en-us/azure/virtual-machines/windows/infrastructure-automation>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 21

Unattempted

Domain :Design a business continuity strategy

A company wants to create a set of virtual machines in Azure. These would be Windows virtual machines that would be hosting a web-based application. The company wants to achieve an availability of 99.99% for the virtual machines.

They decide to implement availability sets for the virtual machines.

Would this satisfy the requirement?

A. Yes

B. No 

Explanation:

Answer – B

Using Availability sets would help achieve an availability of 99.95% and not 99.99%.

For more information on availability sets, please visit the below URL

<https://docs.microsoft.com/en-us/azure/virtual-machines/windows/manage-availability#use-availability-zones-to-protect-from-datacenter-level-failures>

[View Queries](#)

open ▾

Question 22**Unattempted****Domain :Design a business continuity strategy**

A company wants to create a set of virtual machines in Azure. These would be Windows virtual machines that would be hosting a web-based application. The company wants to achieve an availability of 99.99% for the virtual machines.

They decide to implement availability zones for the virtual machines.

Would this satisfy the requirement?

- A. Yes
- B. No

Explanation:

Answer – A

Yes, by deploying the virtual machines across availability zones, you can achieve an availability of 99.99%

The Microsoft documentation mentions the following

What are Availability Zones in Azure?

06/20/2019 • 3 minutes to read • +17

Availability Zones is a high-availability offering that protects your applications and data from datacenter failures. Availability Zones are unique physical locations within an Azure region. Each zone is made up of one or more datacenters equipped with independent power, cooling, and networking. To ensure resiliency, there's a minimum of three separate zones in all enabled regions. The physical separation of Availability Zones within a region protects applications and data from datacenter failures. Zone-redundant services replicate your applications and data across Availability Zones to protect from single-points-of-failure. With Availability Zones, Azure offers industry best 99.99% VM uptime SLA. The full [Azure SLA](#) explains the guaranteed availability of Azure as a whole.

For more information on availability zones, please visit the below URL

<https://docs.microsoft.com/en-us/azure/availability-zones/az-overview>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 23

Unattempted

Domain :Design a business continuity strategy

A company wants to create a set of virtual machines in Azure. These would be Windows virtual machines that would be hosting a web-based application. The company wants to achieve an availability of 99.99% for the virtual machines.

They decide to implement Azure Load balancers for the virtual machines.

Would this satisfy the requirement?

A. Yes

B. No 

Explanation:

Answer – B

Azure Load balancers are primarily used to distribute traffic across virtual machines. They can add high availability to your architecture but does not guarantee an availability of 99.99% for the virtual machines.

For more information on Azure Load balancers, please visit the below URL

<https://docs.microsoft.com/en-us/azure/load-balancer/load-balancer-overview>

Note:

A standard load balancer as a service has 99.99% availability, this provides the same availability to the application, but NOT the virtual machines behind it.

https://azure.microsoft.com/en-ca/support/legal/sla/load-balancer/v1_0/

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 24

Unattempted

Domain :Determine workload requirements

A company has setup an Azure subscription. They also have an on-premise data center. They have both Windows and Linux virtual machines located in their on-premise datacenter and on Azure. They wanted to

design their logging solution in such a way

- **Requirement1** - Collect IIS logs from the Windows virtual machines located in Azure and send them across to Azure Storage
- **Requirement2** -Collect metrics from the on-premise Windows virtual machines
- **Requirement3** -Collect metrics from the on-premise Linux virtual machines

Which of the following could be used for Requirement1?

- A. Log Analytics Agent
- B. Azure Diagnostics Agent
- C. Network Watcher
- D. Application Insights

Explanation:

Answer – B

This can be accomplished with the help of the Azure Diagnostics Agent

The Microsoft documentation mentions the following

Azure Diagnostic extension

The [Azure Diagnostics extension](#) (commonly referred to as the Windows Azure Diagnostic (WAD) or Linux Azure Diagnostic (LAD) extension), which has been provided for Azure Cloud Services since it became generally available in 2010, is an agent that delivers simple collection of diagnostic data from an Azure compute resource like a VM, and persist it to Azure storage. Once in storage, you choose to view with one of several available tools, such as [Server Explorer in Visual Studio](#) and [Azure Storage Explorer](#).

Option A is incorrect because this is used to send logs to a Log Analytics Workspace

Option C is incorrect because this is used to diagnose network issues

Option D is incorrect because this used to debug application related issues

For more information on Azure Diagnostics agent, please visit the below URL

<https://docs.microsoft.com/en-us/azure/azure-monitor/platform/agents-overview>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Domain :Determine workload requirements

A company has setup an Azure subscription. They also have an on-premise data center. They have both Windows and Linux virtual machines located in their on-premise datacenter and on Azure. They wanted to design their logging solution in such a way

- **Requirement1** - Collect IIS logs from the Windows virtual machines located in Azure and send them across to Azure Storage
- **Requirement2** -Collect metrics from the on-premise Windows virtual machines
- **Requirement3** -Collect metrics from the on-premise Linux virtual machines

Which of the following could be used for Requirement2?

- A. Log Analytics Agent ✓
- B. Azure Diagnostics Agent
- C. Network Watcher
- D. Application Insights

Explanation:

Answer – A

You can use the Log Analytics agent to collect data from an on-premise machine.

The Microsoft documentation mentions the following

Collect log data with the Log Analytics agent

07/01/2019 • 7 minutes to read • 

The Azure Log Analytics agent, previously referred to as the Microsoft Monitoring Agent (MMA) or OMS Linux agent, was developed for comprehensive management across on-premises machines, computers monitored by [System Center Operations Manager](#), and virtual machines in any cloud. The Windows and Linux agents attach to an Azure Monitor and store collected log data from different sources in your Log Analytics workspace, as well as any unique logs or metrics as defined in a monitoring solution.

Option B is incorrect because this is used for Azure based virtual machines

Option C is incorrect because this is used to diagnose network issues

Option D is incorrect because this used to debug application related issues

For more information on Azure Diagnostics agent, please visit the below URL

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 26

Unattempted

Domain :Determine workload requirements

A company has setup an Azure subscription. They also have an on-premise data center. They have both Windows and Linux virtual machines located in their on-premise datacenter and on Azure. They wanted to design their logging solution in such a way

- **Requirement1** - Collect IIS logs from the Windows virtual machines located in Azure and send them across to Azure Storage
- **Requirement2** -Collect metrics from the on-premise Windows virtual machines
- **Requirement3** -Collect metrics from the on-premise Linux virtual machines

Which of the following could be used for Requirement3?

- A. Log Analytics Agent 
- B. Azure Diagnostics Agent
- C. Network Watcher
- D. Application Insights

Explanation:

Answer – A

You can use the Log Analytics agent to collect data from an on-premise machine.

The Microsoft documentation mentions the following

Collect log data with the Log Analytics agent

07/01/2019 • 7 minutes to read • 

The Azure Log Analytics agent, previously referred to as the Microsoft Monitoring Agent (MMA) or OMS Linux agent, was developed for comprehensive management across on-premises machines, computers monitored by [System Center Operations Manager](#), and virtual machines in any cloud. The Windows and Linux agents attach to an Azure Monitor and store collected log data from different sources in your Log Analytics workspace, as well as any unique logs or metrics as defined in a monitoring solution.

Option B is incorrect because this is used for Azure based virtual machines

Option C is incorrect because this is used to diagnose network issues

Option D is incorrect because this used to debug application related issues

For more information on Azure Diagnostics agent, please visit the below URL

<https://docs.microsoft.com/en-us/azure/azure-monitor/platform/agents-overview>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 27

Unattempted

Domain :Design a data platform solution

You have to manage the backup of several Azure SQL databases in your company's Azure subscription. You need to configure long term retention for backup's as well. The Azure subscription has a collection of single, pooled databases and Azure Managed Instances as well.

For which of the following can you configure long term retention? Choose 2 answers from the options given below

- A. Azure SQL database – Single instance 
- B. Azure SQL database – Pooled instance 
- C. Azure SQL Managed Instance
- D. Azure SQL datawarehouse

Explanation:

Answer – A and B

Long term retention can be configured for Azure SQL database – Single and pooled instances.

The Microsoft documentation mentions the following

Store Azure SQL Database backups for up to 10 years

05/18/2019 • 4 minutes to read • 5 comments +3

Many applications have regulatory, compliance, or other business purposes that require you to retain database backups beyond the 7-35 days provided by Azure SQL Database [automatic backups](#). By using the long-term retention (LTR) feature, you can store specified SQL database full backups in [RA-GRS](#) blob storage for up to 10 years. You can then restore any backup as a new database.

Note



LTR can be enabled for single and pooled databases. It is not yet available for instance databases in Managed Instances. You can use SQL Agent jobs to schedule [copy-only database backups](#) as an alternative to LTR beyond 35 days.

Since this is clearly mentioned in the Microsoft documentation, all other options are incorrect

For more information on long term retention, please visit the below URL

<https://docs.microsoft.com/en-us/azure/sql-database/sql-database-long-term-retention>

Ask our Experts

Rate this Question?

View Queries

open ▾

Question 28

Unattempted

Domain :Design a data platform solution

You have to manage the backup of several Azure SQL databases in your company's Azure subscription. You need to configure long term retention for backup's as well. The Azure subscription has a collection of single, pooled databases and Azure Managed Instances as well.

Which of the following can be used to create long term backup's for Azure SQL Managed Instances?

A. Az Copy

B. Data Explorer

C. SQL Agent jobs ✓

D. Data Factory

Explanation:

Answer – C

You can achieve this with SQL agent jobs.

The Microsoft documentation mentions the following

Store Azure SQL Database backups for up to 10 years

05/18/2019 • 4 minutes to read • 3 comments +3

Many applications have regulatory, compliance, or other business purposes that require you to retain database backups beyond the 7-35 days provided by Azure SQL Database [automatic backups](#). By using the long-term retention (LTR) feature, you can store specified SQL database full backups in [RA-GRS](#) blob storage for up to 10 years. You can then restore any backup as a new database.

① Note



LTR can be enabled for single and pooled databases. It is not yet available for instance databases in Managed Instances. You can use SQL Agent jobs to schedule [copy-only database backups](#) as an alternative to LTR beyond 35 days.

Since this is clearly mentioned in the Microsoft documentation, all other options are incorrect

For more information on long term retention, please visit the below URL

<https://docs.microsoft.com/en-us/azure/sql-database/sql-database-long-term-retention>

Ask our Experts

Rate this Question?

View Queries

open ▾

Question 29

Unattempted

Domain :Design a business continuity strategy

A company needs to setup a web application on to the Azure Web App Service. One of the key requirements is for the platform to scale automatically when there is sustainable load that caused the CPU threshold to go

beyond a particular threshold.

Which of the following is the most cost-effective App Service Plan to use for the Azure Web App?

- A. Basic
- B. Standard
- C. Premium
- D. Isolated

Explanation:

Answer - B

The Autoscaling option for Web Apps is available from the Standard App Service plan onwards. Hence this would be the most cost-effective plan.

The Microsoft documentation mentions the following

	FREE	SHARED	BASIC	STANDARD	PREMIUM	ISOLATED *	APP SERVICE LINUX	CONSUMPTION PLAN (FUNCTIONS)
Java	✓	✓	✓	✓	✓	✓		alpha
Node.js	✓	✓	✓	✓	✓	✓	✓	✓
PHP	✓	✓	✓	✓	✓	✓	✓	alpha
Python	✓	✓	✓	✓	✓	✓		alpha
Ruby							✓	
– Scale								
Auto-scale				✓	✓	✓	✓	✓
Integrated Load Balancer	✓	✓	✓	✓	✓	✓	✓	✓
Traffic Manager				✓	✓	✓		✓

Since this is clearly mentioned in the Microsoft documentation, all other options are incorrect

For more information on Azure App Service Plans, please visit the below URL

<https://azure.microsoft.com/en-us/pricing/details/app-service/plans/>

View Queries

open ▾

Question 30**Unattempted****Domain :Design a business continuity strategy**

A company needs to setup a web application on to the Azure Web App Service. One of the key requirements is for the platform to scale automatically when there is sustainable load that caused the CPU threshold to go beyond a particular threshold.

The IT Admin team needs to be notified if there is an autoscale failure. Which of the following can be done to alert the team accordingly?

- A. Create an Analytics Log Alert
- B. Create an Activity Log Alert 
- C. Create a metrics Log Alert
- D. Create a Threshold Alert

Explanation:

Answer – B

You can create an Activity Log Alert for this.

The Microsoft documentation mentions the following

Autoscale concepts

- A resource can have only *one* autoscale setting
- An autoscale setting can have one or more profiles and each profile can have one or more autoscale rules.
- An autoscale setting scales instances horizontally, which is *out* by increasing the instances and *in* by decreasing the number of instances. An autoscale setting has a maximum, minimum, and default value of instances.
- An autoscale job always reads the associated metric to scale by, checking if it has crossed the configured threshold for scale-out or scale-in. You can view a list of metrics that autoscale can scale by at [Azure Monitor autoscaling common metrics](#).
- All thresholds are calculated at an instance level. For example, "scale out by one instance when average CPU > 80% when instance count is 2", means scale-out when the average CPU across all instances is greater than 80%.
- All autoscale failures are logged to the Activity Log. You can then configure an activity log alert so that you can be notified via email, SMS, or webhooks whenever there is an autoscale failure.
- Similarly, all successful scale actions are posted to the Activity Log. You can then configure an activity log alert so that you can be notified via email, SMS, or webhooks whenever there is a successful autoscale action. You can also configure email or webhook notifications to get notified for successful scale actions via the notifications tab on the autoscale setting.

Since this is clearly mentioned in the Microsoft documentation, all other options are incorrect

- For more information on best practices for autoscaling, please visit the below URL
 - <https://docs.microsoft.com/en-us/azure/azure-monitor/platform/autoscale-best-practices>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 31

Unattempted

Domain :Design for identity and security

A company has setup an Azure subscription and deployed a set of Virtual Machines to the subscription. They want to enable replication for Azure Virtual Machine and for machines located in their on-premise network. The company has defined the following groups in their Azure AD tenant along with the required relevant permission description

Name

Permissions required

whizlabgrpA

Create recovery service vault

whizlabgrpB

Create and manage replication jobs

whizlabgrpC

Perform failovers

As an architect you need to assign the right role to each group. You have to ensure that you don't assign more permissions than what is required.

Which role should be given to the group whizlabgrpA?

- A. Owner
- B. Contributor
- C. Site Recovery Contributor
- D. Site Recovery Operator

Explanation:

Answer – B

Both the Site Recovery Contributor and the Operator don't have the privilege to create a recovery services vault, so you will need to choose the Contributor role.

Contributor

Description	Lets you manage everything except access to resources.
Id	b24988ac-6180-42a0-ab88-20f7382dd24c
Actions	 Create and manage resources of all types
*	

Option A is incorrect since this role would give elevated permissions which are not required

For more information on in-built roles, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/role-based-access-control/built-in-roles>

View Queries

open ▾

Question 32

Unattempted

Domain :Design for identity and security

A company has setup an Azure subscription and deployed a set of Virtual Machines to the subscription. They want to enable replication for Azure Virtual Machine and for machines located in their on-premise network. The company has defined the following groups in their Azure AD tenant along with the required relevant permission description

Name	Permissions required
whizlabgrpA	Create recovery service vault
whizlabgrpB	Create and manage replication jobs
whizlabgrpC	Perform failovers

As an architect you need to assign the right role to each group. You have to ensure that you don't assign more permissions than what is required.

Which role should be given to the group whizlabgrpB?

- A. Owner
- B. Contributor
- C. Site Recovery Contributor 
- D. Site Recovery Operator

Explanation:

Answer – C

The Site Recovery Contributor has the required privilege for this. The permissions for the role are mentioned in the Microsoft documentation.

Microsoft.RecoveryServices/vaults/replicationFabrics/*	Create and manage replication fabrics	
Microsoft.RecoveryServices/vaults/replicationJobs/*	Create and manage replication jobs	1
Microsoft.RecoveryServices/vaults/replicationPolicies/*	Create and manage replication policies	2
Microsoft.RecoveryServices/vaults/replicationRecoveryPlans/*	Create and manage recovery plans	
Microsoft.RecoveryServices/Vaults/storageConfig/*	Create and manage storage configuration of Recovery Services vault	

Options A and B are incorrect since these roles would give elevated permissions which are not required

Option D is incorrect since this role does not have the required privilege

For more information on in-built roles, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/role-based-access-control/built-in-roles>

Ask our Experts

Rate this Question?

View Queries

open ▾

Question 33

Unattempted

Domain :Design for identity and security

A company has setup an Azure subscription and deployed a set of Virtual Machines to the subscription. They want to enable replication for Azure Virtual Machine and for machines located in their on-premise network. The company has defined the following groups in their Azure AD tenant along with the required relevant permission description

Name	Permissions required
whizlabgrpA	Create recovery service vault
whizlabgrpB	Create and manage replication jobs
whizlabgrpC	Perform failovers

As an architect you need to assign the right role to each group. You have to ensure that you don't assign more permissions than what is required.

Which role should be given to the group whizlabgrpC?

Resource Policy Contributor (Preview)
Scheduler Job Collections Contributor
Search Service Contributor
Security Admin
Security Manager (Legacy)
Security Reader
Site Recovery Contributor
Site Recovery Operator

- A. Owner
- B. Contributor
- C. Site Recovery Contributor
- D. Site Recovery Operator ✓

Explanation:

Answer – D

The Site Recovery operator has the required privilege for this. The permissions for the role are mentioned in the Microsoft documentation.

Microsoft.RecoveryServices/vaults/replicationFabrics/replicationProtectio nContainers/replicationProtectedItems/unplannedFailover/action	Failover	2
Microsoft.RecoveryServices/vaults/replicationFabrics/replicationProtectio nContainers/replicationProtectedItems/updateMobilityService/action	Update Mobility Service	1
Microsoft.RecoveryServices/vaults/replicationFabrics/replicationProtectio nContainers/replicationProtectionContainerMappings/read	Read any Protection Container Mappings	

Security Admin
 Security Manager (Legacy)
 Security Reader
 Site Recovery Contributor
Site Recovery Operator
 Site Recovery Reader
 SQL DB Contributor
 SQL Security Manager

All other roles would give elevated permissions which are not required.

For more information on in-built roles, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/role-based-access-control/built-in-roles>

Ask our Experts

Rate this Question?

View Queries

open ▾

Question 34

Unattempted

Domain :Determine workload requirements

A company needs to ensure that a network security group is defined for every subnet in all virtual networks defined in their Azure subscription. Which of the following could be used to achieve this requirement?

- A. Azure Advisor
- B. Azure Policies ✓
- C. Azure RBAC
- D. Azure AD

Explanation:

Answer – B

An example of this is given in the Microsoft documentation

Sample - Network security group x on every subnet

01/23/2019 • 2 minutes to read • 

This policy requires that a specific network security group is used with every virtual subnet. You specify the ID of the network security group to use.

If you don't have an [Azure subscription](#), create a [free account](#) before you begin.

Option A is incorrect because this is basically a recommendations engine

Option C is incorrect because this is used to set access to resources

Option D is incorrect because this is an identity provider in Azure

For more information on the example itself, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/governance/policy/samples/nsg-on-subnet>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 35

Unattempted

Domain :Design an infrastructure strategy

A company wants to setup the following architecture in Azure

- A web application hosted on a set of two virtual machines
- Users would be accessing the web application via the URL <https://whizlabs.com>
- A database hosted on a set of two virtual machines
- The database would be listening for connections on port 1433
- A Public facing load balancer for the web application layer
- An Internal load balancer for the database layer

How many frontend IP addresses need to be assigned to the public load balancer?

- A. 0
- B. 1
- C. 2
- D. 4

Explanation:

Answer – B

Since the users would be accessing the application via a secure URL, that means the load balancer needs to be listening on port 443. Hence here we just need one frontend IP configuration.

Since this is the ideal approach, all other options are incorrect

For more information on the Azure Load balancer, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/load-balancer/load-balancer-overview>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 36

Unattempted

Domain :Design an infrastructure strategy

A company wants to setup the following architecture in Azure

- A web application hosted on a set of two virtual machines
- Users would be accessing the web application via the URL <https://www.whizlabs.com>
- A database hosted on a set of two virtual machines
- The database would be listening for connections on port 1433
- A Public facing load balancer for the web application layer
- An Internal load balancer for the database layer

How many IP addresses need to be assigned to the Internal load balancer?

- A. 0
- B. 1

C. 2

D. 4

Explanation:

Answer – B

The question here is referring to IP address that needs to be assigned to the Internal Load Balancer and when we work with it, we need to assign an Private IP address, and this address should be with-in the range of your virtual network and subnet.

I request you to refer the following link to know further.

<https://docs.microsoft.com/en-us/azure/load-balancer/tutorial-load-balancer-basic-internal-portal>

<https://docs.microsoft.com/en-us/azure/virtual-network/virtual-network-ip-addresses-overview-arm#private-ip-addresses>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 37

Unattempted

Domain :Design an infrastructure strategy

A company needs to design and deploy an application onto Azure. They want to base the design the application on a microservice design pattern. For deployment, they want to have a microservices based platform on Azure for hosting their application. The application itself would be a ecommerce application with modules for managing user accounts, shopping carts, managing delivery.

Which of the following could they use for hosting the application?

- A. Azure Functions
- B. Azure Service Fabric 
- C. Azure Web Apps
- D. Azure Logic Apps

Explanation:

Answer – B

The ideal solution for this is Azure Service Fabric. The Microsoft documentation mentions the following

Service Fabric as a microservices platform

Azure Service Fabric emerged when Microsoft transitioned from delivering boxed products, which were typically monolithic, to delivering services. The experience of building and operating large services, like Azure SQL Database and Azure Cosmos DB, shaped Service Fabric. The platform evolved over time as more services adopted it. Service Fabric had to run not only in Azure but also in standalone Windows Server deployments.

The aim of Service Fabric is to solve the hard problems of building and running a service and to use infrastructure resources efficiently, so teams can solve business problems by using a microservices approach.

Service Fabric helps you build applications that use a microservices approach by providing:

- A platform that provides system services to deploy, upgrade, detect, and restart failed services, discover services, route messages, manage state, and monitor health.
- The ability to deploy applications either running in containers or as processes. Service Fabric is a container and process orchestrator.
- Productive programming APIs to help you build applications as microservices: [ASP.NET Core](#), [Reliable Actors](#), and [Reliable Services](#). For example, you can get health and diagnostics information, or you can take advantage of built-in high availability.

Option A is incorrect because this is a serverless platform for hosting code

Option C is incorrect because this is primarily used as a platform for hosting web applications

Option D is incorrect because this is a workflow-based service

For more information on Azure Service Fabric, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/service-fabric/service-fabric-overview-microservices>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 38

Unattempted

Domain :Design an infrastructure strategy

A company needs to design and deploy an application onto Azure. They want to base the design on the application on a microservice design pattern. For deployment, they want to have a microservices based platform on Azure for hosting their application. The application itself would be an ecommerce application with modules for managing user accounts, shopping carts, managing delivery.

Based on the application need, what type of design pattern would you use to build the application?

- A. Stateless Microservices
- B. Stateful Microservices 
- C. Function Microservices
- D. Module Microservices

Explanation:

Answer – B

Since the application needs to maintain the state across its lifecycle, (for example from shopping cart items to checkout), you can use the Stateful microservices design pattern.

The Microsoft documentation mentions the following

Stateless and stateful microservices for Service Fabric

Service Fabric enables you to build applications that consist of microservices or containers. Stateless microservices (such as protocol gateways and web proxies) do not maintain a mutable state outside a request and its response from the service. Azure Cloud Services worker roles are an example of a stateless service. Stateful microservices (such as user accounts, databases, devices, shopping carts, and queues) maintain a mutable, authoritative state beyond the request and its response. Today's Internet-scale applications consist of a combination of stateless and stateful microservices.

Since this is clearly mentioned in the documentation, all other options are incorrect

For more information on Azure Service Fabric, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/service-fabric/service-fabric-content-roadmap>

Ask our Experts

Rate this Question?  

View Queries

open 

Question 39

Unattempted

Domain :Design for deployment, migration, and integration

A company needs to transfer a data set to Azure. The size of the data set is around 10TB. The data needs to be copied to an Azure Storage account. The currently line bandwidth that the company has is around 45 Mbps. Which of the following would be the ideal solution to use for the data transfer?

- A. Az copy tool

- B. Azure Data Factory
- C. Azure Data Box
- D. Azure Powershell

Explanation:

Answer – C

Here the size of the data set is large, and the bandwidth is small. Hence doing an offline transfer with the Azure Data Box service is the ideal approach.

The Microsoft documentation also gives a recommendation on when to use the service

Offline transfer or network transfer

Large datasets imply that you have few TBs to few PBs of data. You have limited to no network bandwidth, your network is slow, or it is unreliable. Also:

- You are limited by costs of network transfer from your Internet Service Providers (ISPs).
- Security or organizational policies do not allow outbound connections when dealing with sensitive data.

In all the above instances, use a physical device to do a one-time bulk data transfer. Choose from Data Box Disk, Data Box, Data Box Heavy devices which are supplied by Microsoft, or Import/Export using your own disks.

To confirm whether a physical device is the right option, use the following table. It shows the projected time for network data transfer, for various available bandwidths (assuming 90% utilization). If network transfer is projected to be too slow, you should use a physical device.

Data size ↓ Network bandwidth →	45 Mbps (T3)	100 Mbps	1 Gbps	10 Gbps
1 TB	2 days	1 day	3 hours	15 minutes
10 TB	23 days	11 days	1 day	3 hours
35 TB	82 days	37 days	4 days	9 hours
80 TB	187 days	84 days	8 days	20 hours
100 TB	234 days	105 days	11 days	1 day
200 TB	1 year	211 days	21 days	2 days
500 TB	3 years	1 year	53 days	5 days
1 PB	7 years	3 years	108 days	11 days
2 PB	13 years	6 years	216 days	22 days
5 PB	33 years	15 years	1 year	54 days

Key:
Use a Data Box Disk instead
Use a Data Box instead
Use a Data Box Heavy Instead
Use the network

The other options are incorrect since an online transfer would just take too long.

For more information on transferring large data sets over a limited bandwidth network, please go ahead and visit the below URL

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 40

Unattempted

Domain :Design for deployment, migration, and integration

A company needs to transfer a data set to Azure. The size of the data set is around 10TB. The data needs to be copied to an Azure Storage account. The currently line bandwidth that the company has is around 100 MBps. After the initial data transfer, which of the following can be used for periodic transfers of MB's of data onto the Azure storage account? Choose 2 answers from the options given below

- A. Az Copy tool 
- B. Azure Data Factory 
- C. Azure EventHubs
- D. SQL Elastic Pools

Explanation:

Answer – A and B

The Microsoft documentation also mentions options on periodic data transfers

Recommended options

The recommended options for periodic data transfer fall into two categories depending on whether the transfer is recurring or continuous.

- **Scripted/programmatic tools** – For data transfer that occurs at regular intervals, use the scripted and programmatic tools such as AzCopy and Azure Storage REST APIs. These tools are targeted towards IT professionals and developers.
 - **AzCopy** - Use this command-line tool to easily copy data to and from Azure Blobs, Files, and Table storage with optimal performance. AzCopy supports concurrency and parallelism, and the ability to resume copy operations when interrupted.
 - **Azure Storage REST APIs/SDKs** – When building an application, you can develop the application against Azure Storage REST APIs and use the Azure SDKs offered in multiple languages. The REST APIs can also leverage the Azure Storage Data Movement Library designed especially for the high-performance copying of data to and from Azure.
- **Continuous data ingestion tools** – For continuous, ongoing data ingestion, you can select one of Data Box online transfer device or Azure Data Factory. These tools are set up by IT professionals and can transparently automate data transfer.
 - **Azure Data Factory** – Data Factory should be used to scale out a transfer operation, and if there is a need for orchestration and enterprise grade monitoring capabilities. Use Azure Data Factory to set up a cloud pipeline that regularly transfers files between several Azure services, on-premises, or a combination of the two. Azure Data Factory lets you orchestrate data-driven workflows that ingest data from disparate data stores and automate data movement and data transformation.
 - **Azure Data Box family for online transfers** - Data Box Edge and Data Box Gateway are online network devices that can move data into and out of Azure. Data Box Edge uses artificial intelligence (AI)-enabled Edge compute to pre-process data before upload. Data Box Gateway is a virtual version of the device with the same data transfer capabilities.

Since all the options are clearly mentioned in the documentation, all other options are incorrect

For more information on solutions for periodic data transfers, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/storage/common/storage-solution-periodic-data-transfer>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 41

Unattempted

Domain :Determine workload requirements

A company is planning on setting the following architecture in Azure.

- A web application hosted using the Azure Web App service

- The application would be written in .Net core

- A CosmosDB account with the Table API

The company has the following key requirements

- Ensure proper capacity provisioning is implemented for the CosmosDB account
- Generate reports on which parts of the application are most popular amongst users

Which of the following could be used to accomplish the below requirement?

"Ensure proper capacity provisioning is implemented for the CosmosDB account"

- A. Generate alerts based on the service health of the resource
- B. Generate alerts based on CosmosDB metrics
- C. Generate alerts based on Azure Advisor metrics
- D. Generate alerts based on Azure Activity Log Metrics

Explanation:

Answer – B

You can see how the throughput is being managed via the metrics available on the CosmosDB account. You can create an alert based on the metric itself.

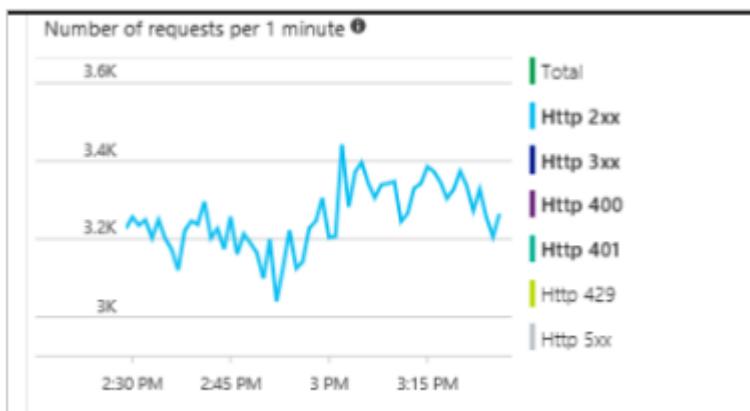
The Microsoft documentation mentions the following on the metrics

Understand how many requests are succeeding or causing errors

To get started, head to the [Azure portal](#) and navigate to the **Metrics** blade. In the blade, find the

**Number of requests exceeded capacity per 1-minute chart. This chart shows a minute by minute total requests segmented by the status code. For more information about HTTP status codes, see [HTTP status codes for Azure Cosmos DB](#).

The most common error status code is 429 (rate limiting/throttling). This error means that requests to Azure Cosmos DB are more than the provisioned throughput. The most common solution to this problem is to [scale up the RUs](#) for the given collection.



Since this is clearly given in the Microsoft documentation, all other options are incorrect

For more information on using metrics for CosmosDB, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/cosmos-db/use-metrics>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 42

Unattempted

Domain :Determine workload requirements

A company is planning on setting the following architecture in Azure.

- A web application hosted using the Azure Web App service
- The application would be written in .Net core
- A CosmosDB account with the Table API

The company has the following key requirements

- Ensure proper capacity provisioning is implemented for the CosmosDB account
- Generate reports on which parts of the application are most popular amongst users

Which of the following could be used to accomplish the below requirement?

"Generate reports on which parts of the application are most popular amongst users"

- A. Use Azure Advisor
- B. Use Application Insights 
- C. Use Azure Monitor
- D. Use Azure metrics

Explanation:

Answer – B

Application Insights can provide the require usage analysis for your web-based application

The Microsoft documentation mentions the following

Usage analysis with Application Insights

10/10/2017 • 5 minutes to read • 

Which features of your web or mobile app are most popular? Do your users achieve their goals with your app? Do they drop out at particular points, and do they return later? [Azure Application Insights](#) helps you gain powerful insights into how people use your app. Every time you update your app, you can assess how well it works for users. With this knowledge, you can make data driven decisions about your next development cycles.

Since this is clearly given in the Microsoft documentation, all other options are incorrect

For more information on usage analysis in Application Insights, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/azure-monitor/app/usage-overview>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 43

Unattempted

Domain :Design for identity and security

A company has setup an Azure subscription and account. They want to provision the following resources in Azure

- A virtual network with a set of virtual machines
- A Cosmos DB account
- An Azure Web App instance

The company wants to set the right roles for different groups to manage the Cosmos DB account. Below are the required permission requirements

Group	Permission
	<p>cosmosdb_grp1</p> <p>Should be able to provision Cosmos DB accounts, but can't access the keys required to access the data</p>

cosmosdb_grp2

Should be able to read the data in a
Cosmos DB account

cosmosdb_grp3

Should be able to submit a request to
restore a Cosmos DB database

You have to advise on the right role-based access to provide to each group based on least privilege.
Which of the following would you assign to cosmosdb_grp1?

- A. DocumentDB Accounts Contributor
- B. Cosmos DB Account Reader
- C. Cosmos DB Backup Operator
- D. Cosmos DB Operator

Explanation:

Answer – D

The right role for this is the Cosmos DB Operator

The Microsoft documentation mentions the following

Role-based access control in Azure Cosmos DB

05/23/2019 • 2 minutes to read • 

Azure Cosmos DB provides built-in role-based access control (RBAC) for common management scenarios in Azure Cosmos DB. An individual who has a profile in Azure Active Directory can assign these RBAC roles to users, groups, service principals, or managed identities to grant or deny access to resources and operations on Azure Cosmos DB resources. Role assignments are scoped to control-plane access only, which includes access to Azure Cosmos accounts, databases, containers, and offers (throughput).

Built-in roles

The following are the built-in roles supported by Azure Cosmos DB:

Built-in role	Description
DocumentDB Accounts Contributor	Can manage Azure Cosmos DB accounts.
Cosmos DB Account Reader	Can read Azure Cosmos DB account data.
Cosmos Backup Operator	Can submit restore request for an Azure Cosmos database or a container.
Cosmos DB Operator	Can provision Azure Cosmos accounts, databases, and containers but cannot access the keys that are required to access the data.

Since this is clearly given in the Microsoft documentation, all other options are incorrect

For more information on role-based access for CosmosDB, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/cosmos-db/role-based-access-control>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 44

Unattempted

Domain :Design for identity and security

A company has setup an Azure subscription and account. They want to provision the following resources in Azure

- A virtual network with a set of virtual machines
- A Cosmos DB account

- An Azure Web App instance

The company wants to set the right roles for different groups to manage the Cosmos DB account. Below are the required permission requirements

Group	Permission
cosmosdb_grp1	Should be able to provision Cosmos DB accounts, but can't access the keys required to access the data
cosmosdb_grp2	Should be able to read the data in a Cosmos DB account
cosmosdb_grp3	Should be able to submit a request to restore a Cosmos DB database

You have to advise on the right role-based access to provide to each group based on least privilege. Which of the following would you assign to cosmosdb_grp2?

- A. DocumentDB Accounts Contributor
- B. Cosmos DB Account Reader
- C. Cosmos DB Backup Operator
- D. Cosmos DB Operator

Explanation:

Answer – B

The right role for this is the Cosmos DB Account Reader

The Microsoft documentation mentions the following

Role-based access control in Azure Cosmos DB

05/23/2019 • 2 minutes to read • 

Azure Cosmos DB provides built-in role-based access control (RBAC) for common management scenarios in Azure Cosmos DB. An individual who has a profile in Azure Active Directory can assign these RBAC roles to users, groups, service principals, or managed identities to grant or deny access to resources and operations on Azure Cosmos DB resources. Role assignments are scoped to control-plane access only, which includes access to Azure Cosmos accounts, databases, containers, and offers (throughput).

Built-in roles

The following are the built-in roles supported by Azure Cosmos DB:

Built-in role	Description
DocumentDB Accounts Contributor	Can manage Azure Cosmos DB accounts.
Cosmos DB Account Reader	Can read Azure Cosmos DB account data.
Cosmos Backup Operator	Can submit restore request for an Azure Cosmos database or a container.
Cosmos DB Operator	Can provision Azure Cosmos accounts, databases, and containers but cannot access the keys that are required to access the data.

Since this is clearly given in the Microsoft documentation, all other options are incorrect

For more information on role-based access for CosmosDB, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/cosmos-db/role-based-access-control>

Ask our Experts

Rate this Question?  

View Queries

open 

Question 45

Unattempted

Domain :Design for identity and security

A company has setup an Azure subscription and account. They want to provision the following resources in Azure

- A virtual network with a set of virtual machines
- A Cosmos DB account

- An Azure Web App instance

The company wants to set the right roles for different groups to manage the Cosmos DB account. Below are the required permission requirements

Group	Permission
cosmosdb_grp1	Should be able to provision Cosmos DB accounts, but can't access the keys required to access the data
cosmosdb_grp2	Should be able to read the data in a Cosmos DB account
cosmosdb_grp3	Should be able to submit a request to restore a Cosmos DB database

You have to advise on the right role-based access to provide to each group based on least privilege. Which of the following would you assign to cosmosdb_grp3?

- A. DocumentDB Accounts Contributor
- B. Cosmos DB Account Reader
- C. Cosmos DB Backup Operator
- D. Cosmos DB Operator

Explanation:

Answer - C

The right role for this is the Cosmos DB Backup Operator

The Microsoft documentation mentions the following

Role-based access control in Azure Cosmos DB

05/23/2019 • 2 minutes to read • 

Azure Cosmos DB provides built-in role-based access control (RBAC) for common management scenarios in Azure Cosmos DB. An individual who has a profile in Azure Active Directory can assign these RBAC roles to users, groups, service principals, or managed identities to grant or deny access to resources and operations on Azure Cosmos DB resources. Role assignments are scoped to control-plane access only, which includes access to Azure Cosmos accounts, databases, containers, and offers (throughput).

Built-in roles

The following are the built-in roles supported by Azure Cosmos DB:

Built-in role	Description
DocumentDB Accounts Contributor	Can manage Azure Cosmos DB accounts.
Cosmos DB Account Reader	Can read Azure Cosmos DB account data.
Cosmos Backup Operator	Can submit restore request for an Azure Cosmos database or a container.
Cosmos DB Operator	Can provision Azure Cosmos accounts, databases, and containers but cannot access the keys that are required to access the data.

Since this is clearly given in the Microsoft documentation, all other options are incorrect

For more information on role-based access for CosmosDB, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/cosmos-db/role-based-access-control>

Ask our Experts

Rate this Question?  

View Queries

open 

Question 46

Unattempted

Domain :Design a business continuity strategy

[View Case Study](#)

The Web application tier needs to be deployed using a service that could automatically scale the web service based on demand. Which of the following would you recommend?

- A. Azure virtual machines with Internet Information Services Installed
- B. Azure Web Apps using the Standard App Service Plan 
- C. An Azure Standard Load Balancer
- D. An Azure Application Gateway with 2 medium instances

Explanation:

Answer – B

The best solution is to use Azure Web Apps as the managed service. With the Standard App Service plan, it would also have the ability to support auto-scaling.

Option A is incorrect since the case study mentions that we need to use Azure managed services wherever possible

Options C and D are incorrect since these services are used to distribute traffic

For more information on Autoscaling, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/architecture/best-practices/auto-scaling>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 47

Unattempted

Domain :Design a data platform solution

View Case Study

The company wants to store the backups taken for the SQL Server instance which would be hosted on an Azure Virtual Machine. Which of the following could be used to store the backups?

- A. Azure SQL Managed Instance
- B. Azure Storage Accounts 
- C. Azure CosmosDB
- D. Azure SQL Data Warehouse

Explanation:

Answer – B

You can store the backup's using the Azure Blob service in Azure storage accounts

The Microsoft documentation mentions the following

Use Azure Storage for SQL Server Backup and Restore

01/31/2017 • 5 minutes to read • 5 people like this +1

Overview

Starting with SQL Server 2012 SP1 CU2, you can now write SQL Server backups directly to the Azure Blob storage service. You can use this functionality to back up to and restore from the Azure Blob service with an on-premises SQL Server database or a SQL Server database in an Azure virtual machine. Backup to cloud offers benefits of availability, limitless geo-replicated off-site storage, and ease of migration of data to and from the cloud. You can issue BACKUP or RESTORE statements by using Transact-SQL or SMO.

The other options are incorrect since they can't be used to store the backups for the SQL Server.

For more information on SQL Server backup, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/virtual-machines/windows/sql/virtual-machines-windows-use-storage-sql-server-backup-restore>

Ask our Experts

Rate this Question?

View Queries

open ▾

Question 48

Unattempted

Domain :Design a business continuity strategy

View Case Study

Which of the following could be implemented the following requirement?

An SLA of 99.95% must be guaranteed on the Infrastructure for the middle tier system

You decide to implement Azure Virtual Machine scale sets

Would this fulfil the requirement?

A. Yes

B. No

Explanation:

Answer – B

refer to VM Scale Set SLA document at https://azure.microsoft.com/en-ca/support/legal/sla/virtual-machine-scale-sets/v1_1/

Virtual Machine Scale Sets is a free service, therefore, it does not have a financially backed SLA itself. However, if the Virtual Machine Scale Sets includes Virtual Machines in at least 2 Fault Domains, the availability of the underlying Virtual Machines SLA for two or more instances applies. If the scale set contains a single Virtual Machine, the availability for a Single Instance Virtual Machine applies.

Virtual Machine scale sets are used more from a scaling perspective and not from an availability perspective

For more information on virtual machine scale sets, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/virtual-machine-scale-sets/overview>

[Ask our Experts](#)

Rate this Question?  

[View Queries](#)

open ▾

Question 49

Unattempted

Domain :Design a business continuity strategy

[View Case Study](#)

Which of the following could be implemented the following requirement?

An SLA of 99.95% must be guaranteed on the Infrastructure for the middle tier system

You decide to implement an Azure Standard Load Balancer

Would this fulfil the requirement?

A. Yes 

B. No

Explanation:

Answer – A

Since we planned to Azure Standard Load Balancer, the SLA provided is 99.99%. And in the given case study, since the required SLA is 99.95% which is acceptable.

For more information on the Azure Load Balancer, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/load-balancer/load-balancer-overview>

https://azure.microsoft.com/en-ca/support/legal/sla/load-balancer/v1_0/

<https://docs.microsoft.com/en-us/azure/load-balancer/load-balancer-standard-overview>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 50

Unattempted

Domain :Design a business continuity strategy

[View Case Study](#)

Which of the following could be implemented the following requirement?

An SLA of 99.95% must be guaranteed on the Infrastructure for the middle tier system

You decide to implement Azure availability sets

Would this fulfil the requirement?

A. Yes

B. No 

Explanation:

Answer – B

You can achieve an availability of 99.95% with the use of availability sets, but since the requirement is "availability of middle tier in case of region failure", this option of working with Availability Sets does not work, reason Availability Sets work within a datacenter.

The Microsoft documentation mentions the following

Configure multiple virtual machines in an availability set for redundancy

To provide redundancy to your application, we recommend that you group two or more virtual machines in an availability set. This configuration within a datacenter ensures that during either a planned or unplanned maintenance event, at least one virtual machine is available and meets the 99.95% Azure SLA. For more information, see the [SLA for Virtual Machines](#).

For more information on availability sets, please go ahead and visit the below URL

Note - This is just for reference only.

The above requirement of availability of middle tier when region or datacentre failure happens can be had by implementing either of the 2 options that Azure provides.

Azure Traffic Manager - is a DNS-based **traffic** load balancer that enables you to distribute **traffic** optimally to services across global **Azure** regions, while providing high availability and responsiveness. ... An endpoint is any Internet-facing service hosted inside or outside of **Azure**.

<https://docs.microsoft.com/en-us/azure/traffic-manager/traffic-manager-overview>

Azure Front Door - enables you to define, manage, and monitor the global routing for your web traffic by optimizing for best performance and instant global failover for high availability. ... An application backend is any Internet-facing service hosted inside or **outside** of **Azure**.

<https://docs.microsoft.com/en-us/azure/frontdoor/front-door-overview>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Question 51

Unattempted

Domain :Determine workload requirements

View Case Study

There is a requirement to get detailed information for any HTTP request to the Azure Web App which results in an HTTP status code of 400 or greater. Which of the following would you use for this requirement?

- A. Metrics available for the Azure Web App
- B. Web server diagnostics for the Azure Web App
- C. Network Watcher for the Azure Web App
- D. Azure Advisor for the Azure Web App

Explanation:

Answer – B

The right approach is to use Web server diagnostics as is also shown in the Microsoft documentation.

Web server diagnostics and application diagnostics

App Service provides diagnostic functionality for logging information from both the web server and the web application. These are logically separated into **web server diagnostics** and **application diagnostics**.

Web server diagnostics

You can enable or disable the following kinds of logs:

- **Detailed Error Logging** - Detailed information for any request that results in HTTP status code 400 or greater. It may contain information that can help determine why the server returned the error code. One HTML file is generated for each error in the app's file system, and up to 50 errors (files) are retained. When the number of HTML files exceed 50, the oldest 26 files are automatically deleted.
- **Failed Request Tracing** - Detailed information on failed requests, including a trace of the IIS components used to process the request and the time taken in each component. It's useful if you want to improve site performance or isolate a specific HTTP error. One folder is generated for each error in the app's file system. File retention policies are the same as the detailed error logging above.
- **Web Server Logging** - Information about HTTP transactions using the [W3C extended log file format](#). It's useful when determining overall site metrics such as the number of requests handled or how many requests are from a specific IP address.

Since this is clearly given in the Microsoft documentation, all other options are incorrect

For more information on diagnostics logs for Azure Web Apps, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/app-service/troubleshoot-diagnostic-logs>

Ask our Experts

Rate this Question?  

View Queries

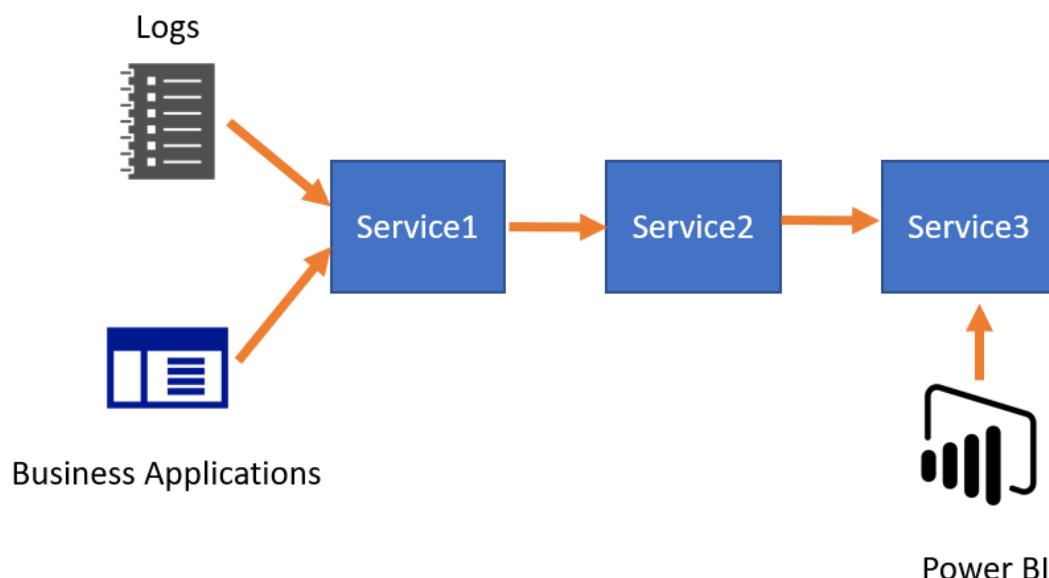
open ▾

Question 52

Unattempted

Domain :Design a data platform solution

A company needs to deploy a system that complies with the below architecture to Azure



Below is the requirement for the various services

- **Service1** – Orchestrate the movement of data from various log and applications sources to a large expansive data store
- **Service2** – Sustainable data store for taking large amounts of data
- **Service3** – Large data store with the ability to analyze and visualize data with Power BI

Which of the following would you use for Service1?

- A. Azure Databricks
- B. Azure Data Factory
- C. Azure Analysis Services
- D. Azure SQL Data Warehouse
- E. Azure Data Lake Storage

Explanation:

Answer – B

You can use the Azure Data Factory for this purpose.

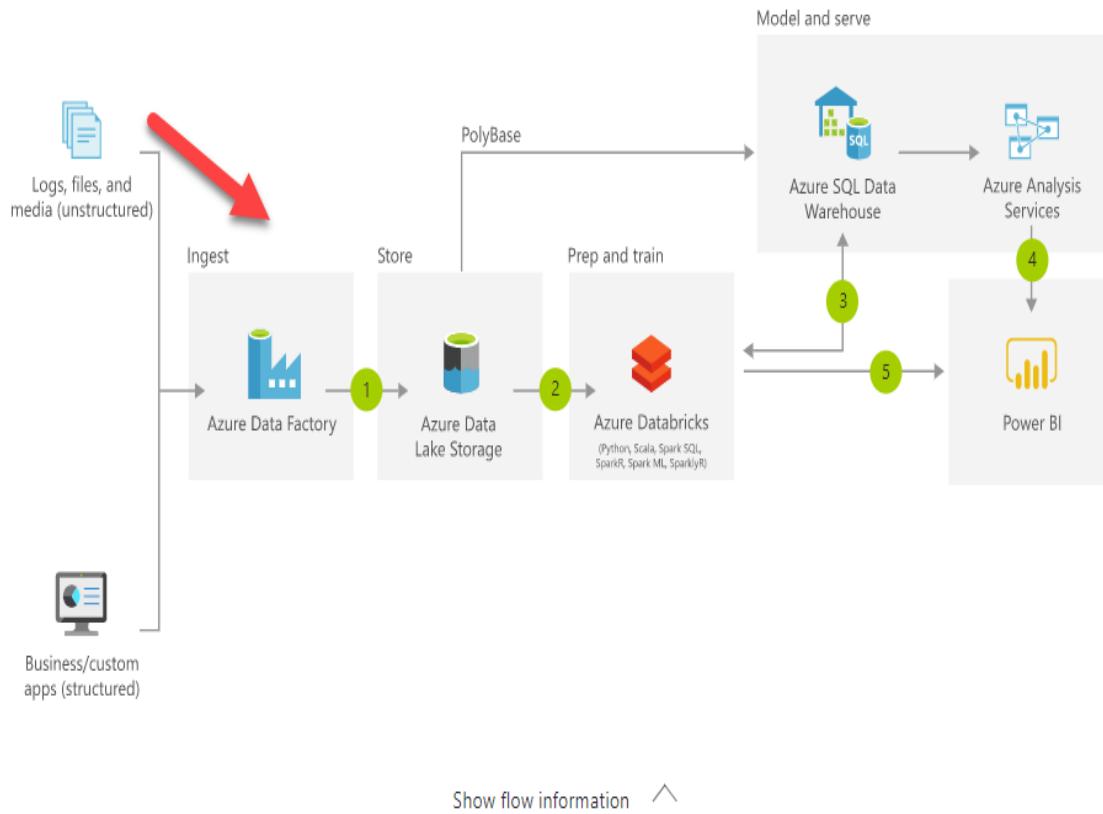
The Microsoft documentation mentions the following on Azure Data Factory

Azure Data Factory is the platform that solves such data scenarios. It is a *cloud-based data integration service* that allows you to create data-driven workflows in the cloud for orchestrating and automating data movement and data transformation. Using Azure Data Factory, you can create and schedule data-driven workflows (called pipelines) that can ingest data from disparate data stores. It can process and transform the data by using compute services such as Azure HDInsight Hadoop, Spark, Azure Data Lake Analytics, and Azure Machine Learning.

A sample architecture for such a scenario is also given in the Microsoft documentation

Modern data warehouse

A modern data warehouse lets you bring together all your data at any scale easily and to get insights through analytical dashboards, operational reports or advanced analytics for all your users.



- 1 Combine all your structured, unstructured and semi-structured data (logs, files and media) using Azure Data Factory to Azure Blob Storage.
- 2 Leverage data in Azure Blob Storage to perform scalable analytics with Azure Databricks and achieve cleansed and transformed data.
- 3 Cleansed and transformed data can be moved to Azure SQL Data Warehouse to combine with existing structured data, creating one hub for all your data.
- 4 Build operational reports and analytical dashboards on top of Azure Data Warehouse to derive insights from the data and use Azure Analysis Services to serve thousands of end users.
- 5 Run ad hoc queries directly on data within Azure Databricks.

Since this is clearly given as an example, all other options are incorrect

For more information on Azure Data Factory and for the sample architecture, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/data-factory/introduction>

<https://azure.microsoft.com/en-in/solutions/architecture/modern-data-warehouse/>

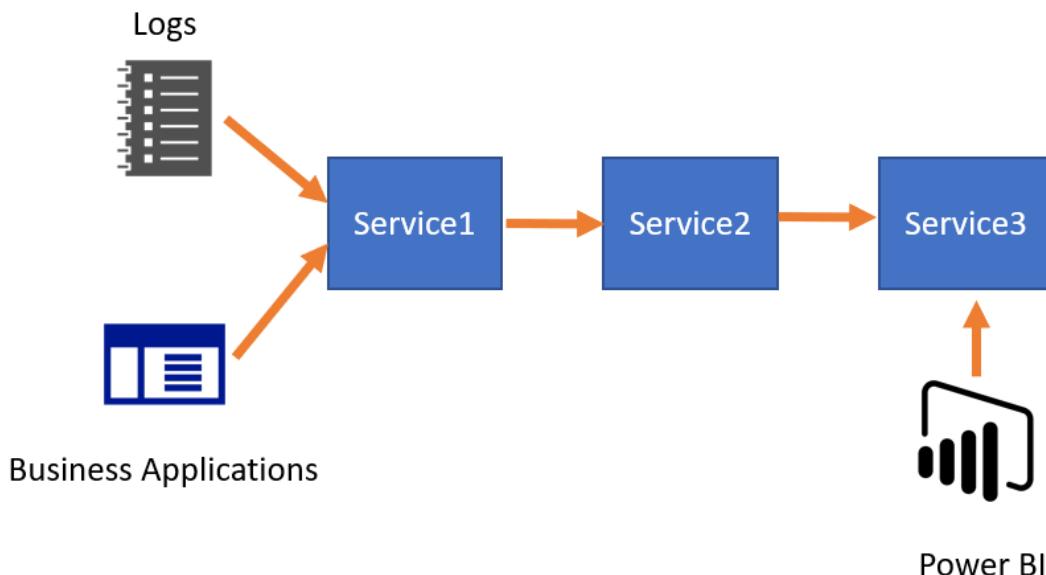
Ask our Experts

Rate this Question? 😊 😐

View Queries

open ▾

A company needs to deploy a system that complies with the below architecture to Azure



Below is the requirement for the various services

- **Service1** – Orchestrate the movement of data from various log and applications sources to a large expansive data store
- **Service2** – Sustainable data store for taking large amounts of data
- **Service3** – Large data store with the ability to analyze and visualize data with Power BI

Which of the following would you use for Service2?

- A. Azure Databricks
- B. Azure Data Factory
- C. Azure Analysis Services
- D. Azure SQL Data Warehouse
- E. Azure Data Lake Storage

Explanation:

Answer – E

An ideal solution for an expansive data storage solution is Data Lake storage

The Microsoft documentation mentions the following on Azure Data Lake Storage

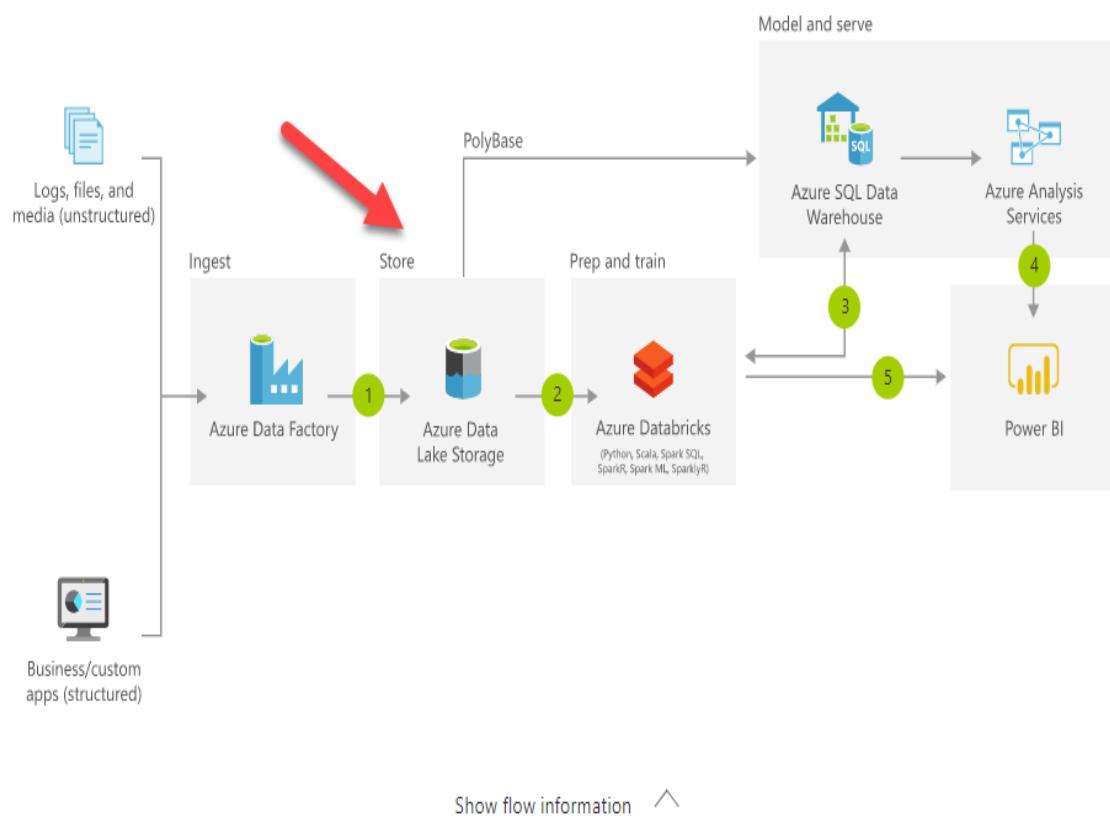
Designed for enterprise big data analytics

Data Lake Storage Gen2 makes Azure Storage the foundation for building enterprise data lakes on Azure. Designed from the start to service multiple petabytes of information while sustaining hundreds of gigabits of throughput, Data Lake Storage Gen2 allows you to easily manage massive amounts of data.

A sample architecture for such a scenario is also given in the Microsoft documentation

Modern data warehouse

A modern data warehouse lets you bring together all your data at any scale easily and to get insights through analytical dashboards, operational reports or advanced analytics for all your users.



- 1 Combine all your structured, unstructured and semi-structured data (logs, files and media) using Azure Data Factory to Azure Blob Storage.
- 2 Leverage data in Azure Blob Storage to perform scalable analytics with Azure Databricks and achieve cleansed and transformed data.
- 3 Cleansed and transformed data can be moved to Azure SQL Data Warehouse to combine with existing structured data, creating one hub for all your data.
- 4 Build operational reports and analytical dashboards on top of Azure Data Warehouse to derive insights from the data and use Azure Analysis Services to serve thousands of end users.
- 5 Run ad hoc queries directly on data within Azure Databricks.

Since this is clearly given as an example, all other options are incorrect

For more information on data lake storage and for the sample architecture, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/storage/blobs/data-lake-storage-introduction>

Ask our Experts

Rate this Question? 😊 😞

View Queries

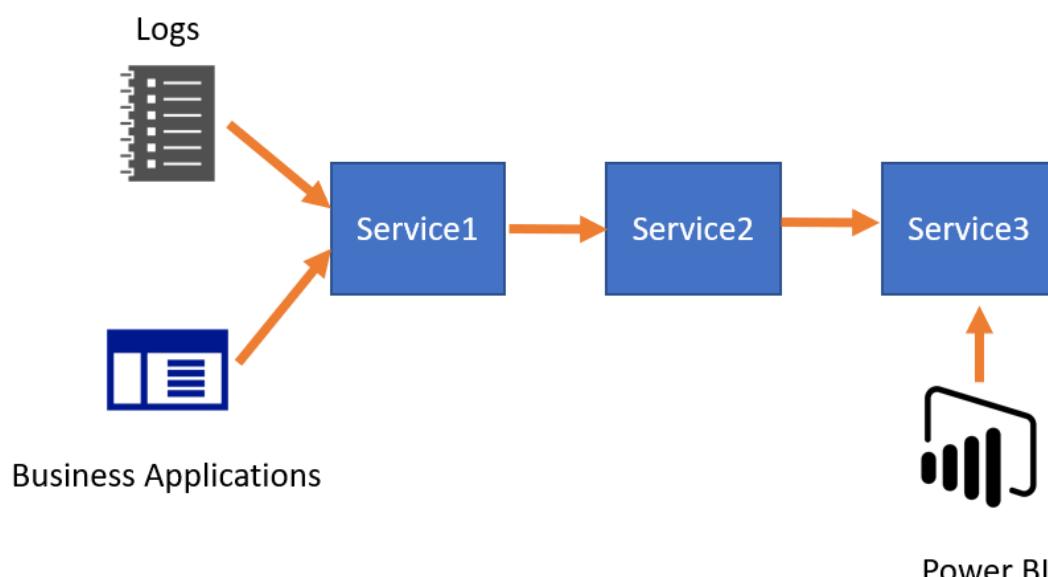
open ▾

Question 54

Unattempted

Domain :Design a data platform solution

A company needs to deploy a system that complies with the below architecture to Azure



Below is the requirement for the various services

- **Service1** – Orchestrate the movement of data from various log and applications sources to a large expansive data store
- **Service2** – Sustainable data store for taking large amounts of data
- **Service3** – Large data store with the ability to analyze and visualize data with Power BI

Which of the following would you use for Service3?

- A. Azure Databricks
- B. Azure Data Factory
- C. Azure Analysis Services
- D. Azure SQL Data Warehouse
- E. Azure Data Lake Storage

Explanation:

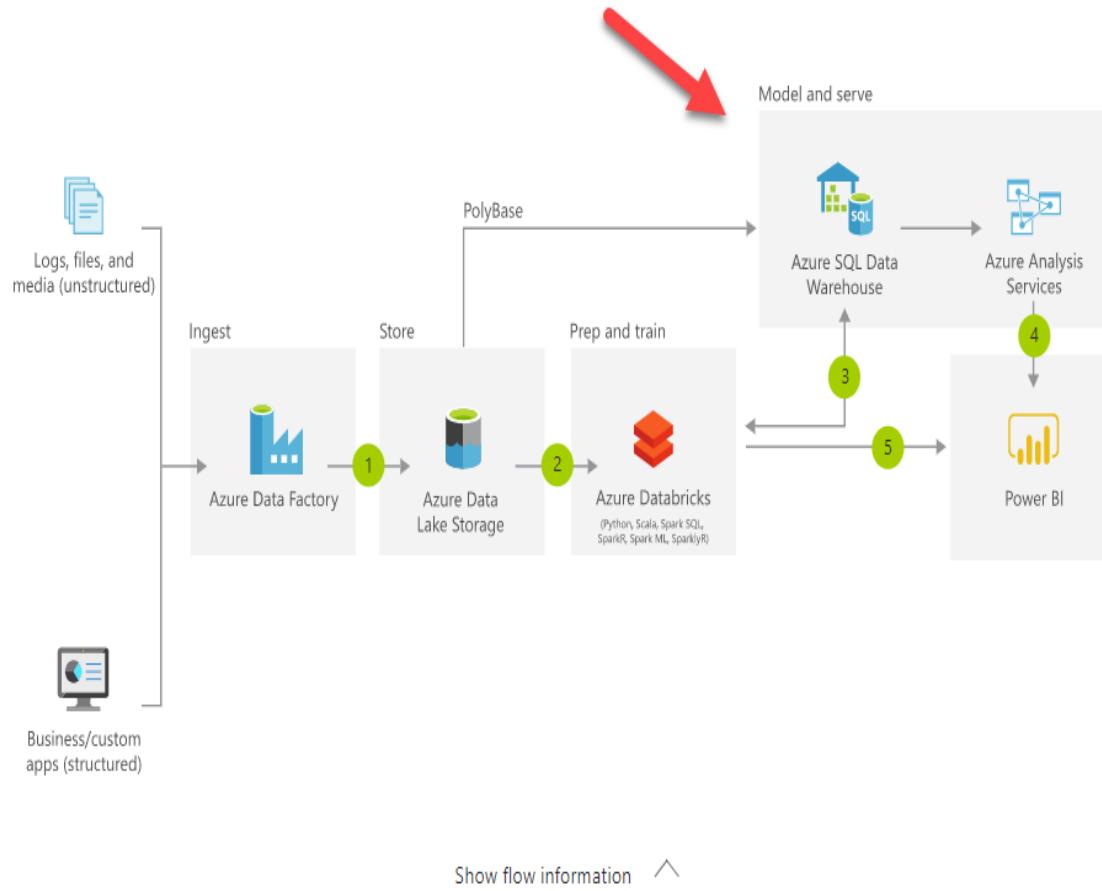
Answer D

An ideal data store that can store large amounts of data with an integration with Power BI is the SQL Data Warehouse

A sample architecture for such a scenario is also given in the Microsoft documentation

Modern data warehouse

A modern data warehouse lets you bring together all your data at any scale easily and to get insights through analytical dashboards, operational reports or advanced analytics for all your users.



Since this is clearly given as an example, all other options are incorrect

For more information on the integration between SQL Data warehouse and Power BI and for the sample architecture, please go ahead and visit the below URL

<https://powerbi.microsoft.com/en-us/integrations/azure-sql-data-warehouse/>

<https://azure.microsoft.com/en-in/solutions/architecture/modern-data-warehouse/>

View Queries

open ▾

Question 55

Unattempted

Domain :Determine workload requirements

A company is developing an ecommerce web application. One of the modules of the application will be built using a messaging solution architecture. The modules will have the following features

- A Workflow run for several items published on the web application.
- The Workflow would be built using Azure Logic Apps.
- The item data would be stored in Azure BLOB storage.

Which of the following would you additionally incorporate for the module?

- A. Azure Event Grid
- B. Azure Event Hub
- C. Azure HDInsight
- D. Azure Service Bus

Explanation:

Answer – D

You can use the Azure Service Bus queue system as the messaging service for the module. The Microsoft documentation mentions the following

What is Azure Service Bus?

09/22/2018 • 4 minutes to read • Contributors  all

Microsoft Azure Service Bus is a fully managed enterprise [integration](#) message broker. Service Bus is most commonly used to decouple applications and services from each other, and is a reliable and secure platform for asynchronous data and state transfer. Data is transferred between different applications and services using *messages*. A message is in binary format, which can contain JSON, XML, or just text.

Some common messaging scenarios are:

- Messaging: transfer business data, such as sales or purchase orders, journals, or inventory movements.
- Decouple applications: improve reliability and scalability of applications and services (client and service do not have to be online at the same time).
- Topics and subscriptions: enable 1:*n* relationships between publishers and subscribers.
- Message sessions: implement workflows that require message ordering or message deferral.

Option A is incorrect since this is normally used for event processing

Option B is incorrect since this is a big data ingestion service

Option C is incorrect since this is an analytics service

For more information on Azure Service Bus, please go to the below URL

<https://docs.microsoft.com/en-us/azure/service-bus-messaging/service-bus-messaging-overview>

Ask our Experts

Rate this Question?  

View Queries

open ▾

Finish Review

Certification

Cloud Certification

Java Certification

PM Certification

Big Data Certification

Company

Support

Discussions

Blog

Business

Follow us



© Copyright 2020. Whizlabs Software Pvt. Ltd. All Right Reserved.