Practice Test 3

Completed on 17-June-2020

Attempt	Marks Obtained	Your score	Time Taken	Result
03	55 / 55	100%	00 H 04 M 45 S	Congratulations! Passed

Share your Result with your friends

Domains wise Quiz Performance Report

No	Domain	Total Question	Correct	Incorrect	Unattempted	Marked as Review
1	Implement data storage solutions	25	25	0	0	0
2	Monitor and optimize data solutions	10	10	0	0	0
3	Manage and develop data processing	20	20	0	0	0
Total	All Domain	55	55	0	0	0

Review the Answers

Sorting by All

Question 1 **Correct**

Domain :Implement data storage solutions

You need to import data into a Microsoft SQL Data Warehouse. The data that needs to be ingested resides in parquet files. These files are stored in an Azure Data Lake Gen 2 storage account. You need to load the data from the storage account into the data warehouse.
You decide to implement the following steps

Create an external data source pointing to the Azure storage account

Create a workload group using the Azure storage account name as the pool name

Load the data using the CREATE TABLE AS SELECT statement

Would these steps fulfil the requirement?

A. Yes

B. No

Explanation:

Answer – B

You need to create an external data source that maps to an Azure Data Lake Gen 2 storage account.

For more information on integration between SQL data warehouse and Azure Data Lake stores, please visit the following URL

https://azure.microsoft.com/es-es/blog/sql-data-warehouse-now-supports-seamless-integration-with-azure-data-lake-store/

Ask our Experts

Rate this Question?

View Queries

open

Question 2

Correct

Domain :Implement data storage solutions

1. You need to import data into a Microsoft SQL Data Warehouse. The data that needs to be ingested resides in parquet files. These files are stored in an Azure Data Lake Gen 2

storage account. You need to load the data from the storage account into the data warehouse.

You decide to implement the following steps

Create a remote service binding pointing to the Azure Data Lake Gen 2 storage account

Create an external file format and external table using the external data source

Load the data using the CREATE TABLE AS SELECT statement

Would these steps fulfil the requirement?

A. Yes

B. No

Explanation:

Answer – B

You need to create an external data source that maps to an Azure Data Lake Gen 2 storage account.

For more information on integration between SQL data warehouse and Azure Data Lake stores, please visit the following URL

https://azure.microsoft.com/es-es/blog/sql-data-warehouse-now-supports-seamless-integration-with-azure-data-lake-store/

Ask our Experts

Rate this Question?

View Queries

open

Question 3

Correct

Domain :Implement data storage solutions

You need to import data into a Microsoft SQL Data Warehouse. The data that needs to be ingested resides in parquet files. These files are stored in an Azure Data Lake Gen 2

storage account. You need to load the data from the storage account into the data warehouse.
You decide to implement the following steps

Use Azure Data Factory to convert the parquet files to CSV files

Create an external data source pointing to the Azure storage account

Create an external file format and external table using the external data source

Load the data using the INSERT..SELECT statement

Would these steps fulfil the requirement?

A. Yes

B. No

Explanation:

Answer – B

You can directly ingest parquet-based files and you don't need to convert the files.

For more information on integration between SQL data warehouse and Azure Data Lake stores, please visit the following URL

https://azure.microsoft.com/es-es/blog/sql-data-warehouse-now-supports-seamless-integration-with-azure-data-lake-store/

Ask our Experts

Rate this Question?

View Queries

open

Question 4

Correct

Domain :Implement data storage solutions

You need to import data into a Microsoft SQL Data Warehouse. The data that needs to be ingested resides in parquet files. These files are stored in an Azure Data Lake Gen 2

storage account. You need to load the data from the storage account into the data warehouse.
You decide to implement the following steps

Create an external data source that points to the Azure Data Lake Gen 2 storage account

Then create and external file format and an external table by making use of an external data source

Then load the data using the CREATE TABLE AS SELECT statement

Would these steps fulfil the requirement?

- A. **Yes** ✓
- B. **No**

Explanation:

Answer – A

Yes, this is given as an example in one of the Microsoft blog articles

```
-- B: Create a database scoped credential
-- IDENTITY: Pass the client id and OAuth 2.0 Token Endpoint taken from your Azure Active Directory Application
-- SECRET: Provide your AAD Application Service Principal key.
-- For more information on Create Database Scoped Credential: https://msdn.microsoft.com/en-us/library/mt270260.aspx
CREATE DATABASE SCOPED CREDENTIAL ADLUser
WITH IDENTITY = '<clientID>@<OAuth2.0TokenEndPoint>',
     SECRET = '<KEY>' ;


-- C: Create an external data source
-- TYPE: HADOOP - PolyBase uses Hadoop APIs to access data in Azure Data Lake Store.
-- LOCATION: Provide Azure Data Lake accountname and URI
-- CREDENTIAL: Provide the credential created in the previous step.
CREATE EXTERNAL DATA SOURCE AzureDataLakeStore
WITH ( TYPE = HADOOP, LOCATION = 'adl://<ADLS Account Name>.azuredatalake.net',
       CREDENTIAL = ADLUser );

 -- D: Create an external file format
 -- FIELDTERMINATOR: Marks the end of each field (column) in a delimited text file
 -- STRINGDELIMITER: Specifies the field terminator for data of type string in the text-delimited file.
 -- DATE_FORMAT: Specifies a custom format for all date and time data that might appear in a delimited text file.
 -- UseTypeDefault: Store all Missing values as NULL

CREATE EXTERNAL FILE FORMAT TextFileFormat
WITH (
FORMATTYPE = DELIMITEDTEXT,
FORMATOPTIONS ( FIELDTERMINATOR = '|',
STRINGDELIMITER = '' ,
DATEFORMAT = 'yyyy-MM-dd HH:mm:ss.fff');

-- E: Create an External Table

CREATE EXTERNAL TABLE [dbo].[DimProduct_external]
( [ProductKey] [int] NOT NULL,
[ProductLabel] nvarchar NULL,
[ProductName] nvarchar NULL )
WITH ( LOCATION='/DimProduct/' ,
       DATA_SOURCE = AzureDataLakeStore ,
       FILE_FORMAT = TextFileFormat ,
       REJECT_TYPE = VALUE ,
       REJECT_VALUE = 0 ) ;

-- F: Load Data with CTAS
CREATE TABLE [dbo].[DimProduct]
WITH (DISTRIBUTION = HASH([ProductKey] ) )
AS SELECT * FROM
[dbo].[DimProduct_external] ;
```

For more information on integration between SQL data warehouse and Azure Data Lake stores, please visit the following URL

https://azure.microsoft.com/es-es/blog/sql-data-warehouse-now-supports-seamless-integration-with-azure-data-lake-store/

Ask our Experts

Rate this Question?

View Queries open

Question 5

Correct

Domain :Implement data storage solutions

A company wants to use Azure Storage accounts for file storage purposes. A single storage account would be required to perform all read, write and delete operations. The company also needs to keep a copy of all historical operations in an on-premise server.
Which of the following actions need to be performed to accomplish this requirement?
Choose 2 answers from the options given below

- [x] A. **Configure a storage account to log read, write and delete operations with the service type of Blob**
- [ ] B. **Configure a storage account to log read, write and delete operations with the service type of table**
- [ ] C. **Configure a storage account to log read, write and delete operations with the service type of queue**
- [x] D. **Use the AzCopy tool to download the log data from $logs/blob**
- [ ] E. **Use the AzCopy tool to download the log data from $logs/table**
- [ ] F. **Use the AzCopy tool to download the log data from $logs/queue**

Explanation:

Answer – A and D

Since the companies wants to store files, one can use the service type of Blob. The Microsoft documentation mentions the following

Introduction to Azure Blob storage

11/04/2019 • 4 minutes to read •  +5

Azure Blob storage is Microsoft's object storage solution for the cloud. Blob storage is optimized for storing massive amounts of unstructured data. Unstructured data is data that does not adhere to a particular data model or definition, such as text or binary data.

About Blob storage

Blob storage is designed for:

- Serving images or documents directly to a browser.
- Storing files for distributed access.
- Streaming video and audio.
- Writing to log files.
- Storing data for backup and restore, disaster recovery, and archiving.
- Storing data for analysis by an on-premises or Azure-hosted service.

The logs are stored in the $logs container

How logs are stored

All logs are stored in block blobs in a container named `$logs`, which is automatically created when Storage Analytics is enabled for a storage account. The `$logs` container is located in the blob namespace of the storage account, for example: `http://<accountname>.blob.core.windows.net/$logs`. This container cannot be deleted once Storage Analytics has been enabled, though its contents can be deleted. If you use your storage-browsing tool to navigate to the container directly, you will see all the blobs that contain your logging data.

All other options are incorrect since the Blob service type is used to store objects such as files

For more information on Azure Blob storage and storage logging, please visit the following URL

https://docs.microsoft.com/en-us/azure/storage/blobs/storage-blobs-introduction

https://docs.microsoft.com/en-us/azure/storage/common/storage-analytics-logging

Ask our Experts

Rate this Question? ☺ ☹

View Queries

open

Question 6

Correct

Domain :Implement data storage solutions

You have to create a new single database instance of Microsoft Azure SQL database. You must ensure that client connections are accepted via a workstation. The workstation will use SQL Server Management Studio to connect to the database instance.
Which of the following Powershell commands would you execute to create and configure the database? Choose 3 answers from the options below

- [ ] A. **New-AzureRmSqlElasticPool**
- [x] B. **New-AzureRmSqlServerFirewallRule**
- [x] C. **New-AzureRmSqlServer**
- [ ] D. **New-AzureRmSqlServerVirtualNetworkRule**
- [x] E. **New-AzureRmSqlDatabase**

Explanation:

Answer – B,C and E

The Microsoft documentation clearly gives the steps to create and configure the database. Please note the below snippet shows the new powershell commands, but you can also use the older Azure PowerShell commands.

```
# Set subscription
Set-AzContext -SubscriptionId $subscriptionId

# Create a resource group
$resourceGroup = New-AzResourceGroup -Name $resourceGroupName -Location $locati

# Create a server with a system wide unique server name
$server = New-AzSqlServer -ResourceGroupName $resourceGroupName `
    -ServerName $serverName `
    -Location $location `
    -SqlAdministratorCredentials $(New-Object -TypeName System.Management.Autom

# Create a server firewall rule that allows access from the specified IP range
$serverFirewallRule = New-AzSqlServerFirewallRule -ResourceGroupName $resourceG
    -ServerName $serverName `
    -FirewallRuleName "AllowedIPs" -StartIpAddress $startIp -EndIpAddress $endI

# Create a blank database with an S0 performance level
$database = New-AzSqlDatabase  -ResourceGroupName $resourceGroupName `
    -ServerName $serverName `
    -DatabaseName $databaseName `
    -RequestedServiceObjectiveName "S0" `
    -SampleName "AdventureWorksLT"

# Clean up deployment
# Remove-AzResourceGroup -ResourceGroupName $resourceGroupName
```

Since this is clearly given in the documentation, all other options are incorrect

For more information on using PowerShell to create and configure an Azure SQL database, please visit the following URL

https://docs.microsoft.com/en-us/azure/sql-database/scripts/sql-database-create-and-configure-database-powershell

Question 7 **Correct**

Domain :Implement data storage solutions

A company is planning on creating an Azure SQL database to support a mission critical application. The application needs to be highly available and not have any performance degradation during maintenance windows. Which of the following technologies can be used to implement this solution? Choose 3 answers from the options given below

- [x] A. **Premium Service Tier**
- [ ] B. **Virtual Machine Scale Sets**
- [ ] C. **Basic Service Tier**
- [ ] D. **SQL Data Sync**
- [x] E. **Always On Availability Groups**
- [x] F. **Zone-redundant configuration**

Explanation:

Answer – A, E and F

The Premium Service tier has better availability for Azure SQL Server databases

The Microsoft documentation mentions the following

Premium and Business Critical service tier availability

Premium and Business Critical service tiers leverage the Premium availability model, which integrates compute resources (SQL Server Database Engine process) and storag (locally attached SSD) on a single node. High availability is achieved by replicating bot compute and storage to additional nodes creating a three to four-node cluster.

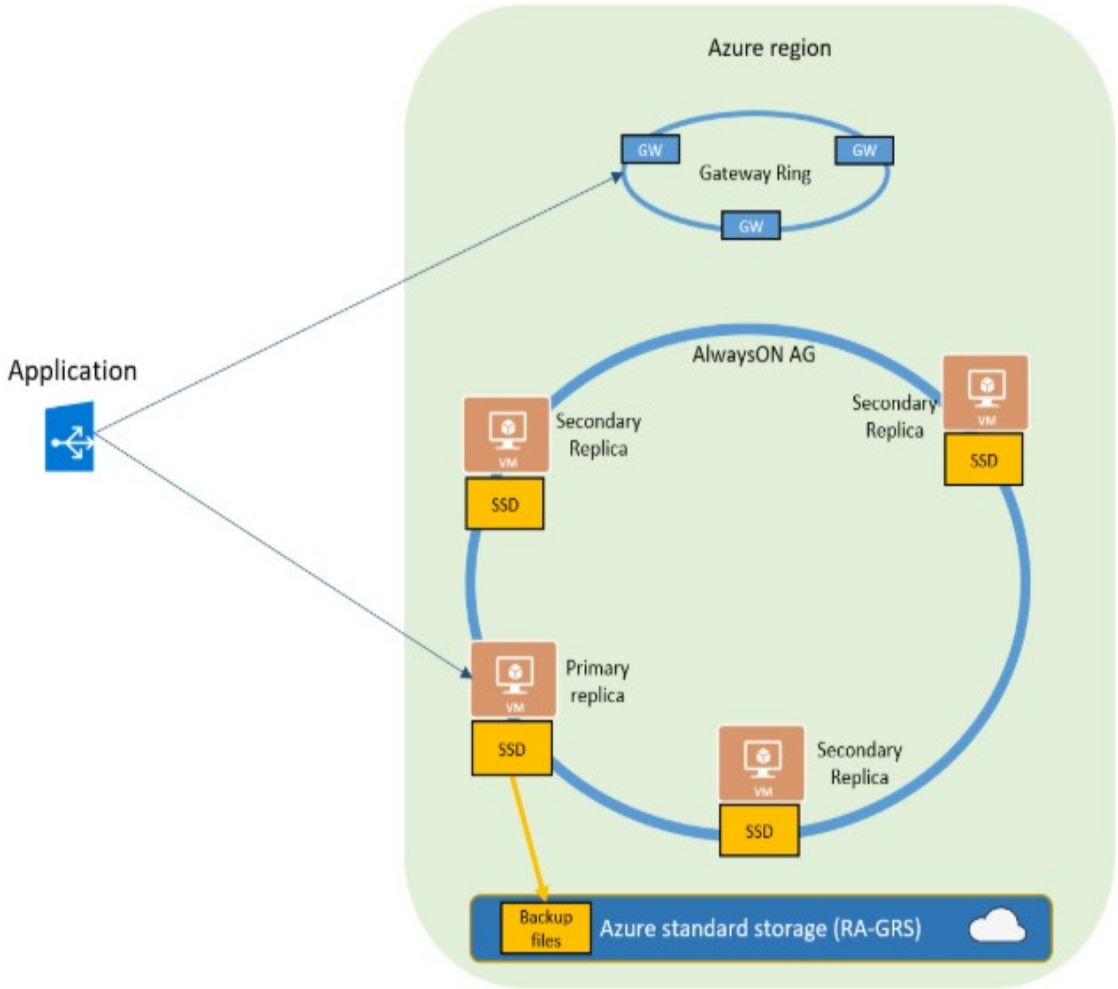

You can use Zone-redundant configuration to increase the availability of nodes in the database

The Microsoft documentation mentions the following

Zone redundant configuration

By default, the cluster of nodes for the premium availability model is created in the same datacenter. With the introduction of Azure Availability Zones, SQL Database can place different replicas of the Business Critical database to different availability zones in the same region. To eliminate a single point of failure, the control ring is also duplicated across multiple zones as three gateway rings (GW). The routing to a specific gateway ring is controlled by Azure Traffic Manager (ATM). Because the zone redundant configuration in the Premium or Business Critical service tiers does not create additional database redundancy, you can enable it at no extra cost. By selecting a zone redundant configuration, you can make your Premium or Business Critical databases resilient to a much larger set of failures, including catastrophic datacenter outages, without any changes to the application logic. You can also convert any existing Premium or Business Critical databases or pools to the zone redundant configuration.

If you have SQL servers on Azure Virtual Machines, you can also achieve high availability with the help of On Always Availability Groups.

Option B is incorrect since this is used to scale virtual machines based on the load or demand

Option C is incorrect since this service tier does not provide high availability options

Option D is incorrect since this feature is used to keep multiple databases in sync

For more information on achieving high availability for SQL Server databases on Azure, please visit the following URL

https://docs.microsoft.com/en-us/azure/sql-database/sql-database-high-availability

https://docs.microsoft.com/en-us/azure/virtual-machines/windows/sql/virtual-machines-windows-portal-sql-availability-group-tutorial

View Queries open

Question 8 Correct

Domain :Monitor and optimize data solutions

View Case Study

Which of the following would go into Slot 1?

- A. delete
- B. blockBlob
- C. baseBlob
- D. snapshot
- E. tierToCool
- F. tierToArchive

Explanation:

Answer – C

Here we need to place an action on “baseBlob”. This is also given as an example in the Microsoft documentation when it comes to lifecycle policies

JSON Copy

```
{
  "rules": [
    {
      "name": "ruleFoo",
      "enabled": true,
      "type": "Lifecycle",
      "definition": {
        "filters": {
          "blobTypes": [ "blockBlob" ],
          "prefixMatch": [ "container1/foo" ]
        },
        "actions": {
          "baseBlob": {
            "tierToCool": { "daysAfterModificationGreaterThan": 30 },
            "tierToArchive": { "daysAfterModificationGreaterThan": 90
            "delete": { "daysAfterModificationGreaterThan": 2555 }
          },
          "snapshot": {
            "delete": { "daysAfterCreationGreaterThan": 90 }
          }
        }
      }
    }
  ]
}
```

Since this is clearly given in the Microsoft documentation, all other options are incorrect

For more information on lifecycle policies, please visit the following URL

https://docs.microsoft.com/en-us/azure/storage/blobs/storage-lifecycle-management-concepts?tabs=azure-portal

Ask our Experts

Rate this Question?

View Queries

open

Question 9 Correct

Domain :Monitor and optimize data solutions

View Case Study

Which of the following would go into Slot 2?

- A. **delete**
- B. **blockBlob**
- C. **baseBlob**
- D. **snapshot**
- E. **tierToCool**
- F. **tierToArchive** ✓

Explanation:

Answer – F

Now since after a year, the question mentions that the data would be used for auditing purposes for 10 years. And that you can have 10 days to make the data available, the most effective measure is to move the data to the Archive tier.

Since this is the most suitable option, all other options are incorrect.

For more information on lifecycle policies, please visit the following URL

https://docs.microsoft.com/en-us/azure/storage/blobs/storage-lifecycle-management-concepts?tabs=azure-portal

Ask our Experts Rate this Question? ☺ ☹

View Queries open

Question 10 Correct

Domain :Monitor and optimize data solutions

View Case Study

Which of the following would go into Slot 3?

- A. **delete** ✓
- B. **blockBlob**
- C. **baseBlob**
- D. **snapshot**
- E. **tierToCool**
- F. **tierToArchive**

Explanation:

Answer – A

Since after 3650 days or 10 years, the data is no longer required, the data can be deleted.

Since this is the most suitable option, all other options are incorrect.

For more information on lifecycle policies, please visit the following URL

https://docs.microsoft.com/en-us/azure/storage/blobs/storage-lifecycle-management-concepts?tabs=azure-portal

Ask our Experts Rate this Question? ☺ ☹

View Queries open

Question 11 Correct

View Case Study

Which of the following would go into Slot 4?

- A. **Delete**
- B. **blockBlob**
- C. **baseBlob**
- D. **snapshot** ✓
- E. **tierToCool**
- F. **tierToArchive**

Explanation:

Answer – D

Here we need to add a rule for the snapshot data.

This is also given as an example in the Microsoft documentation when it comes to lifecycle policies

JSON Copy

```
{
  "rules": [
    {
      "name": "ruleFoo",
      "enabled": true,
      "type": "Lifecycle",
      "definition": {
        "filters": {
          "blobTypes": [ "blockBlob" ],
          "prefixMatch": [ "container1/foo" ]
        },
        "actions": {
          "baseBlob": {
            "tierToCool": { "daysAfterModificationGreaterThan": 30 },
            "tierToArchive": { "daysAfterModificationGreaterThan": 90
            "delete": { "daysAfterModificationGreaterThan": 2555 }
          },
          "snapshot": {
            "delete": { "daysAfterCreationGreaterThan": 90 }
          }
        }
      }
    }
  ]
}
```

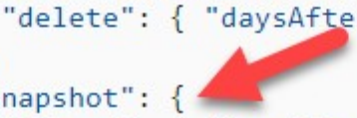

Since this is clearly given in the Microsoft documentation, all other options are incorrect

For more information on lifecycle policies, please visit the following URL

https://docs.microsoft.com/en-us/azure/storage/blobs/storage-lifecycle-management-concepts?tabs=azure-portal

Ask our Experts

Rate this Question?

View Queries

open

Question 12 Correct

Domain :Monitor and optimize data solutions

View Case Study

Which of the following would go into Slot 5?

- A. delete
- B. blockBlob
- C. baseBlob
- D. snapshot
- E. tierToCool
- F. tierToArchive

Explanation:

Answer – E

Now after 60 days, the question mentions that the data is just required for weekly reports , we can move the data to the cool tier which is more cost effective in nature.

Since this is the most suitable option, all other options are incorrect.

For more information on lifecycle policies, please visit the following URL

https://docs.microsoft.com/en-us/azure/storage/blobs/storage-lifecycle-management-concepts?tabs=azure-portal

Ask our Experts Rate this Question?

View Queries open

Question 13 **Correct**

Domain :Manage and develop data processing

A company has an Azure SQL data warehouse. They want to use PolyBase to retrieve data from an Azure Blob storage account and ingest into the Azure SQL data warehouse. The files are stored in parquet format. The data needs to be loaded into a table called whizlab_sales. Which of the following actions need to be performed to implement this requirement? Choose 4 answers from the options given below

- [ ] A. **Create an external file format that would map to the parquet-based files**
- [x] B. **Load the data into a staging table**
- [x] C. **Create an external table called whizlab_sales_details**
- [x] D. **Create an external data source for the Azure Blob storage account**
- [x] E. **Create a master key on the database**
- [ ] F. **Configure Polybase to use the Azure Blob storage account**

Explanation:

Answer – B,C,D,E

There is an article on github as part of the Microsoft documentation that provides details on how to load data into an Azure SQL data warehouse from an Azure Blob storage account. The key steps are

Creating a master key in the database

Creating an external data source for the Azure Blob storage account

3. Create a master key for the MySampleDataWarehouse database. You only need to create a master key once per database.

```
CREATE MASTER KEY;
```

4. Run the following CREATE EXTERNAL DATA SOURCE statement to define the location of the Azure blob. This is the location of the external taxi cab data. To run a command that you have appended to the query window, highlight the commands you wish to run and click **Execute**.

```
CREATE EXTERNAL DATA SOURCE NYTPublic
WITH
(
    TYPE = Hadoop,
    LOCATION = 'wasbs://2013@nytaxiblob.blob.core.windows.net/'
);
```

Next you need to create an external table

7. Create the external tables. The table definitions are stored in SQL Data Warehouse, but the tables reference data that is stored in Azure blob storage. Run the following T-SQL commands to create several external tables that all point to the Azure blob we defined previously in our external data source.

Next you load the data. But it is always beneficial to load the data into a staging table first

Load the data into your data warehouse

This section uses the external tables you just defined to load the sample data from Azure Storage Blob to SQL Data Warehouse.

> [!NOTE] This tutorial loads the data directly into the final table. In a production environment, you will usually use CREATE TABLE AS SELECT to load into a staging table. While data is in the staging table you can perform any necessary transformations. To append the data in the staging table to a production table, you can use the INSERT...SELECT statement. For more information, see Inserting data into a production table.

Since this is clearly provided in the documentation, all other options are incorrect

For more information on loading data from Azure Blob to Azure SQL data warehouse, please visit the following URL

https://github.com/MicrosoftDocs/azure-docs/blob/master/articles/sql-data-warehouse/load-data-from-azure-blob-storage-using-polybase.md

Ask our Experts

Rate this Question?

View Queries

open

Question 14

Correct

Domain :Manage and develop data processing

A company wants to implement a solution that would perform batch processing of geospatial data.
Which of the following could be used as a native client to run interactive queries and batch processes?

- A. HDInsight Tools for Visual Studio
- B. Hive View
- C. HDInsight REST API
- D. Azure Data Factory

Explanation:

Answer – A

Azure HDInsight can be used to perform batch processing. The Microsoft documentation mentions the following

HDInsight

HDInsight is a managed Hadoop service. Use it deploy and manage Hadoop clusters in Azure. For batch processing, you can use Spark, Hive, Hive LLAP, MapReduce.

- Languages: R, Python, Java, Scala, SQL
- Kerberos authentication with Active Directory, Apache Ranger based access control
- Gives you full control of the Hadoop cluster

You can then use HDInsight Tools for Visual Studio as a native client to run interactive queries

Option B is incorrect since this is Web UI tool that can be used to execute queries against Apache Hadoop in HDInsight

Option C is incorrect since this is used to execute REST API calls against HDInsight

Option D is incorrect since this can be used as an ETL or ELT tool

For more information on batch processing solutions and HDInsight tools for visual studio, please visit the following URL

https://docs.microsoft.com/en-us/azure/architecture/data-guide/technology-choices/batch-processing

https://azure.microsoft.com/en-us/resources/videos/hdinsight-tools-for-visual-studio/

Question 15

Correct

Domain :Manage and develop data processing

A company wants to implement a solution that would perform batch processing of geospatial data.
Which of the following can be used as a web browser to run interactive queries and batch processes?

- A. **HDInsight Tools for Visual Studio**
- B. **Hive View** ✅
- C. **HDInsight REST API**
- D. **Azure Data Factory**

Explanation:

Answer – B

You can use Apache Hive view to run interactive queries via the web browser

The Microsoft documentation mentions the following

Option A is incorrect since this is a native tool that can be used to run interactive queries against HDInsight

Option C is incorrect since this is used to execute REST API calls against HDInsight

Option D is incorrect since this can be used as an ETL or ELT tool

For more information on Apache Hive View, please visit the following URL

https://docs.microsoft.com/en-us/azure/hdinsight/hadoop/apache-hadoop-use-hive-ambari-view

Ask our Experts

Rate this Question?

View Queries

open

Question 16

Correct

Domain :Manage and develop data processing

A company wants to implement a solution that would perform batch processing of geospatial data.
Which of the following can be used to develop batch processing applications that use Azure HDInsight?

A. HDInsight Tools for Visual Studio

B. Hive View

C. HDInsight REST API

D. Azure Data Factory

Explanation:

Answer – C

You can use the HDInsight REST API to issue calls to the HDInsight cluster.

The Microsoft documentation mentions the following

Azure HDInsight REST API

11/16/2016 • 2 minutes to read •

Use these APIs to create and manage HDInsight resources through Azure Resource Manager. All task operations conform to the HTTP/1.1 protocol specification and each operation returns an x-ms-request-id header that can be used to obtain information about the request. Ensure that requests made to these resources are secure. For more information, see Authenticating Azure Resource Manager requests.

Option A is incorrect since this is a native tool that can be used to run interactive queries against HDInsight

Option B is incorrect since this is Web UI tool that can be used to execute queries against Apache Hadoop in HDInsight

Option D is incorrect since this can be used as an ETL or ELT tool

For more information on Azure HDInsight REST API, please visit the following URL

https://docs.microsoft.com/en-us/rest/api/hdinsight/

Ask our Experts

Rate this Question?

View Queries

open

Question 17

Correct

Domain :Manage and develop data processing

A company is designing a new lambda architecture on Microsoft Azure. Below are the requirements for each architecture layer

Ingestion

Have the ability to receive millions of events per second

Be a fully managed platform-as-a-service solution

Integrate with Azure Functions

Stream Processing

Have the ability to process data on a per-job basis

Have the ability to connect seamlessly to Azure services

Use a SQL based query language

Analytical data store

Perform as a managed service

Ability to behave as a document store

Provide data encryption at rest

Which of the following would you consider using at the Ingestion layer?

A. **HDInsight Kafta**

B. **Azure Event Hubs** ✓

C. **HDInsight Storm**

D. **HDInsight Spark**

Explanation:

Answer - B

Azure Event Hubs is a platform as a service that can be used to ingest millions of events per second

The Microsoft documentation mentions the following

Azure Event Hubs — A big data streaming platform and event ingestion service

12/06/2018 • 3 minutes to read • +2

Azure Event Hubs is a big data streaming platform and event ingestion service. It can receive and process millions of events per second. Data sent to an event hub can be transformed and stored by using any real-time analytics provider or batching/storage adapters.

The following scenarios are some of the scenarios where you can use Event Hubs:

- Anomaly detection (fraud/outliers)
- Application logging
- Analytics pipelines, such as clickstreams
- Live dashboarding
- Archiving data
- Transaction processing
- User telemetry processing
- Device telemetry streaming

Since this is clear from the Microsoft documentation, all other options are incorrect

For more information on Azure Event Hubs, please visit the following URL

https://docs.microsoft.com/en-us/azure/event-hubs/event-hubs-about

Question 18 **Correct**

Domain :Manage and develop data processing

A company is designing a new lambda architecture on Microsoft Azure. Below are the requirements for each architecture layer

Ingestion

Have the ability to receive millions of events per second

Be a fully managed platform-as-a-service solution

Integrate with Azure Functions

Stream Processing

Have the ability to process data on a per-job basis

Have the ability to connect seamlessly to Azure services

Use a SQL based query language

Analytical data store

Perform as a managed service

Ability to behave as a document store

Provide data encryption at rest

Which of the following would you consider using at the Stream Processing layer?

- A. **Azure Stream Analytics** ✓
- B. **HDInsight with Spark Streaming**
- C. **Azure Cosmos DB Change feed**
- D. **Azure Analysis Services**

Explanation:

Answer – A

Azure Stream Analytics can be used for the streaming layer.

The Microsoft documentation mentions the following

What is Azure Stream Analytics?

06/21/2019 • 6 minutes to read • +18

Azure Stream Analytics is a real-time analytics and complex event-processing engine that is designed to analyze and process high volumes of fast streaming data from multiple sources simultaneously. Patterns and relationships can be identified in information extracted from a number of input sources including devices, sensors, clickstreams, social media feeds, and applications. These patterns can be used to trigger actions and initiate workflows such creating alerts, feeding information to a reporting tool, or storing transformed data for later use. Also, Stream Analytics is available on Azure IoT Edge runtime, and supports the same exact language or syntax as cloud.

You can also configure Azure Stream Analytics on a per job basis and you can use SQL query language to query the data. The Microsoft documentation mentions the following on this

How does Stream Analytics work?

An Azure Stream Analytics job consists of an input, query, and an output. Stream Analytics ingests data from Azure Event Hubs, Azure IoT Hub, or Azure Blob Storage. The query, which is based on SQL query language, can be used to easily filter, sort, aggregate, and join streaming data over a period of time. You can also extend this SQL language with JavaScript and C# user defined functions (UDFs). You can easily adjust the event ordering options and duration of time windows when preforming aggregation operations through simple language constructs and/or configurations.

Since this is clear from the Microsoft documentation, all other options are incorrect

For more information on Azure Stream Analytics, please visit the following URL

https://docs.microsoft.com/en-us/azure/stream-analytics/stream-analytics-introduction

Ask our Experts

Rate this Question?

View Queries

open

Question 19

Correct

Domain :Manage and develop data processing

A company is designing a new lambda architecture on Microsoft Azure. Below are the requirements for each architecture layer

Ingestion

Have the ability to receive millions of events per second

Be a fully managed platform-as-a-service solution

Integrate with Azure Functions

Stream Processing

Have the ability to process data on a per-job basis

Have the ability to connect seamlessly to Azure services

Use a SQL based query language

Analytical data store

Perform as a managed service

Ability to behave as a document store

Provide data encryption at rest

Which of the following would you consider using at the Analytical data store layer?

A. **Hive LLAP on HDInsight**

B. **Azure Analysis Services**

C. **Azure Cosmos DB**

D. **SQL Data Warehouse**

Explanation:

Answer – C

Azure Cosmos DB acts a document store and is a fully managed service

The Microsoft documentation mentions the following

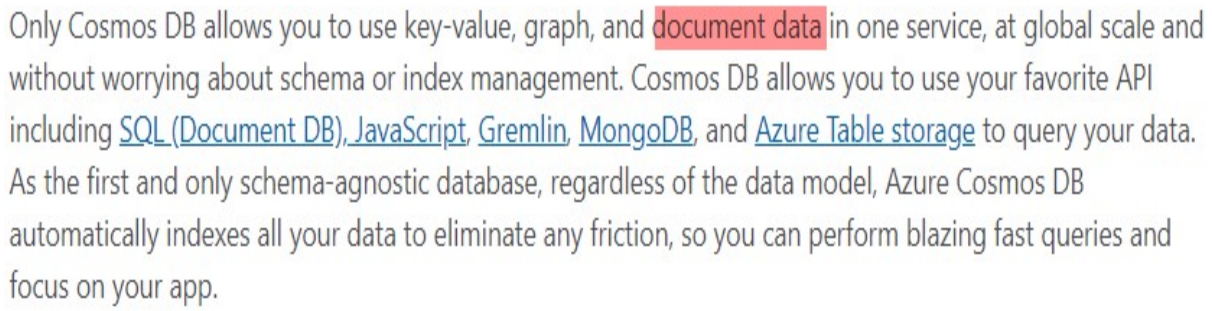
Only Cosmos DB allows you to use key-value, graph, and document data in one service, at global scale and without worrying about schema or index management. Cosmos DB allows you to use your favorite API including SQL (Document DB), JavaScript, Gremlin, MongoDB, and Azure Table storage to query your data. As the first and only schema-agnostic database, regardless of the data model, Azure Cosmos DB automatically indexes all your data to eliminate any friction, so you can perform blazing fast queries and focus on your app.

Since this is clear from the Microsoft documentation, all other options are incorrect

For more information on Azure Cosmos DB, please visit the following URL

https://azure.microsoft.com/es-es/blog/dear-documentdb-customers-welcome-to-azure-cosmos-db/

Question 20

Correct

Domain :Manage and develop data processing

A company wants to create an in-memory batch processing solution. They want to provision an HDInsight cluster for the batch processing solution.
You need to complete the below PowerShell snippet for implementing the solution

Area 1 `

-Name $whizlabclustername -Context $defaultStorageContext

$whizlabclusterSizeInNodes = "4"

$whizlabclusterVersion = "3.6"

$whizlabclusterType = Area 2

$whizlabclusterOS = "Linux"

Area 3 `

-ResourceGroupName "whizlab-rg" `

-ClusterName $whizlabclusterName `

-Location $location `

-ClusterSizeInNodes $whizlabclusterSizeInNodes `

-ClusterType $whizlabclusterType `

-OSType $whizlabclusterOS `

-Version $whizlabclusterVersion `

-HttpCredential $whizlabhttpCredential `

-DefaultStorageAccountName "$defaultStorageAccountName.blob.core.windows.net" `

-DefaultStorageAccountKey $defaultStorageAccountKey `

-DefaultStorageContainer $whizlabclusterName `

-SshCredential $whizlabsshCredentials

Which of the following would go into Area 1?

- A. **New-AzStorageContainer** ✓
- B. **New-AzHDInsightCluster**
- C. **Hadoop**
- D. **Spark**

Explanation:

Answer – A

An example of this is given in the Microsoft documentation where we first have to create a storage container

```
# Default cluster size (# of worker nodes), version, type, and OS
$clusterSizeInNodes = "4"
$clusterVersion = "3.6"
$clusterType = "Hadoop"
$clusterOS = "Linux"
# Set the storage container name to the cluster name
$defaultBlobContainerName = $clusterName

# Create a blob container. This holds the default data store for the cluster.
New-AzStorageContainer `
    -Name $clusterName -Context $defaultStorageContext

# Create the HDInsight cluster
New-AzHDInsightCluster `
    -ResourceGroupName $resourceGroupName `
    -ClusterName $clusterName `
    -Location $location `
    -ClusterSizeInNodes $clusterSizeInNodes `
    -ClusterType $clusterType `
    -OSType $clusterOS `
    -Version $clusterVersion `
    -HttpCredential $httpCredential `
    -DefaultStorageAccountName "$defaultStorageAccountName.blob.core.windows.ne
    -DefaultStorageAccountKey $defaultStorageAccountKey `
    -DefaultStorageContainer $clusterName `
    -SshCredential $sshCredentials
```

Since this is clear from the Microsoft documentation, all other options are incorrect

For more information on an example on creating Azure HDInsight clusters, please visit the following URL

https://docs.microsoft.com/en-us/azure/hdinsight/hdinsight-hadoop-create-linux-clusters-azure-powershell

Ask our Experts

Rate this Question?

View Queries

open

Question 21

Correct

Domain :Manage and develop data processing

A company wants to create an in-memory batch processing solution. They want to provision an HDInsight cluster for the batch processing solution.
You need to complete the below PowerShell snippet for implementing the solution

Area 1 `

-Name $whizlabclustername -Context $defaultStorageContext

$whizlabclusterSizeInNodes = "4"

$whizlabclusterVersion = "3.6"

$whizlabclusterType = Area 2

$whizlabclusterOS = "Linux"

Area 3 `

-ResourceGroupName "whizlab-rg" `

-ClusterName $whizlabclusterName `

-Location $location `

-ClusterSizeInNodes $whizlabclusterSizeInNodes `

-ClusterType $whizlabclusterType `

-OSType $whizlabclusterOS `

-Version $whizlabclusterVersion `

-HttpCredential $whizlabhttpCredential `

-DefaultStorageAccountName "$defaultStorageAccountName.blob.core.windows.net" `

-DefaultStorageAccountKey $defaultStorageAccountKey `

-DefaultStorageContainer $whizlabclusterName `

-SshCredential $whizlabsshCredentials

Which of the following would go into Area 2?

- A. New-AzStorageContainer
- B. New-AzHDInsightCluster
- C. Hadoop
- D. Spark

Explanation:

Answer - D

For the batch processing solution on HDInsight, you use Spark. So we need to mention the cluster type as Spark

The Microsoft documentation mentions the following

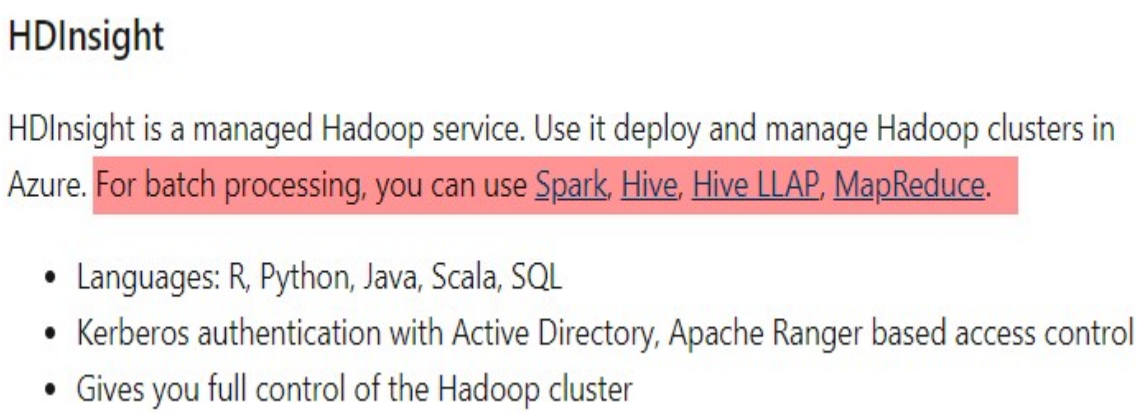

Since this is clear from the Microsoft documentation, all other options are incorrect

For more information on an example on batch processing options, please visit the following URL

https://docs.microsoft.com/en-us/azure/architecture/data-guide/technology-choices/batch-processing

Ask our Experts

Rate this Question?

View Queries

open

Question 22

Correct

Domain :Manage and develop data processing

A company wants to create an in-memory batch processing solution. They want to provision an HDInsight cluster for the batch processing solution.

You need to complete the below PowerShell snippet for implementing the solution

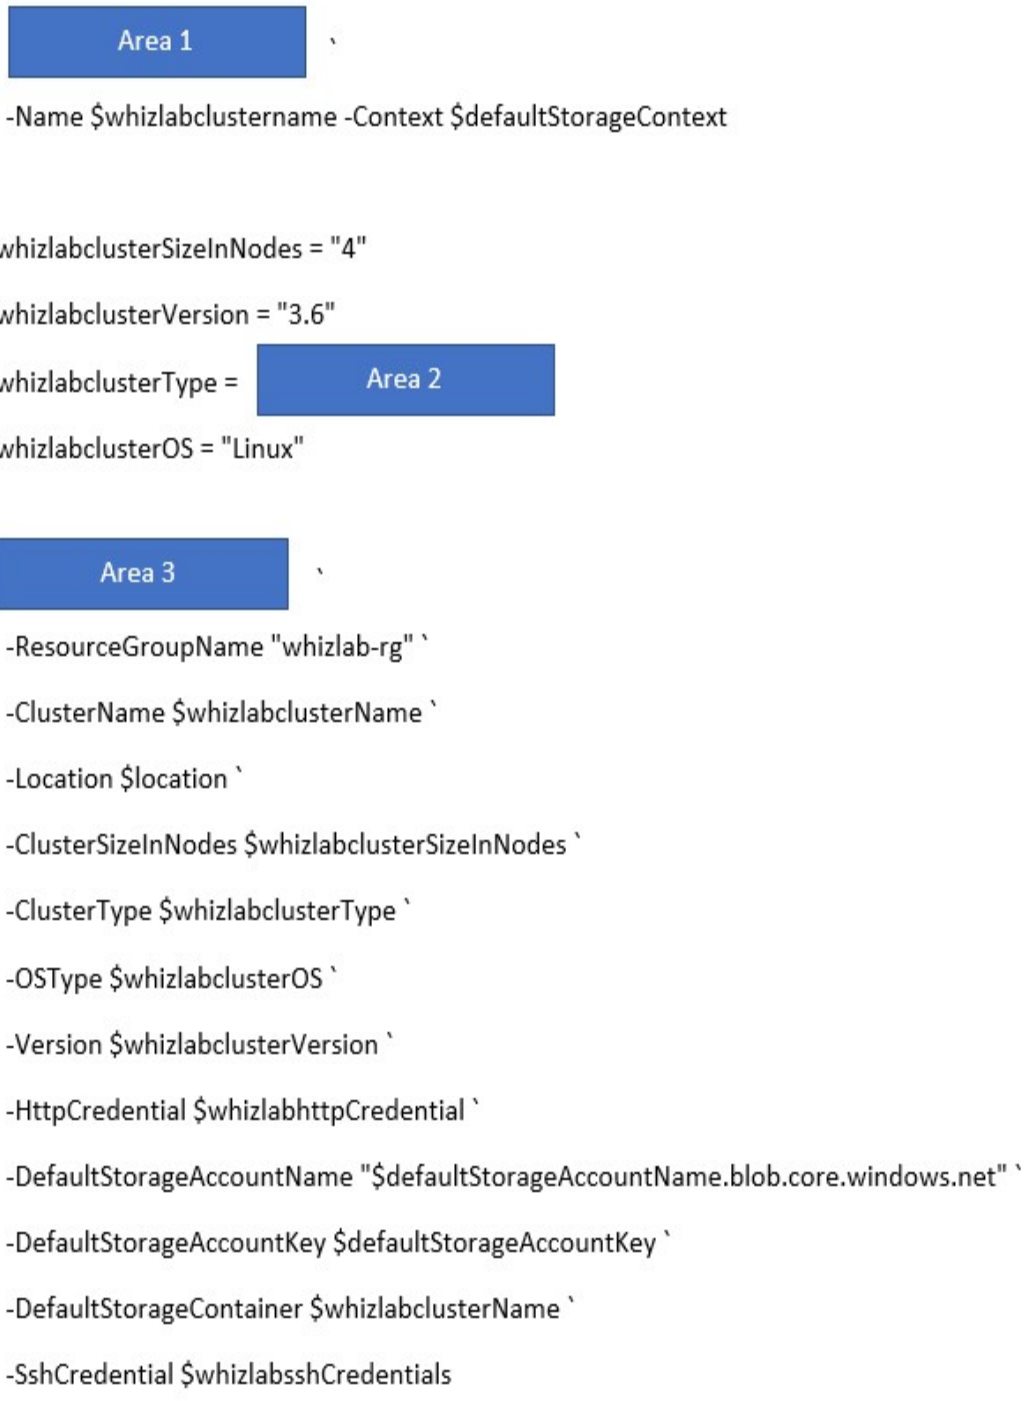

Which of the following would go into Area 3?

A. **New-AzStorageContainer**

B. **New-AzHDInsightCluster**

C. **Hadoop**

D. **Spark**

Explanation:

Answer – B

We finally would go ahead and create the Azure HDInsight cluster.

An example of this is given in the Microsoft documentation

```
# Default cluster size (# of worker nodes), version, type, and OS
$clusterSizeInNodes = "4"
$clusterVersion = "3.6"
$clusterType = "Hadoop"
$clusterOS = "Linux"
# Set the storage container name to the cluster name
$defaultBlobContainerName = $clusterName

# Create a blob container. This holds the default data store for the cluster.
New-AzStorageContainer `
    -Name $clusterName -Context $defaultStorageContext

# Create the HDInsight cluster
New-AzHDInsightCluster `
    -ResourceGroupName $resourceGroupName `
    -ClusterName $clusterName `
    -Location $location `
    -ClusterSizeInNodes $clusterSizeInNodes `
    -ClusterType $clusterType `
    -OSType $clusterOS `
    -Version $clusterVersion `
    -HttpCredential $httpCredential `
    -DefaultStorageAccountName "$defaultStorageAccountName.blob.core.windows.ne
    -DefaultStorageAccountKey $defaultStorageAccountKey `
    -DefaultStorageContainer $clusterName `
    -SshCredential $sshCredentials
```

Since this is clear from the Microsoft documentation, all other options are incorrect

For more information on an example on creating Azure HDInsight clusters, please visit the following URL

https://docs.microsoft.com/en-us/azure/hdinsight/hdinsight-hadoop-create-linux-clusters-azure-powershell

Ask our Experts

Rate this Question?

View Queries

open

Question 23

Correct

Domain :Monitor and optimize data solutions

You have to deploy resources on Azure HD Insight for a batch processing job. The batch processing must run daily and must scale to minimize costs. You also be able to monitor cluster performance.
You need to decide on a tool that will monitor the clusters and provide information on suggestions on how to scale.
You decide to download the Azure HDInsight cluster logs by using Azure PowerShell.
Would this fulfil the requirement?

A. Yes

B. No

Explanation:

Answer – B

This will not give you a complete picture and give you the ability to decide on how to scale the cluster.

You have to use Azure HDInsight cluster management solutions

For more information on HDInsight cluster management solutions, please visit the following URL

https://docs.microsoft.com/en-us/azure/hdinsight/hdinsight-hadoop-oms-log-analytics-tutorial#install-hdinsight-cluster-management-solutions

Ask our Experts

Rate this Question?

View Queries

open

Question 24

Correct

Domain :Monitor and optimize data solutions

You have to deploy resources on Azure HD Insight for a batch processing job. The batch processing must run daily and must scale to minimize costs. You also be able to monitor cluster performance.
You need to decide on a tool that will monitor the clusters and provide information on suggestions on how to scale.
You decide on using Azure Log Analytics and HDInsight cluster management solutions.
Would this fulfil the requirement?

A. Yes

B. No

Explanation:

Answer – A

Yes, you can use Azure Log Analytics along with cluster management solutions. The Microsoft documentation mentions the following

Install HDInsight cluster management solutions

HDInsight provides cluster-specific management solutions that you can add for Azure Monitor logs. Management solutions add functionality to Azure Monitor logs, providing additional data and analysis tools. These solutions collect important performance metrics from your HDInsight clusters and provide the tools to search the metrics. These solutions also provide visualizations and dashboards for most cluster types supported in HDInsight. By using the metrics that you collect with the solution, you can create custom monitoring rules and alerts.

These are the available HDInsight solutions:

- HDInsight Hadoop Monitoring
- HDInsight HBase Monitoring
- HDInsight Interactive Query Monitoring
- HDInsight Kafka Monitoring
- HDInsight Spark Monitoring
- HDInsight Storm Monitoring

For more information on HDInsight cluster management solutions, please visit the following URL

https://docs.microsoft.com/en-us/azure/hdinsight/hdinsight-hadoop-oms-log-analytics-tutorial#install-hdinsight-cluster-management-solutions

Question 25

Correct

Domain :Monitor and optimize data solutions

You have to deploy resources on Azure HD Insight for a batch processing job. The batch processing must run daily and must scale to minimize costs. You also be able to monitor cluster performance.
You need to decide on a tool that will monitor the clusters and provide information on suggestions on how to scale.
You decide on monitoring the cluster load by using the Ambari Web UI
Would this fulfil the requirement?

- A. Yes ✓
- B. No

Explanation:

Answer – A

Yes, this will give you a good idea on the load on the Azure HD Insight cluster

The Microsoft documentation mentions the following

Monitor cluster load

Hadoop clusters can deliver the most optimal performance when the load on cluster is evenly distributed across all the nodes. This enables the processing tasks to run without being constrained by RAM, CPU, or disk resources on individual nodes.

To get a high-level look at the nodes of your cluster and their loading, sign in to the Ambari Web UI, then select the **Hosts** tab. Your hosts are listed by their fully qualified domain names. Each host's operating status is shown by a colored health indicator:

Color	Description
Red	At least one master component on the host is down. Hover to see a tooltip that lists affected components.
Orange	At least one secondary component on the host is down. Hover to see a tooltip that lists affected components.
Yellow	Ambari Server has not received a heartbeat from the host for more than 3 minutes.
Green	Normal running state.

For more information on using the Ambari UI, please visit the following URL

https://docs.microsoft.com/en-us/azure/hdinsight/hdinsight-key-scenarios-to-monitor

Question 26

Correct

Domain :Manage and develop data processing

A company wants to develop a solution in Azure. The solution would handle streaming data from Twitter. Azure Event Hubs would be used to ingest the streaming data. Azure Data bricks would then be used to receive the data from Event Hubs. Which of the following actions would you implement for this requirement? Choose 3 answers from the options given below

- [x] A. **Create and configure a Notebook that would be used to consume the streaming data** ✓
- [ ] B. **Use Environment variables that would be used to define the Apache Spark connection**
- [ ] C. **Configure an ODBC or JDBC connector**
- [x] D. **Deploy the Azure Databricks service** ✓
- [x] E. **Deploy a Spark cluster and also attach the required libraries** ✓

Explanation:

Answer – A,D and E

There is a tutorial in the Microsoft documentation that shows how to stream data from Azure Event Hubs to Azure Databricks

You first need to deploy the Azure Databricks service.

Then deploy the Spark cluster

And then finally configure a notebook in the databricks service

The Microsoft documentation mentions the following

The following illustration shows the application flow:

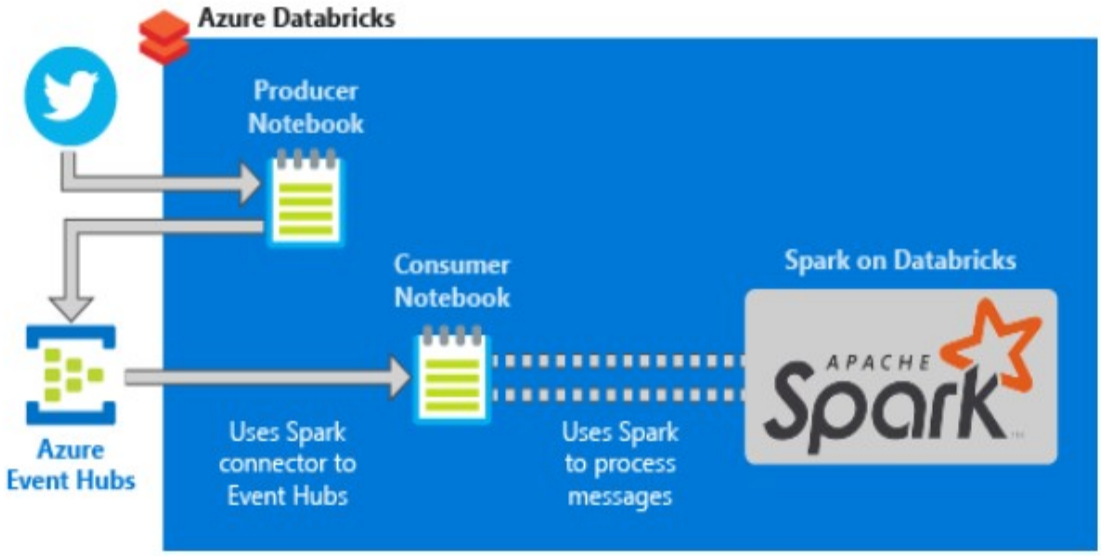

This tutorial covers the following tasks:

- ✓ Create an Azure Databricks workspace
- ✓ Create a Spark cluster in Azure Databricks
- ✓ Create a Twitter app to access streaming data
- ✓ Create notebooks in Azure Databricks
- ✓ Attach libraries for Event Hubs and Twitter API
- ✓ Send tweets to Event Hubs
- ✓ Read tweets from Event Hubs

Since this is clear from the Microsoft documentation, all other options are incorrect

For more information on streaming data from Event Hubs, please visit the following URL

https://docs.microsoft.com/en-us/azure/azure-databricks/databricks-stream-from-eventhubs

Question 27 **Correct**

Domain :Manage and develop data processing

A company has an on-premise Microsoft SQL Server database. They want to copy the data from the instance to Azure Blob storage. They want to configure Azure Data Factory to connect to the on-premise SQL Server instance. Which of the following steps must be carried out for this requirement? Choose 3 answers from the options given below

- [x] A. **Configure a linked service that would connect to the SQL Server Instance**
- [x] B. **On the on-premise network, install and configure a self-hosted runtime**
- [x] C. **Deploy an instance of Azure Data Factory**
- [ ] D. **Create a backup of the SQL Server database**
- [ ] E. **In the SQL Server database, create the database master key**

Explanation:

Answer – A,B and C

First you deploy an instance of Azure Data Factory

Then install and configure a self-hosted runtime

And then finally configure a linked service

You can create a self-hosted runtime in your on-premise network. This can be used for copy activities between a cloud data store and a data store in your private network.

The Microsoft documentation mentions the following

Create and configure a self-hosted integration runtime

06/18/2019 • 22 minutes to read • +14

The integration runtime (IR) is the compute infrastructure that Azure Data Factory uses to provide data-integration capabilities across different network environments. For details about IR, see Integration runtime overview.

A self-hosted integration runtime can run copy activities between a cloud data store and a data store in a private network. It also can dispatch transform activities against compute resources in an on-premises network or an Azure virtual network. The installation of a self-hosted integration runtime needs an on-premises machine or a virtual machine inside a private network.

Option D is incorrect since you don't need to create a backup when using Azure Data Factory

Option E is incorrect since you don't need a master key for the copy operation

For more information on the self-hosted integration runtime, please visit the following URL

https://docs.microsoft.com/en-us/azure/data-factory/create-self-hosted-integration-runtime

Ask our Experts

Rate this Question?

View Queries

open

Question 28

Correct

Domain :Manage and develop data processing

A company wants to make use of Azure Stream Analytics. The incoming data is formatted as one record per row. You need to complete the following REST API segment which would be used to create the input stream

Which of the following would go into Area 1?

- A. **“type” : “CSV”** ✓
- B. **“type” : “Avro”**

C. **"type" : "JSON"**

D. **o**

Explanation:

Answer – A

Since the records are going to sent one per row and we have a field delimiter, the input format would be CSV.

For the format of the REST API, there is also an example in the Microsoft documentation

```
{
    "name":"MyBlobSource",
    "properties":{
       "type":"reference",
       "serialization":{
          "type":"CSV",
          "properties":{
             "fieldDelimiter":",",
             "encoding":"UTF8"
          }
       },
```

All other options are incorrect since they would not be the ideal match for the format of data being sent.

For more information on Stream Analytics jobs, please visit the following URL

https://docs.microsoft.com/en-us/rest/api/streamanalytics/stream-analytics-job

A company wants to make use of Azure Stream Analytics. The incoming data is formatted as one record per row. You need to complete the following REST API segment which would be used to create the input stream

```
"inputs":[
    {
      "name":"whizlabsource",
      "properties":{
        "type":"stream",
        "serialization":{
          Area 1 ,
          "properties":{
           “fieldDelimiter” :”,”,
            "encoding":"UTF8"
          }
        },
        "datasource":{
        Area 2 ,
          "properties":{
            "serviceBusNamespace":"whizlabbus2020",
            "sharedAccessPolicyName":"whizlabpolicy",
            "sharedAccessPolicyKey":"***/**********/******************************",
            "eventHubName":"whizlabhub"
          }
        }
      }
    }
```

Which of the following would go into Area 2?

- A. **“type” : “Microsoft.Storage/Blob”**

B. **"type" : "Microsoft.ServiceBus/EventHub"**

C. **"type" : "Microsoft.Devices/IoTHubs"**

D. **o**

Explanation:

Answer – B

Here since there is a mention of a service namespace and hub name, we need to use the type of "Microsoft.ServiceBus/EventHub"

Since this is the ideal type to use, all other options are incorrect

For more information on Stream Analytics jobs, please visit the following URL

https://docs.microsoft.com/en-us/rest/api/streamanalytics/stream-analytics-job

Ask our Experts

Rate this Question?

View Queries

open

Question 30

Correct

Domain :Manage and develop data processing

A company wants to implement a solution using Azure Stream Analytics. Below are the key requirements of the solution

Ingest data from Blob storage

Be able to analyze data in real time

Be able to store the processed data in Azure Cosmos DB

Which of the following actions would you implement for this requirement? Choose 3 answers from the options given below

A. **Setup Cosmos DB as the output**

- [ ] B. **Setup Blob Storage as the output**
- [x] C. **Configure Blob Storage as the input**
- [ ] D. **Create a query statement with the ORDER by clause**
- [x] E. **Create a query statement with the SELECT INTO statement**

Explanation:

Answer – A, C and E

First, we need to set Blob storage as the input in Azure Stream Analytics. Then set Cosmos DB as the output. And then finally have a query statement based on the SELECT INTO statement.

Option B is incorrect since Blob storage should be set as the input

Option D is incorrect since the query should be based on the SELECT INTO statement

For more information on Stream analytics for Cosmos DB, please visit the following URL

https://docs.microsoft.com/en-us/azure/stream-analytics/stream-analytics-documentdb-output

Ask our Experts

Rate this Question?

View Queries

open

Question 31

Correct

Domain :Implement data storage solutions

You are going to create an Azure Databricks environment. You will be accessing data in an Azure Blob storage account. The data must be available to all Azure Databricks workspaces. Which of the following actions would you perform for this requirement? Choose 3 answers from the options given below

- [ ] A. **Make sure to upload a certificate**
- [x] B. **Ensure to add secrets to the scope**
- [ ] C. **Ensure to use Blob storage access keys**
- [x] D. **Create a secret scope**
- [x] E. **Mount the Azure Blob storage container**

Explanation:

Answer – B,D and E

If you want to ensure blob storage access is available across all workspaces, then you need to setup the access keys in the scope accordingly

The Microsoft documentation mentions the following

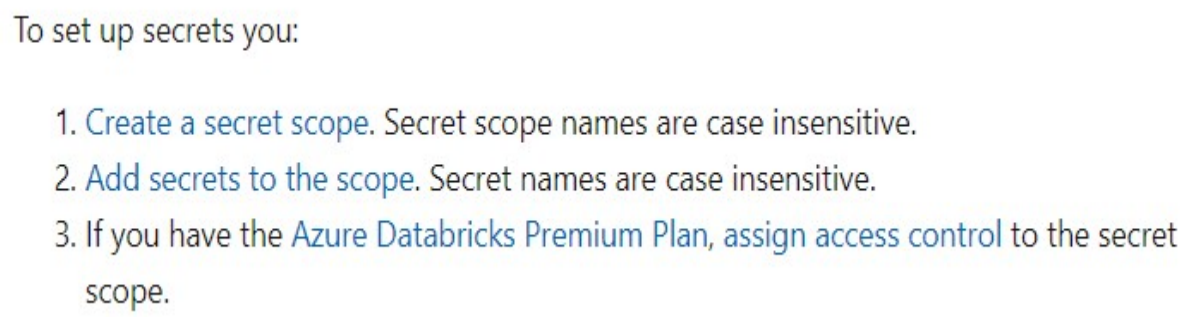

You then need to mount the Azure Blob storage container

The Microsoft documentation mentions the following

Mount an Azure Blob storage container

You can mount Blob storage containers using Databricks Runtime 4.0 or higher. Once a Blob storage container is mounted, you can use Databricks Runtime 3.4 or higher to access the mount point.

1. To mount a Blob storage container or a folder inside a container, use the following command:

Scala

Option A is incorrect because you don't need a certificate to access Blob storage

Option C is incorrect because you need to use secrets at the scope level to ensure all workspaces can use the storage accounts

For more information on using Azure Blob storage in Azure Databricks, please visit the following URL

https://docs.microsoft.com/en-us/azure/databricks/data/data-sources/azure/azure-storage

Ask our Experts

Rate this Question?

View Queries

open

Question 32 **Correct**

Domain :Manage and develop data processing

A company wants to develop a data processing system. The system should meet the following requirements

The solution must scale up and down to ensure cost is effectively managed

It should use an in-memory data processing engine to speed up the machine learning operations

The system should provide streaming capabilities

The system should provide the ability to code in SQL, Python and Scala

Which of the following would you consider implementing for this solution?

- A. **HDInsight Spark Cluster** ✓
- B. **Azure Stream Analytics**
- C. **HDInsight Hadoop Cluster**
- D. **Azure SQL Data warehouse**

Explanation:

Answer – A

You can use Spark Clusters to carry out in-memory processing

The Microsoft documentation mentions the following

What is Apache Spark in Azure HDInsight

10/01/2019 • 6 minutes to read • +4

Apache Spark is a parallel processing framework that supports in-memory processing to boost the performance of big-data analytic applications. Apache Spark in Azure HDInsight is the Microsoft implementation of Apache Spark in the cloud. HDInsight makes it easier to create and configure a Spark cluster in Azure. Spark clusters in HDInsight are compatible with Azure Storage and Azure Data Lake Storage. So you can use HDInsight Spark clusters to process your data stored in Azure. For the components and the versioning information, see Apache Hadoop components and versions in Azure HDInsight.

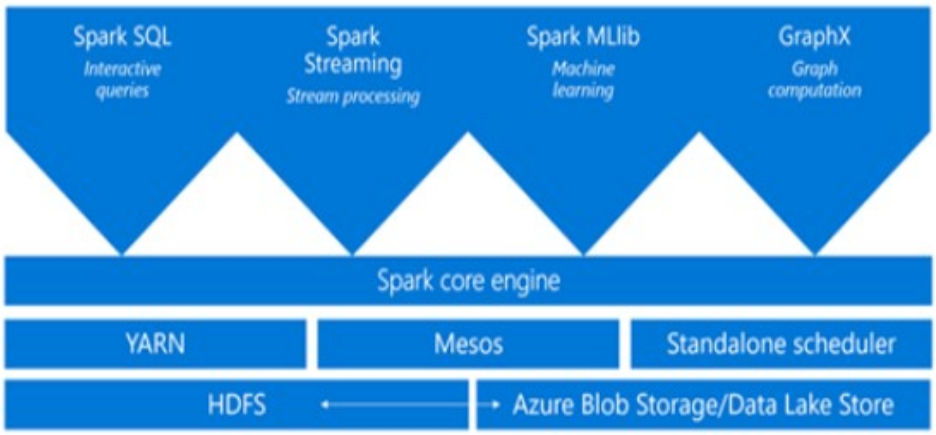

Option B is incorrect since this is a fully managed service designed to help you analyze and process fast moving streams of data

Option C is incorrect since you need to use Spark to provide streaming capabilities

Option D is incorrect since this is more of a data warehousing solution

For more information on using HDInsight Spark, please visit the following URL

https://docs.microsoft.com/en-us/azure/hdinsight/spark/apache-spark-overview

Ask our Experts

Rate this Question?

View Queries

open

Question 33

Correct

Domain :Manage and develop data processing

A company is planning on using Azure SQL Database along with Elastic Database jobs. You have to analyze, troubleshoot and reports on the various components that are responsible for running Elastic database jobs.
Which of the following would be used for the task of storing "Execution results and diagnostics"?

- A. **Control database**
- B. **Azure Service Bus**
- C. **Azure Storage**
- D. **Job Service**

Explanation:

Answer – A

The Job or control database is used for storing the Execution results and diagnostics

The Microsoft documentation mentions the following

Job database

The *Job database* is used for defining jobs and tracking the status and history of job executions. The *Job database* is also used to store agent metadata, logs, results, job definitions, and also contains many useful stored procedures and other database objects for creating, running, and managing jobs using T-SQL.

Since this is clear from the Microsoft documentation, all other options are incorrect

For more information on SQL job automation, please visit the following URL

https://docs.microsoft.com/en-us/azure/sql-database/sql-database-job-automation-overview

Ask our Experts

Rate this Question?

View Queries

open

Question 34

Correct

Domain :Manage and develop data processing

A company is planning on using Azure SQL Database along with Elastic Database jobs. You have to analyze, troubleshoot and reports on the various components that are responsible for running Elastic database jobs.
Which of the following would be used for the task of handling the "Job launcher and tracker"?

A. **Control database**

B. **Azure Service Bus**

C. **Azure Storage**

D. **Job Service**

Explanation:

Answer – D

The Job service or elastic job agent is responsible for launching and tracking the job

The Microsoft documentation mentions the following

Elastic Job agent

An Elastic Job agent is the Azure resource for creating, running, and managing jobs. The Elastic Job agent is an Azure resource you create in the portal (PowerShell and REST are also supported).

Since this is clear from the Microsoft documentation, all other options are incorrect

For more information on SQL job automation, please visit the following URL

https://docs.microsoft.com/en-us/azure/sql-database/sql-database-job-automation-overview

Ask our Experts

Rate this Question?

View Queries

open

Question 35

Correct

Domain :Manage and develop data processing

A company is planning on using Azure SQL Database along with Elastic Database jobs. You have to analyze, troubleshoot and reports on the various components that are responsible for running Elastic database jobs.
Which of the following would be used for the task of storing the "Job metadata and state"?

- A. **Control database**
- B. **Azure Service Bus**
- C. **Azure Storage**
- D. **Job Service**

Explanation:

Answer – A

The job state and metadata are stored in the Control or Job database

The Microsoft documentation mentions the following

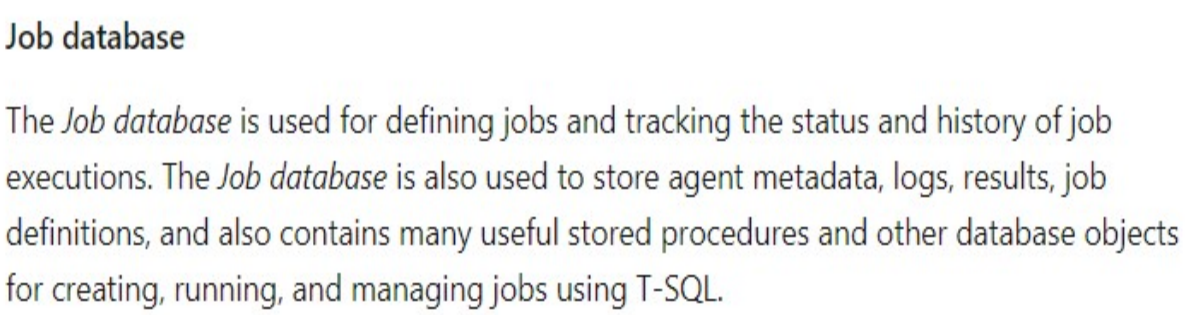

Job database

The *Job database* is used for defining jobs and tracking the status and history of job executions. The *Job database* is also used to store agent metadata, logs, results, job definitions, and also contains many useful stored procedures and other database objects for creating, running, and managing jobs using T-SQL.

Since this is clear from the Microsoft documentation, all other options are incorrect

For more information on SQL job automation, please visit the following URL

https://docs.microsoft.com/en-us/azure/sql-database/sql-database-job-automation-overview

Ask our Experts

Rate this Question?

View Queries

open

Question 36

Correct

Domain :Implement data storage solutions

A company wants to set up a NoSQL database in Azure to store data. They want to have a database that can be used to store key-value pairs. They also want to have a database that can store wide-column data values. Which of the following API types would you choose for Cosmos DB for these requirements? Choose 2 answers from the options given below

- [x] A. **Table API**
- [ ] B. **MongoDB API**

- [ ] C. **Gremlin API**
- [ ] D. **SQL API**
- [x] E. **Cassandra API**

Explanation:

Answer – A and E

The Table API can be used to store key-value pairs

The Microsoft documentation mentions the following

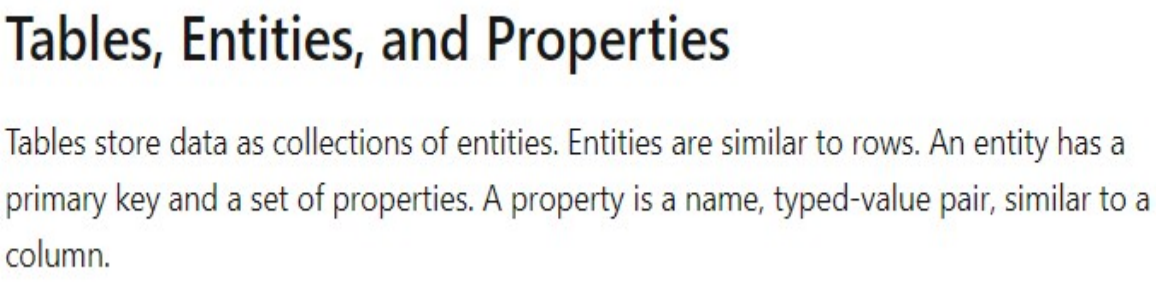

The Cassandra database type can be used to store wide-column data values

Option B is incorrect since this is a document-based API

Option C is incorrect since this is a graph-based database

Option D is incorrect since this is used to store JSON based data

For more information on the Table and Cassandra API, please visit the following URL

https://docs.microsoft.com/en-us/azure/cosmos-db/cassandra-introduction

https://docs.microsoft.com/en-us/azure/cosmos-db/table-introduction

Ask our Experts

Rate this Question?

Question 37

Correct

Domain :Monitor and optimize data solutions

A company wants to deploy a sales application as part of their application portfolio. The application needs to store its data in an Azure SQL Database. The data for sales will be stored in Azure SQL database from multiple regions. After a week, the sales data needs to be stored in another Azure SQL database to perform analytics. The Analytics is a very resource intensive process and can generate up to 40 TB of data. You have to provision the right type of database for each operation. The database provisioning must ensure performance is maximized and cost is minimized.
Which of the following would you choose as the database instance type for uploading daily sales data?

- A. **Azure SQL Database elastic pools** ✓
- B. **Azure SQL Database Premium**
- C. **Azure SQL Database Managed Instance**
- D. **Azure SQL Database Hyperscale**

Explanation:

Answer – A

For daily use and optimization of cost, you can opt to use Azure SQL Database elastic pools

The Microsoft documentation mentions the following

Elastic pools help you manage and scale multiple Azure SQL databases

08/06/2019 • 13 minutes to read • +16

SQL Database elastic pools are a simple, cost-effective solution for managing and scaling multiple databases that have varying and unpredictable usage demands. The databases in an elastic pool are on a single Azure SQL Database server and share a set number of resources at a set price. Elastic pools in Azure SQL Database enable SaaS developers to optimize the price performance for a group of databases within a prescribed budget while delivering performance elasticity for each database.

Option B is incorrect since this should ideally be chosen when you want better capabilities for your database

Option C is incorrect since this should be used when you want an easy way to migrate databases onto Azure and when you don't want to manage the underlying infrastructure.

Option D is incorrect since this should be used when you want higher performance for your database.

For more information on Azure SQL Database Elastic Pools, please visit the following URL

https://docs.microsoft.com/en-us/azure/sql-database/sql-database-elastic-pool

Question 38 **Correct**

Domain :Monitor and optimize data solutions

A company wants to deploy a sales application as part of their application portfolio. The application needs to store its data in an Azure SQL Database. The data for sales will be stored in Azure SQL database from multiple regions. After a week, the sales data needs to be stored in another Azure SQL database to perform analytics. The Analytics is a very resource intensive process and can generate up to 40 TB of data. You have to provision the right type of database for each operation. The database provisioning must ensure performance is maximized and cost is minimized.
Which of the following would you choose as the database instance type for the weekly sales data on which analytics needs to be performed?

- A. **Azure SQL Database elastic pools**
- B. **Azure SQL Database Premium**
- C. **Azure SQL Database Managed Instance**
- D. **Azure SQL Database Hyperscale** ✓

Explanation:

Answer – D

If you need high storage capabilities and faster performance, then you should choose Azure SQL Hyperscale

The Microsoft documentation mentions the following

What are the Hyperscale capabilities

The Hyperscale service tier in Azure SQL Database provides the following additional capabilities:

- Support for up to 100 TB of database size
- Nearly instantaneous database backups (based on file snapshots stored in Azure Blob storage) regardless of size with no IO impact on compute resources
- Fast database point-in-time restores (based on file snapshots) in minutes rather than hours or days (not a size of data operation)
- Higher overall performance due to higher log throughput and faster transaction commit times regardless of data volumes
- Rapid scale out - you can provision one or more read-only nodes for offloading your read workload and for use as hot-standbys
- Rapid Scale up - you can, in constant time, scale up your compute resources to accommodate heavy workloads as and when needed, and then scale the compute resources back down when not needed.

Option A is incorrect since this should ideally be chosen for databases which have unpredictable workloads and you want to have a cost optimized approach for hosting databases

Option B is incorrect since this should ideally be chosen when you want better capabilities for your database

Option C is incorrect since this should be used when you want an easy way to migrate databases onto Azure and when you don't want to manage the underlying infrastructure.

For more information on Azure SQL Database Hyperscale, please visit the following URL

https://docs.microsoft.com/en-us/azure/sql-database/sql-database-service-tier-hyperscale

Ask our Experts

Rate this Question?

View Queries

open

Question 39

Correct

Domain :Implement data storage solutions

A company wants to synchronize data from an on-premise Microsoft SQL Server database to an Azure SQL Database. You have to perform an assessment to understand whether the data can actually be moved without any sort of compatibility issues. Which of the following would you use to perform the assessment?

A. **Azure SQL Data Sync**

B. **SQL Server Migration Assistant**

C. **Data Migration Assistant**

D. **Microsoft Assessment and Planning Toolkit**

Explanation:

Answer – C

You can use the Data Migration Assistant for this requirement

The Microsoft documentation mentions the following

Overview of Data Migration Assistant

11/05/2019 • 2 minutes to read • +6

The Data Migration Assistant (DMA) helps you upgrade to a modern data platform by detecting compatibility issues that can impact database functionality in your new version of SQL Server or Azure SQL Database. DMA recommends performance and reliability improvements for your target environment and allows you to move your schema, data, and uncontained objects from your source server to your target server.

Since this is clearly mentioned in the documentation, all other options are incorrect

For more information on Data Migration Assistant, please visit the following URL

https://docs.microsoft.com/en-us/sql/dma/dma-overview?view=sql-server-ver15

Ask our Experts

Rate this Question?

View Queries

open

Question 40

Correct

Domain :Implement data storage solutions

A company wants to start using Azure Cosmos DB. They would be using the Cassandra API for the database.
Which of the following would they need to choose as the container type?

- A. **table**
- B. **collection**

C. graph

D. entities

E. rows

Explanation:

Answer – A

The Microsoft documentation gives the different types of containers when it comes to the different API's

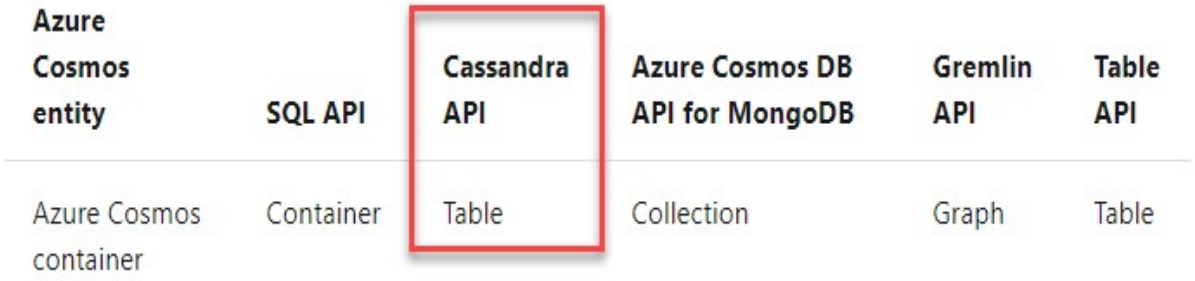

Azure Cosmos entity	SQL API	Cassandra API	Azure Cosmos DB API for MongoDB	Gremlin API	Table API
Azure Cosmos container	Container	Table	Collection	Graph	Table

Since this is clearly mentioned in the documentation, all other options are incorrect

For more information on Azure Cosmos DB container and items, please visit the following URL

https://docs.microsoft.com/en-us/azure/cosmos-db/databases-containers-items

Ask our Experts

Rate this Question?

View Queries

open

Question 41

Correct

Domain :Implement data storage solutions

A company wants to start using Azure Cosmos DB. They would be using the Cassandra API for the database.
Which of the following would they need to choose as the item type?

- A. **table**
- B. **collection**
- C. **graph**
- D. **entities**
- E. **rows** ✓

Explanation:

Answer – E

The Microsoft documentation gives the different types of items when it comes to the different API's

Azure Cosmos items

Depending on which API you use, an Azure Cosmos item can represent either a document in a collection, a row in a table, or a node or edge in a graph. The following table shows the mapping of API-specific entities to an Azure Cosmos item:

Cosmos entity	SQL API	Cassandra API	Azure Cosmos DB API for MongoDB	Gremlin API	Table API
Azure Cosmos item	Document	Row	Document	Node or edge	Item

Since this is clearly mentioned in the documentation, all other options are incorrect

For more information on Azure Cosmos DB container and items, please visit the following URL

https://docs.microsoft.com/en-us/azure/cosmos-db/databases-containers-items

Ask our Experts

Rate this Question?

View Queries

open

Question 42

Correct

Domain :Implement data storage solutions

A company is planning on setting up an Azure SQL data warehouse. They want to setup different tables. The different tables have different requirements as stated below

whizlab_sales – Here the rows should be distributed in such a way that it offers high performance

whizlab_offers – Here data should be available on all nodes to achieve better performance on table joins

whizlab_orders – Here data should be loaded faster on the underlying table

Which of the following table type would you use for the table whizlab_sales?

- A. **Hash-distributed tables**
- B. **Primary tables**
- C. **Replicated tables**
- D. **Round-robin tables**

Explanation:

Answer – A

Here you would use Hash-distributed tables

The Microsoft documentation mentions the following

Hash-distributed tables

A hash distributed table distributes rows based on the value in the distribution column. A hash distributed table is designed to achieve high performance for queries on large tables. There are several factors to consider when choosing a distribution column.

Since this is clearly mentioned in the documentation, all other options are incorrect

For more information on Azure SQL data warehouse tables, please visit the following URL

https://docs.microsoft.com/en-us/azure/sql-data-warehouse/sql-data-warehouse-tables-overview

Ask our Experts

Rate this Question?

View Queries

open

Question 43

Correct

Domain :Implement data storage solutions

A company is planning on setting up an Azure SQL data warehouse. They want to setup different tables. The different tables have different requirements as stated below

whizlab_sales – Here the rows should be distributed in such a way that it offers high performance

whizlab_offers – Here data should be available on all nodes to achieve better performance on table joins

whizlab_orders – Here data should be loaded faster on the underlying table

Which of the following table type would you use for the table whizlab_offers?

A. **Hash-distributed tables**

B. **Primary tables**

C. **Replicated tables**

D. **Round-robin tables**

Explanation:

Answer – C

Here you would use Replicated tables

The Microsoft documentation mentions the following

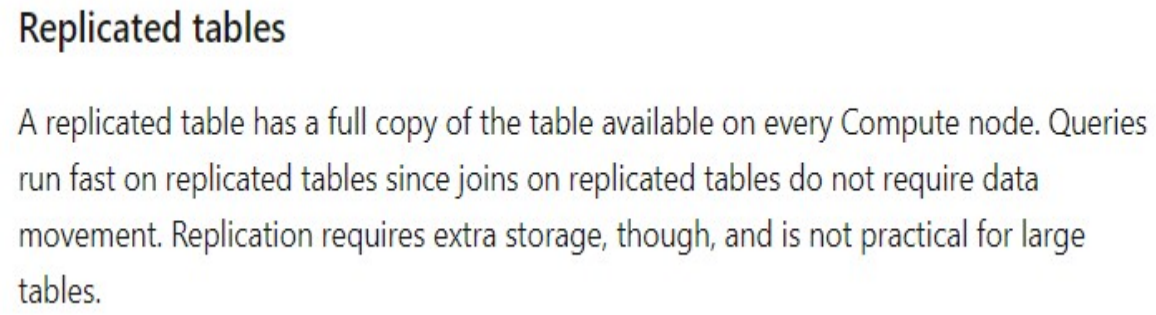

Since this is clearly mentioned in the documentation, all other options are incorrect

For more information on Azure SQL data warehouse tables, please visit the following URL

https://docs.microsoft.com/en-us/azure/sql-data-warehouse/sql-data-warehouse-tables-overview

Ask our Experts

Rate this Question?

View Queries

open

Question 44

Correct

Domain :Implement data storage solutions

A company is planning on setting up an Azure SQL data warehouse. They want to setup different tables. The different tables have different requirements as stated below

whizlab_sales – Here the rows should be distributed in such a way that it offers high performance

whizlab_offers – Here data should be available on all nodes to achieve better performance on table joins

whizlab_orders – Here data should be loaded faster on the underlying table
Which of the following table type would you use for the table whizlab_orders?

- A. **Hash-distributed tables**
- B. **Primary tables**
- C. **Replicated tables**
- D. **Round-robin tables** ✓

Explanation:

Answer – D

Here you would use Round-robin tables

The Microsoft documentation mentions the following

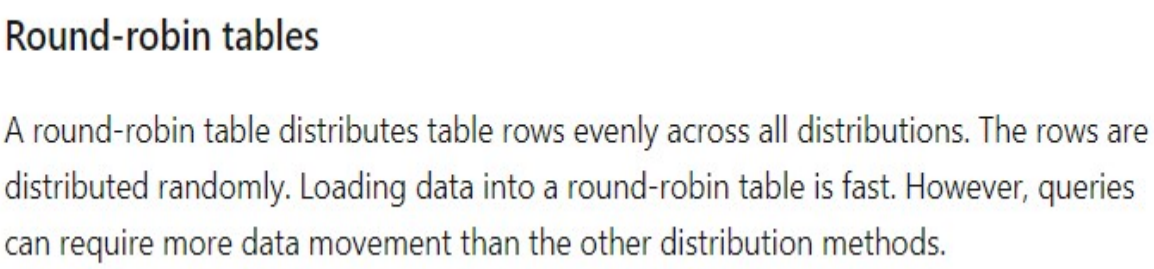

Since this is clearly mentioned in the documentation, all other options are incorrect

For more information on Azure SQL data warehouse tables, please visit the following URL

https://docs.microsoft.com/en-us/azure/sql-data-warehouse/sql-data-warehouse-tables-overview

Ask our Experts

Rate this Question?

View Queries

open

Question 45

Correct

Domain :Implement data storage solutions

A company wants to migrate several on-premise Microsoft SQL Server databases to Azure. They want to migrate to Azure using the backup process available for Microsoft SQL servers. Which of the following is the data technology they should use on Azure?

- A. **Azure SQL Database Managed Instance**
- B. **Azure SQL Data warehouse**
- C. **Azure Cosmos DB**
- D. **Azure SQL Database single database**

Explanation:

Answer – A

If you are looking at migrating on-premise SQL databases, then the best option is to use Azure SQL Database Managed Instance

The Microsoft documentation mentions the following

What is Azure SQL Database managed instance?

11/27/2019 • 20 min. za čitanje • +11

Managed instance is a new deployment option of Azure SQL Database, providing near 100% compatibility with the latest SQL Server on-premises (Enterprise Edition) Database Engine, providing a native virtual network (VNet) implementation that addresses common security concerns, and a business model favorable for on-premises SQL Server customers. The managed instance deployment model allows existing SQL Server customers to lift and shift their on-premises applications to the cloud with minimal application and database changes. At the same time, the managed instance deployment option preserves all PaaS capabilities (automatic patching and version updates, automated backups, high-availability), that drastically reduces management overhead and TCO.

Database migration

The managed instance deployment option targets user scenarios with mass database migration from on-premises or IaaS database implementations. Managed instance supports several database migration options:

Back up and restore

The migration approach leverages SQL backups to Azure Blob storage. Backups stored in Azure storage blob can be directly restored into a managed instance using the T-SQL RESTORE command.

Option B is incorrect since this is a data warehousing solution

Option C is incorrect since this is a NoSQL database

Option D is incorrect since using Azure SQL Database Managed Instance is a better option when you want to migrate on-premise Microsoft SQL Server databases

For more information on Azure SQL database managed instance, please visit the following URL

https://docs.microsoft.com/en/azure/sql-database/sql-database-managed-instance

Ask our Experts

Rate this Question?

View Queries

open

Question 46

Correct

Domain :Implement data storage solutions

A company wants to create an Azure Cosmos DB account. They need to ensure the database uses the SQL API and the latency is minimized.

You have to complete the following Azure CLI command for this requirement

az cosmosdb create --name "whizlabaccount2020" --resource-group "whizlab-rg"

--kind Area 1 --default-consistency-level Area 2

Which of the following would go into Area 1?

A. **MongoDB**

B. **GlobalDocumentDB**

C. **SQL**

D. **JSON**

Explanation:

Answer – B

The value to specify when you want the SQL API is GlobalDocumentDB

This is the only option viable for creating the SQL API account, hence all other options are incorrect

For more information on creating an Azure Cosmos SQL API account, please visit the following URL

https://docs.microsoft.com/en-us/azure/cosmos-db/scripts/cli/sql/create

Ask our Experts

Rate this Question?

View Queries

open

Question 47

Correct

Domain :Implement data storage solutions

A company wants to create an Azure Cosmos DB account. They need to ensure the database uses the SQL API and the latency is minimized.

You have to complete the following Azure CLI command for this requirement

az cosmosdb create --name "whizlabaccount2020" --resource-group "whizlab-rg"

--kind Area 1 --default-consistency-level Area 2

Which of the following would go into Area 2?

A. **Strong**

B. **Session**

C. **Eventual**

D. **Bounded staleness**

Explanation:

Answer – C

If you look at the Microsoft documentation and see the consistency levels, the Eventual consistency level will have the least latency.

With Azure Cosmos DB, developers can choose from five well-defined consistency models on the consistency spectrum. From strongest to more relaxed, the models include *strong*, *bounded staleness*, *session*, *consistent prefix*, and *eventual* consistency. The models are well-defined and intuitive and can be used for specific real-world scenarios. Each model provides availability and performance tradeoffs and is backed b the SLAs. The following image shows the different consistency levels as a spectrum.

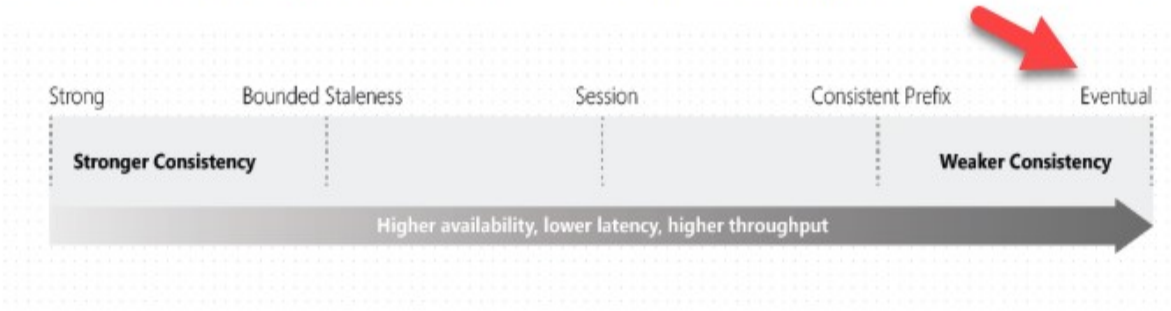

The consistency levels are region-agnostic and are guaranteed for all operations regardless of the region from which the reads and writes are served, the number of regions associated with your Azure Cosmos account, or whether your account is configured with a single or multiple write regions.

Since this is evident from the Microsoft documentation, all other options are incorrect

For more information on Azure Cosmos DB consistency levels, please visit the following URL

https://docs.microsoft.com/en-us/azure/cosmos-db/consistency-levels

Question 48 Correct

Domain :Implement data storage solutions

View Case Study

You have to implement the deployment of Azure Data Factory pipelines. Which of the following would you use for authorization of the deployments?

- A. RBAC
- B. MAC
- C. Claims
- D. Tokens

Explanation:

Answer – A

You can use role-based access control for deployments. The Microsoft documentation mentions the following

Roles and permissions for Azure Data Factory

11/05/2018 • 4 minutes to read • 

This article describes the roles required to create and manage Azure Data Factory resources, and the permissions granted by those roles.

Roles and requirements

To create Data Factory instances, the user account that you use to sign in to Azure must be a member of the *contributor* or *owner* role, or an *administrator* of the Azure subscription. To view the permissions that you have in the subscription, in the Azure portal, select your username in the upper-right corner, and then select **Permissions**. If you have access to multiple subscriptions, select the appropriate subscription.

To create and manage child resources for Data Factory - including datasets, linked services, pipelines, triggers, and integration runtimes - the following requirements are applicable:

- To create and manage child resources in the Azure portal, you must belong to the **Data Factory Contributor** role at the resource group level or above.
- To create and manage child resources with PowerShell or the SDK, the **contributor** role at the resource level or above is sufficient.

Since this is evident from the Microsoft documentation, all other options are incorrect

For more information on Azure Data Factory roles and permissions, please visit the following URL

https://docs.microsoft.com/en-us/azure/data-factory/concepts-roles-permissions

Ask our Experts

Rate this Question?

View Queries

open

Question 49

Correct

Domain :Implement data storage solutions

View Case Study

You have to implement the deployment of Azure Data Factory pipelines. Which of the following would you use for authentication of the deployments?

A. **Service Principal**

B. **Kerberos**

C. **Certificate-based**

D. **Bearer Token**

Explanation:

Answer – A

You can use a service principal for deployment purposes

The Microsoft documentation mentions the following

Overview

When creating a data factory, a managed identity can be created along with factory creation. The managed identity is a managed application registered to Azure Active Directory, and represents this specific data factory.

Managed identity for Data Factory benefits the following features:

- Store credential in Azure Key Vault, in which case data factory managed identity is used for Azure Key Vault authentication.
- Connectors including Azure Blob storage, Azure Data Lake Storage Gen1, Azure Data Lake Storage Gen2, Azure SQL Database, and Azure SQL Data Warehouse.
- Web activity.

Since this is evident from the Microsoft documentation, all other options are incorrect

For more information on Azure Data Factory service identity, please visit the following URL

https://docs.microsoft.com/en-us/azure/data-factory/data-factory-service-identity

Ask our Experts

Rate this Question? ☺ ☹

View Queries

open

Question 50

Correct

Domain :Implement data storage solutions

View Case Study

You have to create the storage account that would be used to store the polling data. Which of the following would you use as the Account type?

A. **Storage**

B. **Storage V2** ✅

C. **LRS**

D. **GRS**

Explanation:

Answer – B

The storage accounts must be General Purpose V2 to take advantage of Azure Data Lake Gen 2 capabilities.

The Microsoft documentation mentions the following

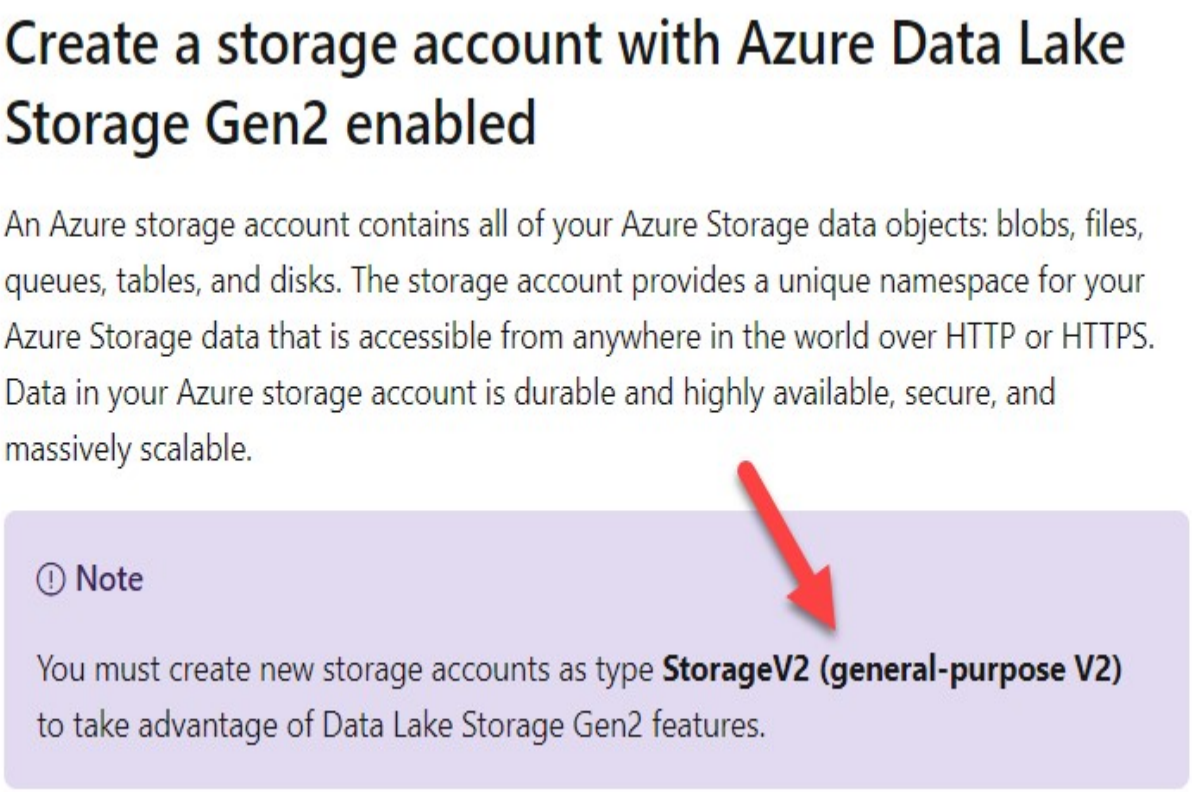

Option A is incorrect since the account kind needs to be Storage V2

Options C and D are incorrect since these are all replication strategies for storage accounts

For more information on creating an Azure Data Lake Gen2 storage account, please visit the following URL

https://docs.microsoft.com/en-us/azure/storage/blobs/data-lake-storage-quickstart-create-account

Ask our Experts

Rate this Question?

View Queries

open

Question 51

Correct

Domain :Implement data storage solutions

View Case Study

You have to create the storage account that would be used to store the polling data. Which of the following would you use as the replication type?

A. **Storage**

B. **Storage V2**

C. **LRS**

D. **GRS**

Explanation:

Answer – D

Since the case study mentions that you need to ensure that services are made available even in the event of a regional outage, you should choose GRS or Geo-redundant storage

The Microsoft documentation mentions the following

Geo-redundant storage (GRS): Cross-regional replication for Azure Storage

10/20/2018 • 5 minutes to read •

Geo-redundant storage (GRS) is designed to provide at least 99.99999999999999% (16 9's) durability of objects over a given year by replicating your data to a secondary region that is hundreds of miles away from the primary region. If your storage account has GRS enabled, then your data is durable even in the case of a complete regional outage or a disaster in which the primary region isn't recoverable.

Options A and B are incorrect since these are storage account kinds

Option C is incorrect since this replication type would not allow for data to be available in the case of a region failure

For more information on Azure Storage replication, please visit the following URL

https://docs.microsoft.com/en-us/azure/storage/common/storage-redundancy

Ask our Experts

Rate this Question?

View Queries

open

Question 52

Correct

Domain :Manage and develop data processing

View Case Study

You have to ensure Azure Data Factory would run to make the polling data available in the polling data database. Which of the following would you configure in Azure Data Factory?

A. An event-based trigger

B. A schedule-based trigger

C. A manually executed trigger

D. A function related trigger

Explanation:

Answer – B

Since you need to ensure the Data Factory pipeline runs after business hours, you need to create a schedule-based trigger

The Microsoft documentation mentions the following

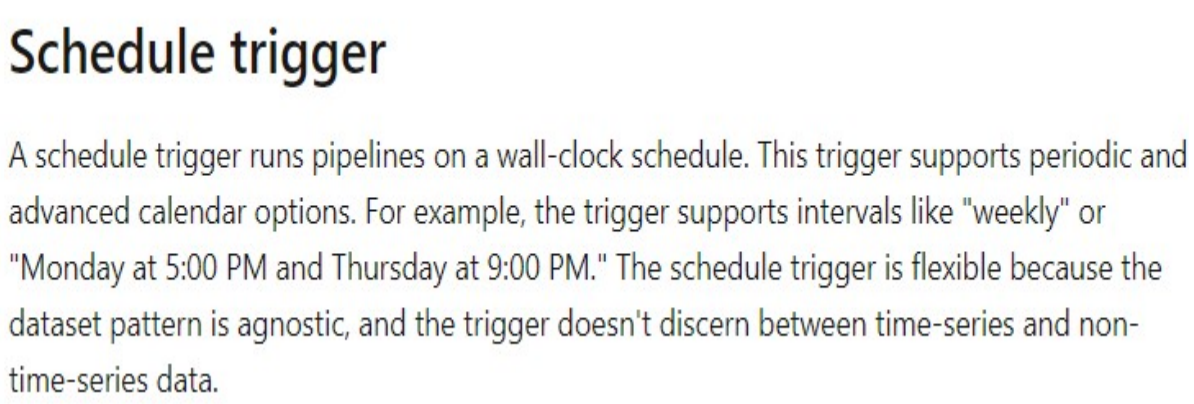

Since the requirement is given clearly in the case study, all other options are incorrect

For more information on Data Factory execution triggers, please visit the following URL

https://docs.microsoft.com/en-us/azure/data-factory/concepts-pipeline-execution-triggers

Ask our Experts

Rate this Question?

View Queries

open

Question 53 Correct

Domain :Implement data storage solutions

View Case Study

You have to ensure the polling data security requirements are met. Which of the following would you set for Polybase?

- A. **A database scoped credential** ✓
- B. **A database encryption key**
- C. **An application role**
- D. **Access Keys**

Explanation:

Answer – A

You would need to use a database scoped credential. An example of loading data from Azure Data Lake Storage is given in the Microsoft documentation

```sql
-- A: Create a Database Master Key.
-- Only necessary if one does not already exist.
-- Required to encrypt the credential secret in the next step.
-- For more information on Master Key: https://msdn.microsoft.com/library/ms174382.a

CREATE MASTER KEY;


-- B (for service principal authentication): Create a database scoped credential
-- IDENTITY: Pass the client id and OAuth 2.0 Token Endpoint taken from your Azure A
-- SECRET: Provide your AAD Application Service Principal key.
-- For more information on Create Database Scoped Credential: https://msdn.microsoft

CREATE DATABASE SCOPED CREDENTIAL ADLSCredential
WITH
    -- Always use the OAuth 2.0 authorization endpoint (v1)
    IDENTITY = '<client_id>@<OAuth_2.0_Token_EndPoint>',
    SECRET = '<key>'
;
```

Since this is clearly given in the Microsoft documentation, all other options are incorrect

For more information on the example, please visit the following URL

https://docs.microsoft.com/en-us/azure/sql-data-warehouse/sql-data-warehouse-load-from-azure-data-lake-store

Ask our Experts

Rate this Question?

View Queries

open

Question 54

Correct

Domain :Implement data storage solutions

You need to grant access to a storage account from a virtual network. Which of the following would you need to enable first for this requirement?

- A. **Create a virtual private network connection**
- B. **Enable SSL for the virtual network**
- C. **Enable a service endpoint for the storage account** ✓
- D. **Enable CORS configuration for the virtual network**

Explanation:

Answer – C

You have to enable a service endpoint for the storage account

This is also given in the Microsoft documentation

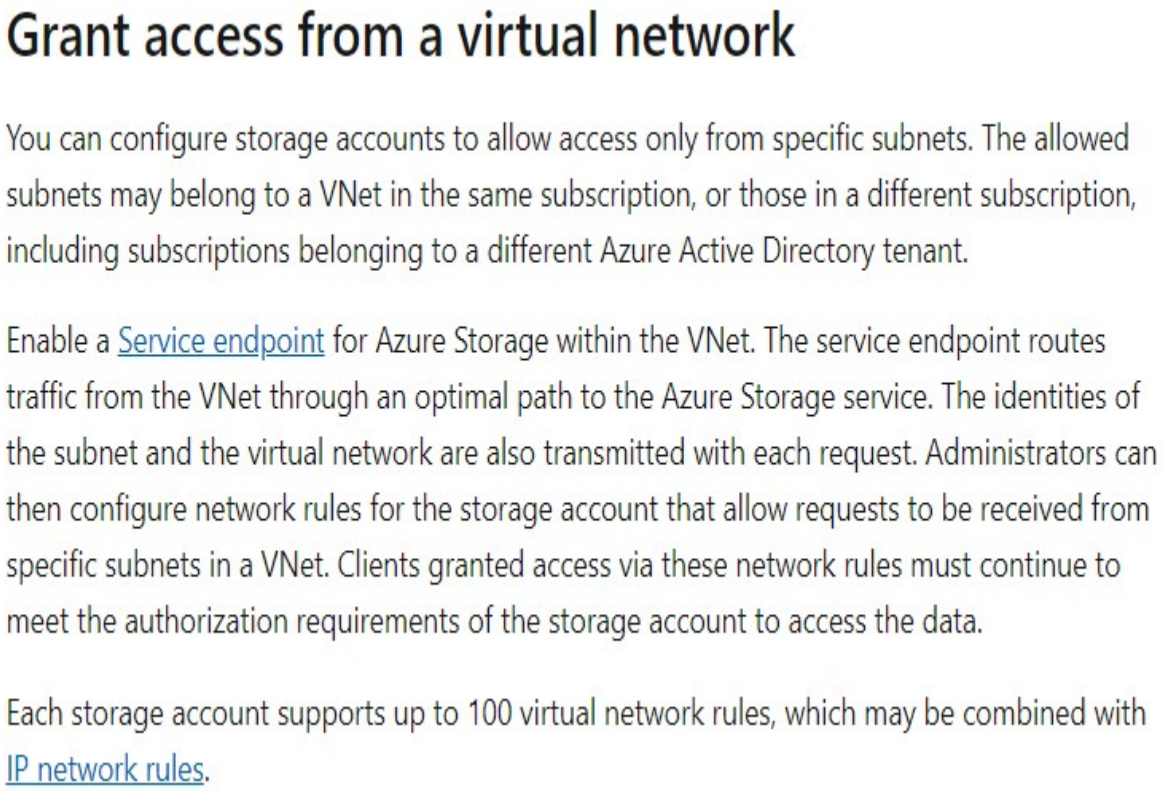

Grant access from a virtual network

You can configure storage accounts to allow access only from specific subnets. The allowed subnets may belong to a VNet in the same subscription, or those in a different subscription, including subscriptions belonging to a different Azure Active Directory tenant.

Enable a Service endpoint for Azure Storage within the VNet. The service endpoint routes traffic from the VNet through an optimal path to the Azure Storage service. The identities of the subnet and the virtual network are also transmitted with each request. Administrators can then configure network rules for the storage account that allow requests to be received from specific subnets in a VNet. Clients granted access via these network rules must continue to meet the authorization requirements of the storage account to access the data.

Each storage account supports up to 100 virtual network rules, which may be combined with IP network rules.

Since this is clearly given in the Microsoft documentation, all other options are incorrect

For more information on storage network security, please visit the following URL

https://docs.microsoft.com/en-us/azure/storage/common/storage-network-security

Ask our Experts

Rate this Question?

View Queries

open

Question 55

Correct

Domain :Implement data storage solutions

A company has a storage account named whizlabstore2020. They want to ensure that they can recover a blob object if it was deleted in the last 10 days. Which of the following would they implement for this requirement?

A. Firewalls and virtual networks

B. CORS

C. Soft Delete

D. Access Keys

Explanation:

Answer – C

They can use the Soft Delete feature

This is also given in the Microsoft documentation

How soft delete works

When enabled, soft delete enables you to save and recover your data when blobs or blob snapshots are deleted. This protection extends to blob data that is erased as the result of an overwrite.

When data is deleted, it transitions to a soft deleted state instead of being permanently erased. When soft delete is on and you overwrite data, a soft deleted snapshot is generated to save the state of the overwritten data. Soft deleted objects are invisible unless explicitly listed. You can configure the amount of time soft deleted data is recoverable before it is permanently expired.

Since this is clearly given in the Microsoft documentation, all other options are incorrect

For more information on the soft delete feature, please visit the following URL

https://docs.microsoft.com/en-us/azure/storage/blobs/storage-blob-soft-delete?tabs=azure-portal

Certification

Cloud Certification

Java Certification

PM Certification

Big Data Certification

Company

Become Our Instructor

Support

Discussions

Blog

Business

Follow us